



FACULTY OF INFORMATION TECHNOLOGY AND ELECTRICAL ENGINEERING

Syed Abdullah Hassan

**MEMORY BUS FOR MULTI-FRAGMENT LLM
WORKFLOWS: CONTEXT COMPRESSION, THE
COORDINATION CLIFF, AND PRIVACY**

Master's thesis
Degree Programme in Computer Science and Engineering
2026-06-01

Copyright © 2026 Syed Abdullah Hassan.

This work is licensed under a Creative Commons “Attribution 4.0 International” license.



Some figures are subject to different terms, such as copyright licenses and publisher permissions. These terms are indicated in the figure captions. Reproduction or modification of such figures is allowed under the conditions specified.

Hassan S. (2026) **Memory Bus for Multi-Fragment LLM Workflows: Context Compression, the Coordination Cliff, and Privacy.** University of Oulu, Faculty of Information Technology and Electrical Engineering, Degree Programme in Computer Science and Engineering, Master's thesis, 94 p.

ABSTRACT

Large language models running in multi-step workflows pay for every token they read: energy, latency, and cost scale with context length, and much of that context carries little information. Context compression is the natural lever, and the single-agent prompt-compression literature is mature. What has not been measured is whether a compressor that preserves single-agent question-answering accuracy also preserves the harder property a planner needs — combining information distributed across several compressed fragments into a correct answer. This thesis measures that on *multi-fragment LLM workflows*, using a deterministic regex solver or a single LLM call over all compressed fragments as the planner; multi-round agent simulation is outside the empirical scope.

The headline contribution is the H1 finding above. Three further contributions support it. C1 is a reproducible 150-instance multi-fragment benchmark across three workload families (cross-document fact aggregation, constraint-satisfaction planning, multi-step retrieval). C2 is the empirical characterisation itself — a sharp *coordination cliff* τ^* across four training-free compressors, together with a compounding-error model whose contribution is explaining *why* the transition is sharp (a recall threshold on critical tokens produces a step in success); position prediction is a *first-order* estimate only (median relative error $\approx 35\%$, predicted band misses the empirical τ^* on 11/11 testable cells). A direct A3 probe with an LLM planner — the only non-circular test of the threshold mechanism — confirms the step shape on the dense family ($k \approx 15$, $r_0 \approx 0,84$) and a graded transition on the distributed family ($k \approx 5,9$, band $\approx 0,54$). C3 is a three-pipeline RAG catalogue with a first-pass EUR-per-workflow cost ledger whose corpus-dominated fidelity limits the P3-vs-P1 comparison (no cost-parity claim is made). C4 is a memory-bus reference implementation together with a summary-level *inference-disclosure* metric; the metric is the C4 *research* contribution and the bus is the engineering artefact that operationalises it. CAAC operationalises the compounding-error bound as a runtime back-off and serves as a constructive demonstration of the cliff measurement rather than a fifth headline contribution.

The central result is H1: at the workload level ($n = 150$), the change in single-agent QA-F1 under compression does not positively predict the change in coordination success for any of the four compressors (Spearman $\rho \in [-0,82, +0,32]$; all 95% bootstrap CIs exclude 0,6). The single-agent benchmarks that dominate the compression literature therefore mis-rank compressors for multi-fragment use. A sharp cliff exists on 11 of 12 (compressor, family) cells. Within a defined calibrated regime the *LLM-planner* cliff shows no detected shift across a single well-identified scale-and-architecture comparison (local Llama-3.1-8B $\tau^* \approx 2,48$ vs frontier Qwen-2.5-72B $\approx 2,79$ on one cell); a third frontier model (DeepSeek V4 Pro) corroborates only weakly with a bootstrap CI too wide to identify the position. The RAG sign-flip hypothesis is not supported — compress-first is preferred over retrieve-first on the single retriever/embedder stack tested — and compression reduces inference disclosure, but in proportion to how aggressively it destroys source tokens rather

than through privacy-specific filtering (an oracle field-redaction control achieves the same ~ 23 pp reduction at near-zero compression, matching the truncation baseline at $-24,6$ pp). The H1 result on family-a should be read with one structural caveat: that family's regex solver succeeds iff enough critical numeric tokens survive compression, and the critical-token-recall metric *is* that surviving fraction, so the model's central equation on family-a relates two near-identical measurements — it is most directly *computable* on family-a, not most independently *validating*.

In the interest of honesty: of the six pre-registered hypotheses, two hold as stated (H1, H2) and the rest were reframed or not supported as originally written (the H3 sign-flip; the H5 monotonicity claim, reframed as ceiling-cliff separation; the H6 transfer claim, reframed as information-density scaling); the compounding-error model is a first-order, not a tight, position predictor. These reframings are reported explicitly. The deliverable is the manuscript and an open-source reference implementation: the local pipeline reproduces from one command given a precomputed compression cache, while the frontier-API arm is reproducible in distribution rather than byte-for-byte.

Keywords: context compression, large language models, coordination cliff, multi-fragment workflows, memory bus, inference disclosure, agentic memory, retrieval-augmented generation

Hassan S. (2026) Muistiväylä monifragmenttisille kielimallityönkuluille: kontekstin tiivistäminen, koordinaation jyrkänne ja yksityisyys. Oulun yliopisto, Tieto- ja sähkötekniikan tiedekunta, Tietotekniikan tutkinto-ohjelma, Diplomityö, 94 s.

TIIVISTELMÄ

Monivaiheisissa työnkulkutilanteissa toimivat suuret kielimallit maksavat jokaisesta lukemastaan merkistä. Kontekstin tiivistäminen on luonteva keino, mutta yksitoimijallinen tiivistyskirjallisuus ei mittaa, säilyttääkö kysymys-vastaus-tarkkuutta säilyttävä tiivistäminen myös vaikeamman ominaisuuden: kyvyn yhdistää useaan tietoaaineistoon hajautettu tieto oikeaksi vastaukseksi. Tämä diplomityö tutkii tätä kysymystä monifragmenttisissä kielimallityönkuluissa. Arvioinnissa käytetään joko determinististä säännöllisten lausekkeiden jäsentäjää tai yhtä kielimallikutsua, jossa kaikki tiivistetyt fragmentit ovat näkyvissä; monikierroksinen agenttisimulaatio on empiirisen tarkastelun ulkopuolella. Työn kontribuutiot ovat (i) toistettavissa oleva koordinaatiokokoelma, jossa on 150 instanssia kolmesta työkuormaperheestä; (ii) empiirinen luonnehdinta koordinaatiojyrkänneestä ja siihen liittyvä yhdistävä-virhe-malli, joka ennustaa jyrkänneen paikan per-kierroksisesta token-saannista, validoitu Qwen-72B- ja DeepSeek V4 Pro -malleilla kalibroidulla alueella; (iii) kolmen RAG-putken katalogi, joka osoittaa että tiivistä-ensin on vakaasti suositeltava molemmissa, sekä tallennustila- että tarkkuusrajoitetuissa olosuhteissa; ja (iv) muistiväylän viiteimplementaatio sekä yhteenvetotason päättelyilmaisumetriikka. Päätulokset: kysymys-vastaus- tarkkuus dekorreloitu koordinaatiosta; jyrkkä koordinaatio- jyrkänne on olemassa; mallin kapasiteetti määrää onnistumiskaton, ei jyrkänneen paikkaa, kalibroidulla alueella; informaatiotiheys skaalaa jyrkännettä aineistojen välillä; ja tiivistäminen vähentää mitattavasti suojattujen faktojen vuotamista. Soviteltava tiivistys (CAAC) realisoi yhdistävän-virheen rajan ajonaikaiseksi turvarajaksi.

Avainsanat: kontekstin tiivistäminen, suuret kielimallit, koordinaatiojyrkänne, monifragmenttiset työnkulut, muistiväylä, päättelyilmaisu, agenttinen muisti, hakuparannettu generointi

TABLE OF CONTENTS

ABSTRACT

TIIVISTELMÄ

TABLE OF CONTENTS

FOREWORD

LIST OF ABBREVIATIONS AND SYMBOLS

1. INTRODUCTION.....	12
1.1. Motivation: Every Token Has a Cost.....	12
1.2. The Multi-Fragment Workflow	12
1.3. Problem Statement	13
1.4. Contributions.....	14
1.4.1. Reading Guide and Contribution Hierarchy	16
1.5. What Is Novel.....	16
1.6. Thesis Outline	17
2. BACKGROUND AND RELATED WORK.....	19
2.1. The Transformer Architecture and the Cost of Context	19
2.2. The Scaling Era of LLMs	20
2.3. The Cost of Running LLMs at Scale.....	21
2.4. Context Compression --- the Breakthrough Works	22
2.4.1. Token-Level Pruning (the LLMLingua Family)	22
2.4.2. Selective and Instruction-Aware Compression	22
2.4.3. Learned Compression (Parameter-Cost Compressors)	23
2.4.4. Extractive Small-LLM Span Selection	23
2.4.5. Comparison Table	24
2.5. Multi-Agent LLM Systems and Agentic Memory	24
2.6. Retrieval-Augmented Generation and Long-Context Benchmarks	25
2.7. Privacy in Agentic Memory and Compressed Retrieval	25
2.8. The Gap This Thesis Closes	26
3. SYSTEM DESIGN AND IMPLEMENTATION	28
3.1. The Memory-Bus Reference Implementation.....	28
3.1.1. Three-Layer Architecture	28
3.1.2. Data Flow for a Write	29
3.1.3. Data Flow for a Read	29
3.1.4. Tags, Classifications, and the Audit Log.....	30
3.2. Retrieval-Augmented Generation Pipeline Catalogue	31
3.3. Compressor Framework	31
3.3.1. LLMLingua-2 (Token-Level Classifier).....	32
3.3.2. Instruction-Aware Filter	32
3.3.3. Phi-3-Mini Extractive Span Selector	33
3.3.4. Truncation (Prefix-Keep Baseline).....	33
3.3.5. Cliff-Aware Adaptive Compression (CAAC)	33
3.4. The C1 Coordination Benchmark.....	34
3.4.1. Why a Synthetic Benchmark	34
3.4.2. Three Workload Families	34
3.4.3. Tag Distributions for the Privacy Axis	35

3.4.4.	Coordination Success and Supporting Metrics.....	35
3.4.5.	The Critical-Token-Recall Metric	35
3.5.	Compression Cache	36
3.6.	Inference Backends	36
3.7.	Scope of Evaluation: the Multi-Fragment Proxy	37
3.8.	Compute Envelope and Reproducibility	37
4.	EXPERIMENT DESIGN AND PROTOCOL	39
4.1.	Experimental Protocol.....	39
4.1.1.	The Cliff Sweep	39
4.1.2.	Statistical Protocol.....	39
4.2.	H1: Question-Answering Accuracy Does Not Transferably Predict Coordination Success	41
4.2.1.	Result	41
4.2.2.	Interpretation	43
4.2.3.	Comparison to Prior Work	43
4.3.	H2: A Sharp Coordination Cliff Exists	43
4.3.1.	Result	43
4.3.2.	Interpretation	46
4.4.	The Compounding-Error Model.....	47
4.4.1.	Setup and Derivation	47
4.4.2.	Assumptions A1--A4.....	48
4.4.3.	The Calibrated Regime	48
4.4.4.	Predicted versus Empirical τ^*	49
4.4.5.	Direct A3 Probe: a Non-Circular Test of the Threshold Mechanism.....	51
4.5.	Corollary 1: Ceiling--Cliff Separation	52
4.6.	Frontier Validation Within the Calibrated Regime	54
4.6.1.	Setup and Statistical Protocol	55
4.6.2.	Qwen-2.5-72B-Instruct: In-Regime, within Bootstrap CI	55
4.6.3.	DeepSeek V4 Pro: within Tolerance but Weakly Identified	56
4.6.4.	GPT-Oss 120B: Out-Of-Regime Diagnostic	56
4.6.5.	Multi-Model Figure and the Cross-Architecture Finding	57
4.6.6.	Significance	57
4.7.	H3: RAG Pipeline Placement --- Compress-First Dominance	59
4.7.1.	Setup	59
4.7.2.	Result	59
4.7.3.	Interpretation: Why the Sign-Flip Did Not Appear	60
4.7.4.	Limitations and the Cost Model Caveat.....	60
4.8.	H4: Summary-Level Inference Disclosure	61
4.8.1.	Threat Model	61
4.8.2.	Why This Measurement Matters	62
4.8.3.	Metric Definition and the Unbiased Benchmark.....	62
4.8.4.	Result	63
4.8.5.	Construct-Validity Finding: Disclosure-Reduction Tracks Information Destruction, Not Privacy-Specific Filtering	64
4.8.6.	Interpretation	64
4.8.7.	The Reader's YES-Bias Caveat	65
4.8.8.	Memory-Bus Integration.....	66

4.8.9.	Memory-Bus Operation Benchmark	66
4.9.	Corollary 2: Information-Density Scaling on Real Benchmarks	67
4.9.1.	θ_{info} versus θ_q	67
4.9.2.	Result on Three Benchmarks.....	68
4.9.3.	Interpretation	68
4.10.	Reproducibility	69
4.11.	Summary of Verdicts	71
5.	DISCUSSION AND SUMMARY	72
5.1.	Synthesis.....	72
5.2.	Cliff-Aware Adaptive Compression: a Constructive Realisation	74
5.2.1.	Algorithm	74
5.2.2.	Operating-Point Framing	75
5.2.3.	Weak-Dominance Result and the θ_q/N_q Ablation.....	75
5.2.4.	What CAAC Contributes to the Thesis.....	76
5.3.	Limitations	76
5.4.	Significance.....	79
5.4.1.	For Practitioners.....	79
5.4.2.	For the Field	80
5.5.	Comparison to Industry Observations.....	81
5.6.	Future Work	81
5.7.	Closing	83
6.	REFERENCES	84
7.	APPENDICES	89
7.1.	Appendix A. Memory-Bus HTTP Contract.....	89
7.2.	Appendix B. Long-Format Results Schema	89
7.3.	Appendix C. Example C1 Workload (Family a)	90
7.4.	Appendix D. Reproducibility Recipe and Memory-Bus Trace	92
7.4.1.	D.1 One-Command Reproduction of Every Chapter Figure.....	92
7.4.2.	D.2 Memory-Bus End-To-End Trace	93
7.4.3.	D.3 Compute Envelope and Software Stack.....	94

FOREWORD

This thesis was carried out at the Future Computing Group, CSE Research Unit, Faculty of Information Technology and Electrical Engineering, University of Oulu, in collaboration with TalentAdore Ltd. The work was supervised academically by Lauri Lovén, with industrial supervision from Asim Nadeem and Oskari Valkama at TalentAdore.

I owe my warmest thanks to Lauri for the latitude to scope a project around an unfamiliar empirical question — the way coordination under compression behaves at institutional scale — and for the discipline he brought to keeping the scope tight. I am grateful to Asim Nadeem and Oskari Valkama for sharing TalentAdore’s production cost structure, which anchored the cost-per-workflow arm of the evaluation in a real industrial setting. I thank Mohammad Abaeiani, Faisal Khan, and Vu Truong for their willingness to coordinate around the demands of this thesis.

The thesis was prepared on a single Apple M4 Pro 48 GB workstation with an RTX 5090 32 GB GPU host on the Tailscale mesh for the compression precomputation and the model-scaling and frontier-API sweeps. The compute envelope, the reproducibility package, and the open-source reference implementation are described in Appendix 7, Section D.

Use of generative-AI tools. Every scientific claim, every experimental design choice, every results interpretation, and the final manuscript text in this thesis are the author’s own work and responsibility. Generative-AI tools — Anthropic’s Claude, accessed via the Claude Code command-line interface — were used under the author’s continuous supervision as a writing and engineering aid only: to draft and revise prose that the author then edited and verified, to assist with code review and debugging of the experiment-orchestration and analysis scripts, and to cross-check verdict-table numbers against the on-disk results CSVs. No hypothesis, experimental result, or interpretation in this thesis was generated by an AI tool without the author’s independent verification. Design decisions are recorded as architecture decision records in the project repository.

Oulu, 2026-06-01

Syed Abdullah Hassan

LIST OF ABBREVIATIONS AND SYMBOLS

ACL	access control list
API	application programming interface
AUC	area under the curve
BCa	bias-corrected and accelerated (bootstrap CI)
CAAC	cliff-aware adaptive compression
CI	confidence interval
C1 . . . C4	contributions one through four of this thesis
CSE	Computer Science and Engineering
CTR	critical-token-recall (per-family task-relevant token recall)
EM	exact match (QA metric)
EUR	euro (currency)
FAISS	Facebook AI Similarity Search
FCG	Future Computing Group
H1 . . . H6	hypotheses one through six
HNSW	hierarchical navigable small world (vector index)
HTTP	hypertext transfer protocol
ICAE	in-context autoencoder
ITEE	Information Technology and Electrical Engineering
JSON	JavaScript Object Notation
KV cache	key–value cache (transformer attention)
LLM	large language model
LoRA	low-rank adaptation
LRU	least recently used
MLX	Apple’s native ML framework for Apple Silicon
MPS	Metal Performance Shaders (Apple PyTorch backend)
NIST AI RMF	NIST AI Risk Management Framework
OAuth	open authorisation protocol
OTel	OpenTelemetry
P1 . . . P3	retrieval–compression pipeline variants one through three
pp	percentage points
QA	question answering
RAG	retrieval-augmented generation
REST	representational state transfer
SHA-256	secure hash algorithm (256-bit)
SQL	structured query language
SQLite	a lightweight, embedded SQL database
SSE	server-sent events
TF-IDF	term frequency–inverse document frequency
TTL	time-to-live
UUID	universally unique identifier
WAL	write-ahead log

Symbols of the compounding-error model (Chapter 4, Section 4.4):

r	nominal compression ratio (source / output token count)
$q(r)$	per-round critical-token-recall at ratio r
θ_q	cliff-recall threshold (per-family) — the minimum recall required for coordination success
θ_{info}	task information-density estimate (per-task, AUC-based; distinct from θ_q)
τ^*	cliff position — the ratio at which coordination success drops below the cliff threshold
p_0	baseline coordination success at $r = 1$ (no-compression baseline)
N	number of sequential compression passes (always 1 in the empirical chapters of this thesis)

Statistical symbols:

ρ	Spearman's rank correlation coefficient
p	two-sided p -value (Holm-corrected within hypothesis family unless noted)
n_{boot}	bootstrap resample count (default 1000, BCa interval)

1. INTRODUCTION

1.1. Motivation: Every Token Has a Cost

Large language models pay for every token they read. The cost is measurable at three layers: the energy and water needed to power a prefill and a decoding pass; the latency the user pays to wait for that pass to complete; and the dollar amount the operator pays per million tokens through an API or amortises through on-premises hardware. Peer-reviewed measurements of inference-side cost are recent but consistent. Luccioni *et al.*’s 2024 *Power Hungry Processing* study reports per-inference Watt-hour numbers for a broad set of open models, and finds that the wall-power cost of a single text-generation inference call on a modestly-sized model is within an order of magnitude of the wall-power cost of an internet search [1]; Samsi *et al.* [2] benchmark the energy cost of LLM inference directly and show it scales with the number of tokens processed, and Pope *et al.* [3] quantify how serving cost grows with sequence length once the prefill dominates. The training-time carbon-and-cost accounting of Patterson *et al.* [4] established the methodology that these inference-time studies extend. When the workload is a multi-step agentic system that issues thousands of internal inference calls to answer a single user request, those per-token costs compound.

A growing fraction of the tokens a deployed LLM reads do not change the answer. They are conversational filler (“Hello”, “Thanks”), repeated system prompts that scaffold every call, restated history that re-anchors a long conversation, boilerplate disclaimers, and structural framing that an attentive reader could elide without loss. Because the peer-reviewed measurements above tie cost directly to the number of tokens processed [1–3], every such token is converted directly into energy, latency, and operating expense regardless of whether it changes the answer. As an industry anecdote consistent with this, the OpenAI compute bill for processing user politeness alone was informally estimated at “tens of millions of dollars” annually in 2025 [5]; we cite it only as corroborating colour, not as evidence.

Context compression is the natural lever. Per-token compute and per-token cost scale roughly linearly in sequence length once the prefill is past the GPU memory-bound regime; halving the input context halves the cost. The single-agent prompt-compression literature has matured rapidly over the past three years [6–11]. What has not yet been measured, and what this thesis sets out to measure, is whether the compression that preserves the standard single-agent question-answering metric — token-overlap F1 against a ground- truth answer — also preserves the harder structural property that a planner needs: combining information distributed across multiple fragments into a correct answer.

1.2. The Multi-Fragment Workflow

The setting this thesis studies is the *multi-fragment LLM workflow*: a task in which the answer cannot be read out of any single fragment in isolation and requires combining information across two or more fragments, each of which may be compressed independently before the planner sees it. Cross-document fact aggregation (sum eight numbers reported by eight different institutional systems), constraint-satisfaction planning (assign N sub-tasks to K capacity-bounded workers given specifications spread across fragments), and multi-step retrieval (chain references through a sequence of documents until a leaf answer is reached) are three representative examples; the

C1 benchmark of Chapter 3 constructs 50 instances of each. The setting is not a corner case. The FCG group at the University of Oulu’s flagship use case — *Vignette 3.7 – Period-end project reporting* [12] — gathers reporting data across eight university systems via roughly eight specialised agents, achieves an estimated 80% cloud-payload reduction, and reduces the per-report cost to approximately EUR 2,15. TalentAdore, the industry partner of this thesis, runs production workloads whose operating expenses are dominated by token cost. Anthropic’s own published report on their multi-agent research system notes that token usage explains approximately 80% of the performance variance they observe [13]; the corresponding context-engineering and context-editing reports describe an 84% token reduction on a 100-turn web-search evaluation [14, 15]. These are industry observations on industry workflows; the master’s-scale problem is the controlled cross-compressor, cross-task, cross-architecture measurement that turns the industry urgency into a technical question.

The technical question this thesis takes up does not include the multi-round agent-negotiation question that the “multi-agent” framing of the FCG and Anthropic systems would imply. The empirical evaluation of this thesis uses either a deterministic regex-based information-extraction solver or a single LLM call with all (compressed) fragments visible as the planner. An AutoGen v0.4 multi-round agent backend integrated with the memory bus is shipped in the source tree (`src/m6/orchestrator/` [16]) but is not used in any reported result. The scope decision is documented in the project repository. The motivation is that round-to-round LLM variance dominates the compression signal at the C1 compression ratios that this thesis sweeps; cleanly attributing coordination-quality drops to compression rather than to multi-round agent variance requires isolating the compressor on the critical path, which a single-LLM-call (or a deterministic regex) planner does. The memory bus is *designed for* a multi-agent integration; the experimental scope is multi-fragment task solvability. The chapter on the system design (3.7 of Chapter 3) restates this scope decision explicitly; every results chapter is calibrated against it.

Operational definition of “coordination success.” Because the title and the central coinage *coordination cliff* both invite a multi-agent reading that the experiments do not test, the operational meaning of “coordination success” is stated once here and used uniformly throughout the rest of the manuscript. **Coordination success is the binary indicator that the planner recovers the correct answer to a multi-fragment task whose answer cannot be read out of any single fragment alone.** It is therefore a measurement of *multi-fragment solvability under compression* — i.e., whether enough task-critical tokens survive compression for a planner to recover the answer when the planner must combine information distributed across two or more independently compressed fragments. It is *not* a measurement of multi-round agent negotiation, message-passing efficiency, role allocation under disagreement, or any of the other phenomena the literature usually labels “coordination.” Where the manuscript subsequently uses “coordination” or “coordination cliff,” it always refers to the multi-fragment solvability quantity defined here; the scope reminder is not repeated.

1.3. Problem Statement

How does context compression affect multi-fragment coordination quality across compressors, task structures, and planner scales — and when and how can a memory-bus service exploit that understanding to compress safely, predictably, and privately?

The question subdivides into four operationally distinct sub-questions, one per contribution. First, a benchmark question: what does a coordination-quality metric on a multi-fragment workflow look like, and how should it differ from single-agent question-answering accuracy? Second, a characterisation question: does compression introduce a *coordination cliff* — a compression ratio beyond which coordination success collapses sharply — and what determines its position? Third, a deployment question: how should compression be combined with retrieval in a retrieval-augmented-generation pipeline, and what is the cost trade-off across pipeline architectures? Fourth, a governance question: can compression be measured as a privacy lever, and can a memory-bus service expose that lever as a deployment-time enforcement mechanism?

The thesis answers all four with four artefacts, each delivered end-to-end and reproducible from the open-source repository.

1.4. Contributions

Headline. The thesis’s single most-citable result is the H1 finding (the methodological backbone of C2): on the multi-fragment benchmark shipped as C1, the change in single-agent question-answering F1 under compression does not positively predict the change in multi-fragment coordination success for *any* of the four compressors tested. The single-agent benchmarks that dominate the prompt-compression literature (LLMLingua-2, LongLLMLingua, LongBench, RULER) therefore systematically mis-rank compressors for multi-fragment deployment. The remaining contributions support this finding: C1 is the benchmark it is measured on, C2 is the structural characterisation it sits inside, C3 is a deployment-oriented catalogue, and C4 is the engineering artefact that operationalises the disclosure-side findings.

C1 — A reproducible multi-fragment coordination benchmark. *The C1 benchmark ships 150 instances across three workload families (cross-document fact aggregation, constraint-satisfaction planning, multi-step retrieval), with synthetic provenance and access-control tag distributions, four coordination-quality metrics (final task success, sub-task assignment accuracy, critical-token-recall, achieved compression ratio), and a single-command regeneration script from a fixed seed. The benchmark is the substrate on which the cliff characterisation, the model-scaling result, and the inference- disclosure measurement all run; releasing it as part of the thesis lets other groups extend the characterisation to new compressors and new planner regimes.*

C2 — Empirical characterisation of the coordination cliff and a compounding-error model that explains it. *The cliff exists and is sharp; its position shows no detectable shift across the heterogeneous planners we could test, on the cells where it is resolvable, within a defined calibrated regime.* The empirical work covers four training-free compressors — LLMLingua-2 [8], the Phi-3- Mini extractive span selector, an instruction-aware filter, and a truncation baseline — at ten compression ratios across three families and across three local planners drawn from **three different architecture families** (Qwen-2.5-1.5B-Instruct, Phi-3-Mini-3.8B-Instruct, Llama-3.1-8B-Instruct; the original “Qwen-2.5 1.5B/3.8B/8B” notation conflated three architectures and is corrected here), with a frontier consistency arm on Qwen-2.5-72B and DeepSeek V4 Pro. Because a fine-grid re-resolution shows the synthetic reference $\tau^* = 2.5$ is an artefact of the deterministic solver (Section 4.3), the load-bearing comparison is LLM-vs-LLM: local Llama-3.1-8B ($\tau^* \approx 2.48$) and frontier Qwen-2.5-72B ($\tau^* \approx 2.79$, $n = 30$) cliff together across a $9\times$ scale-up and a change of architecture family, both at $p_0 = 1.0$, while

DeepSeek V4 Pro corroborates more weakly (point estimate $\tau^* = 2,15$, bootstrap CI [1,76, 7,14] so wide the position is only weakly identified). (The local three-architecture sweep on family-c shows a 24% τ^* spread, within a 50% tolerance but above the strict 20% tolerance matching the H2 spec — necessary but not sufficient; the load-bearing evidence is the cross-architecture comparison. The local arm is itself a cross-architecture invariance result across Qwen, Phi-3, and Llama, not a within-family parameter-count sweep.) The compounding-error model derives the *threshold structure* of the cliff — why a sharp transition exists at all — from per-round **critical-token-recall** and a per-family recall threshold θ_q , and recovers the empirical position only to first order (33% of cells within $\pm 25\%$, median relative error 35%; its bootstrap band does not contain the empirical τ^* on the testable cells — §4.4.4, §5.3). The model’s contribution is the derivation of the threshold mechanism, not a tight position predictor. The same machinery supports **Corollary 2 (information-density scaling)**: the cliff *shape* varies with a task’s information density θ_{info} (an AUC-based per-task quantity, distinct from θ_q), validated on HotpotQA and MultiHopRAG against the C1 family-a synthetic reference. An out-of-regime diagnostic on GPT-oss 120B identifies extended-reasoning planners as a structurally distinct case that the model’s threshold-success assumption does not cover — which is itself a contribution, since prior compression work treats reasoning regime implicitly.

C3 — A catalogue of three retrieval-augmented- generation pipelines with a EUR-per-workflow cost model. Three pipeline architectures (compress-first, retrieve-first, joint relevance-conditional routing) are implemented end-to-end on FAISS-CPU and HNSW and evaluated under matched storage-bounded and accuracy-bounded regimes. The original H3 hypothesis predicted a sign-flip in the optimal pipeline between regimes; the empirical work shows that compress-first (P1) is preferred over retrieve-first (P2) in both regimes (by 3.2 pp F1 storage, 2.0 pp F1 accuracy) *on the single bge-large-en-v1.5 + FAISS-CPU stack evaluated here*, and the joint relevance-conditional routing (P3) scores higher F1 still (0,905 in both regimes). The compress-first-over-retrieve-first finding *does not reproduce*, on this stack, the post-retrieval-compression preference of the LongLLMLingua [7] line of work — this is a single-stack result, not a general falsification. The P3 F1 advantage must be read with an important caveat: a cost-instrumented re-run shows that P3 (and P2) route retrieved fragments *verbatim* (achieved compression ratio 1,00×) while only compress-first P1 actually compresses, and the EUR cost model is corpus-dominated and does not meter compression compute (§4.7), so the cross-pipeline F1 gap is *not* a matched-compression comparison and **no cost-parity or Pareto claim can be supported**. The honest C3 contribution is the compress-first-over-retrieve-first ordering on the tested stack plus the catalogue itself; the P3 result is reported as *verbatim-routing-preserves-F1-at-uncharged-cost*.

C4 — A memory-bus reference implementation and a summary-level inference-disclosure metric. The memory bus is a FastAPI service with a tamper-evident SQLite audit log, a policy middleware that enforces a five-tier classification lattice over a 64-bit access-control mask, an in-memory scratchpad with TTL eviction, and a FAISS-CPU vector store. The compressor abstraction is a Python protocol that all five compressor variants satisfy. The summary-level inference-disclosure metric quantifies how compression at 4× reduces the rate at which a held-out local Llama-3.1-8B reader can recover protected facts from a compressed CONFIDENTIAL fragment. The metric ranks compressors by disclosure rate at iso-ratio, and the memory bus’s policy layer operationalises the ranking as a deployment-time enforcement mechanism; Section 4.8.5 documents that the ranking is monotonic in how aggressively each compressor

destroys source tokens (truncation at the top), which calibrates the privacy interpretation against the destruction baseline. **Scope of the C4 claim.** The *research* content of C4 is the disclosure metric (H4, Section 4.8) and the oracle-redaction control; the bus *service* is the engineering artefact that operationalises it. The bus is benchmarked for single-host throughput and latency (policy check, end-to-end write/read, and audit hash-chain verification; Section 4.8.9), but its *distributed* and multi-tenant scaling — contention under concurrent writers, policy-lattice scaling, and cross-host replication — is not measured here and is named as follow-on work.

The thesis also presents **CAAC** (Cliff-Aware Adaptive Compression), a wrapper that operationalises the compounding- error bound as a runtime back-off: given a per-family recall threshold, CAAC binary-searches downward over compression ratios until the per-fragment critical-token-recall sits on the safe side of the cliff. CAAC is presented in Chapter 5 as a constructive demonstration of the compounding-error model rather than as a fifth headline contribution; the demonstration is that the cliff bound is not only descriptive but operational.

1.4.1. Reading Guide and Contribution Hierarchy

The four contributions are not equally weighted, and the reader should not treat them as a flat list. The thesis has *one* load-bearing result and three supporting artefacts, and it is worth stating the priority explicitly so that the later chapters are read with the right emphasis.

The central result is H1 (Section 4.2): under compression, single-agent question-answering accuracy does not transferably predict multi-fragment coordination success. This is the finding that changes what a practitioner or benchmark designer should do, and everything else in the thesis exists to establish, explain, or exploit it. **The coordination cliff and its compounding-error model (C2)** are the mechanism behind H1: they explain *why* compression breaks coordination sharply rather than gradually, and they are the second priority. **The C1 benchmark** is the enabling instrument — necessary for H1 and C2 to be measured at all, and released so others can reproduce and extend them — rather than a contribution that stands alone. **The RAG catalogue (C3), the inference-disclosure metric and memory bus (C4), and the CAAC wrapper** are deliberately *secondary*: each stress-tests or operationalises the central finding in a specific deployment setting (retrieval placement, privacy governance, runtime back-off), and each is reported with its own scope limits rather than as a headline claim.

Concretely, the reader pressed for time should weight Chapter 4 Sections 4.2–4.4 (H1, the cliff, the model) and Section 4.5 (the cross-architecture consistency check) as the core, and read the RAG, disclosure, and CAAC sections as supporting evidence that the central finding survives contact with three distinct real deployment concerns. This prioritisation also orders the themes the thesis spans: compression-quality measurement is primary; retrieval, privacy, and adaptive-compression engineering are applied consequences of it.

1.5. What Is Novel

The closest published industry analogue to C2 is Anthropic’s report that token usage explains approximately 80% of the performance variance in their multi-agent research workflow [13]. That is an industry observation on a single workflow run by a single team; the C2 contribution is the controlled cross-compressor, cross-family, cross-architecture sweep with a quantitative model, an empirical match-rate measurement, and a frontier-model validation.

The closest published academic analogue is the single-agent question-answering evaluation of the LLMingua line of work [6–8] and the long-context benchmarks LongBench [17] and RULER [18]; H1 shows that those evaluations *over-report* compressor utility for multi-fragment deployment because the single-agent metric they measure does not predict coordination success. The multi-fragment-specific characterisation is the novelty.

The compounding-error model is novel in its application to multi-fragment compression evaluation. The derivation is short and the threshold-success assumption is a first-order approximation, but the combination of a quantitative position-prediction with a bootstrap-CI band on the prediction and a per-cell empirical match-rate measurement is, to the author’s knowledge, not present in the prior compression literature.

The summary-level inference-disclosure metric (C4) is distinct from prior privacy-aware retrieval work [19, 20], which protects the retrieval index from leaking; the metric this thesis contributes measures leakage *through the compressor itself*, which is a separate axis. The memory-bus reference implementation makes the metric auditable as a deployment-time governance budget rather than a per-publication number.

1.6. Thesis Outline

The remaining chapters are organised as follows.

Chapter 2 surveys the prior work the four contributions build on, organised into five strands: the transformer architecture and the cost of context, the modern training-free context-compression literature, multi-agent LLM systems, retrieval-augmented generation, and privacy in agentic memory. Each section closes by naming the gap that this thesis fills relative to the prior work it surveys.

Chapter 3 presents the artefacts the empirical evaluation runs on: the three-layer memory-bus architecture, the compressor framework, the C1 benchmark generator and scoring pipeline, the critical-token-recall metric, the compression cache, and the multi-fragment scope of the evaluation.

Chapter 4 presents the empirical work: H1 (question-answering accuracy is not a transferable predictor of coordination success); H2 (a sharp coordination cliff exists, characterised across 12 compressor-family cells); the compounding-error model that explains the cliff’s threshold structure and gives a first-order position estimate (with a predicted-versus-empirical match-rate band); Corollary 1 (no cliff-position shift detected across planner parameter count and architecture family within a defined calibrated regime, on the cells where the cliff is resolvable — carried by one well-identified frontier cell, Qwen-2.5-72B, with DeepSeek V4 Pro weakly corroborating and GPT-oss 120B reported as an out-of-regime diagnostic at the regime boundary); H3 (the RAG pipeline catalogue with the compress-first dominance result); H4 (the inference- disclosure measurement on the unbiased benchmark); and Corollary 2 (the cliff *shape* varies with task information density across three benchmarks).

Chapter 5 synthesises the findings, presents CAAC as the constructive realisation of the compounding-error bound, enumerates the limitations, discusses the practitioner and field-level significance, compares the findings to industry observations, and names the future-work directions in order of expected impact.

All artefacts — the manuscript source, the code, the configuration files, the results CSVs, and the `docker-compose.yml` that brings the memory-bus reference service up locally — are

released under a permissive licence in the open-source repository. Appendix 7 describes the reproducibility package and the one-command reproduction recipe for every headline figure.

2. BACKGROUND AND RELATED WORK

This chapter surveys the prior work the four contributions of this thesis build on. It is organised into seven strands. Section 2.1 introduces the transformer architecture and the per-token compute and memory cost that context-compression aims to reduce. Section 2.2 sketches the scaling-law trajectory and the inference-cost asymmetry that followed. Section 2.3 surveys the published energy and economic measurements of inference at scale, which underpin the motivation chapter. Section 2.4 covers the modern training-free context-compression literature. Section 2.5 covers multi-agent LLM systems and agentic memory. Section 2.6 covers retrieval-augmented generation, long-context benchmarks, and the evaluation suite that frames C3. Section 2.7 covers privacy-aware retrieval and the inference-disclosure literature. Each section closes by naming the gap that this thesis fills relative to the prior work it surveys; the consolidated gap statement is Section 2.8.

2.1. The Transformer Architecture and the Cost of Context

The artefacts of this thesis are built around large language models whose architectural foundation is the transformer [21]. The transformer’s defining mechanism is *self-attention*: each token in a sequence computes a weighted sum of representations of every other token in the same sequence, with the weights determined by the compatibility of learned query and key projections of the respective tokens. For a sequence of length N and a hidden size d , the self-attention mechanism produces N^2 compatibility scores (a quadratic-in- N memory and compute cost) and a final $N \times d$ output. The transformer block interleaves self-attention with a position-wise feed-forward network whose cost is linear in N but proportional to the square of the hidden size d . Stacking L such blocks gives total prefill compute approximately $L \cdot (N^2 d + N d^2)$ floating-point operations per forward pass and a memory footprint dominated by the key-value cache of size $L \cdot N \cdot d_{kv}$ bytes per request, which grows linearly in N and is paid for every token through the autoregressive decoding loop.

The decoder-only sub-family of the transformer architecture is the one this thesis evaluates. The original transformer [21] was encoder-decoder; the encoder-only line gave us BERT and its descendants; the decoder-only line gave us GPT-2, GPT-3, GPT-4, and the open-source families that this thesis’s planners are drawn from (Qwen, Llama, DeepSeek, Phi-3). The reasons decoder-only architectures came to dominate large-scale deployment are three. First, the training objective is uniform across all token positions (next-token prediction), which scales cleanly with data and parameters. Second, the autoregressive decoding loop allows the key-value cache to be amortised across token positions after the prefill, which makes each generated token cheap relative to the context that was loaded. Third, the in-context-learning behaviour that motivated agentic systems is a decoder-only phenomenon: the model conditions its output on a long prefix without further training. The practical implication for context compression is that the cost a deployed decoder-only model pays per query is roughly $2 \cdot P \cdot N$ floating-point operations during prefill (where P is the parameter count) and a per-decoded-token cost that is dominated by the key-value cache loaded from the prefill. Halving the prefill-length N halves both costs. This is the lever context compression operates on.

The scaling laws of large language models give one final piece of the cost picture. Kaplan *et al.* [22] showed that the loss of a transformer trained on next-token prediction follows

an approximate power-law in parameters, training tokens, and compute, with prescriptive implications for how to allocate a fixed compute budget between model size and training data. Hoffmann *et al.* [23] refined this with the *Chinchilla* prescription that parameters and training tokens should scale proportionally for compute-optimal training. The training-time scaling laws gave the field a recipe to spend training compute productively; what they did not give is a corresponding recipe to spend *inference* compute productively. Inference cost is paid per query for the deployment lifetime of the model, and at production scale it dominates training cost over a small number of months. Context compression is to inference cost what scaling laws are to training cost: a structural lever that converts a linear cost in some axis into a constant cost.

What this thesis adds. The cost framing this section provides is the motivation backdrop. What is missing in prior work is a structural characterisation of how aggressive compression interacts with the per-token cost reduction it delivers when the downstream task requires combining information across multiple compressed fragments rather than reading a single answer out of a single fragment. The cliff measurement of Chapter 4 fills that gap.

2.2. The Scaling Era of LLMs

The trajectory from GPT-2 (1.5B parameters, 2019) to GPT-3 (175B, 2020) to GPT-4 (undisclosed, 2023) to the frontier of 2025–2026 (Llama-3, Qwen-2.5, DeepSeek V3 and V4, Claude Sonnet 4.x, GPT-oss) marks a four-order-of-magnitude growth in deployed model size. Each step was accompanied by a near-linear increase in the cost per query and a corresponding increase in the demand for context-engineering techniques that control how much input the model reads at inference. As of the writing of this thesis, the frontier API price for a high-end model is approximately USD 3 per million input tokens and USD 15 per million output tokens, which corresponds to approximately EUR 2,76 and EUR 13,80 at the time of writing. The on-premises amortised marginal cost on a typical mid-range GPU deployment is roughly EUR 0,05 per million tokens once hardware is paid for. The two- to three-orders-of-magnitude gap between frontier-cloud and on-premises cost is what makes context compression an operational concern: cutting context length by half on a frontier-cloud workload saves real money per query.

The model-side trends are accompanied by a context-window trend. The 2K-context GPT-2 era gave way to 4K, 8K, 32K, 128K, and now multi-million-token contexts on the frontier. The long-context-benchmark literature (Section 2.6) documents that the model’s effective use of those long contexts is less complete than the nominal window size [18, 24]: tokens placed in the middle of a long context are recovered less reliably than tokens at the start or the end. Long contexts make the case for compression even stronger; tokens that the model will not effectively use are pure cost.

What this thesis adds. The model-scale trend motivates the question that Corollary 1 of Chapter 4 answers: does the cliff position depend on which point on the model-scale curve the planner is at? The frontier validation on Qwen-72B and DeepSeek V4 Pro is the cross-architecture answer — no, within the calibrated regime — that the scaling-era trend makes interesting in the first place.

2.3. The Cost of Running LLMs at Scale

The motivation chapter of this thesis frames the work in terms of the per-token cost an operator pays. This subsection surveys the published measurements that anchor that framing.

Energy. Luccioni, Jernite, and Strubell’s 2024 *Power Hungry Processing* study [1] measures the wall-power cost of inference for a broad set of open-weight models on a single A100 GPU server. The headline number is that text generation on a $\sim 6\text{B}$ -parameter open model costs ~ 0.5 Watt-hours per prompt-and-completion of typical length, which is within an order of magnitude of the energy cost of a single web search. The number scales near-linearly with model size, and the per-prompt energy at frontier-scale models (tens to hundreds of billions of parameters) is correspondingly larger. The study also separates the per-token amortised cost from the per-request fixed cost, and shows that for short prompts the fixed cost dominates while for long prompts the per-token cost dominates. Two further peer-reviewed studies corroborate the per-token scaling: Samsi *et al.* [2] benchmark the energy of LLM inference across model sizes and batch configurations and confirm that energy grows with the number of tokens processed, and Pope *et al.* [3] give an analytical and empirical account of how serving cost and latency scale with sequence length once the prefill dominates. The relevant implication for compression is that the operational regime context-compression operates in is the long-prompt regime, where the per-token cost is what compression directly cuts.

Token economics. Industry pricing as of the writing of this thesis sits at approximately USD 3 per million input tokens for the frontier “base” tier and USD 15 per million for the frontier “output” tier, with reasoning-augmented tiers priced higher. At a steady state of one million calls per day and an average prompt of 10,000 tokens, that is USD 30,000 per day of input cost alone for a frontier model. Halving the prompt length saves USD 15,000 per day — the kind of magnitude that funds an entire compression-engineering team. The on-premises amortised marginal cost is two orders of magnitude lower per token, but the absolute hardware and operations cost of running on- premises is high enough that the same compression lever is operationally valuable. The TalentAdore production accounts that this thesis’s cost model is calibrated against confirm the pricing as of the writing of the manuscript.

Politeness as waste. The peer-reviewed per-token cost results above already establish the substantive point: every token processed — including “Hello”, “Thanks”, restated history, and repeated system prompts that do not change the answer — is converted directly into kilowatt-hours and dollars [1, 2]. As a widely-circulated industry anecdote in the same direction, Sam Altman estimated in 2025 that the compute cost of processing user politeness at OpenAI’s frontier products amounted to “tens of millions of dollars” annually [5]; this is corroborating colour rather than evidence. The thesis uses the framing as the opening rhetorical hook (Chapter 1); the substantive question of how aggressively the wasted tokens can be removed without losing the answer is what the rest of the thesis measures.

Multi-agent multiplier. Anthropic’s report on their multi-agent research system [13] observes that token usage explains approximately 80% of the performance variance they see, and that a substantial fraction of that token cost is the shared context that flows between agents in each round. The follow-on report on context engineering for AI agents [15] and the report on context editing and memory tools [14] describe an 84% token reduction on a 100-turn web-search

evaluation. These are industry observations rather than controlled measurements; what they do is establish that the cost-of-context question is operationally pressing for production multi-agent deployments.

What this thesis adds. Prior work establishes that inference cost is real, that compression cuts it, and that deployed multi-agent systems are dominated by token cost. What is missing is the structural answer to *how aggressively compression can be applied before coordination breaks*. The cliff characterisation of Chapter 4 is that answer.

2.4. Context Compression --- the Breakthrough Works

The training-free context-compression literature has matured in the past three years into four recognisable families, surveyed comprehensively by Li *et al.* [25]; the most aggressive recent variants push compression to the extreme of a single-token document representation (xRAG [26]). This subsection surveys the four families, with a comparison table at the end that maps the families against the design dimensions that matter for multi-fragment evaluation. Each family closes with a “what is missing in prior work” line that the rest of the thesis addresses.

2.4.1. Token-Level Pruning (the LLMLingua Family)

The first practical context-compression technique to see broad adoption was LLMLingua [6], which uses a small autoregressive language model to score tokens by their per-step perplexity contribution and drop those whose removal least disturbs the conditional distribution. LongLLMLingua [7] extends the coarse-to-fine ranking with a question-aware step to handle long retrieval-augmented contexts. LLMLingua-2 [8] trains an XLM-RoBERTa token classifier on data distilled from a more capable model, producing a task-agnostic compressor that is faster at inference than the original LLMLingua because it does not require running a generative model at compression time. The LLMLingua-2 classifier is the principal hard-prompt compressor of this thesis; its evaluation centres on NarrativeQA F1 within 2 pp and HotpotQA F1 within 3 pp of reported numbers [27, 28] as the calibration target.

What is missing in prior work. All three LLMLingua papers report single-agent question-answering F1 retention at modest compression ratios. None measures whether the surviving content is sufficient for a downstream planner to combine information across multiple compressed fragments into a correct answer. H1 of Chapter 4 measures exactly that and finds that the single-agent F1 metric does not predict coordination success.

2.4.2. Selective and Instruction-Aware Compression

Li *et al.*’s *Selective Context* [29] prunes tokens with low self-information under the model’s own conditional distribution; the technique requires no training but does require running a scoring forward-pass at compression time. Xu *et al.*’s *RECOMP* [30] compresses long retrieval contexts with a small task-aware extractor trained to summarise the retrieved passages for the downstream question. The instruction-aware filter of this thesis sits in the same family but takes the deliberate simplification of using a TF-IDF pruning step followed by a cross-encoder re-ranker (BAAI/bge-reranker-base) and a stop-word/short-token trim, with no LM in the

loop. The simplification gives a deterministic, training-free, low-cost baseline that the more sophisticated compressors must beat.

What is missing in prior work. The task-awareness in this literature is question-conditioned: the compressor takes the downstream question as input and prunes against it. In a multi-fragment workflow the planner’s question is not the fragment-level task; the fragment-level task is to preserve information that a planner-as-yet-unknown will combine with information from other fragments. The thesis’s instruction-aware filter and the empirical evaluation of all four compressors on multi-fragment families address this gap.

2.4.3. *Learned Compression (Parameter-Cost Compressors)*

Mu *et al.* [9] introduced *gist tokens*, special positions trained with a masked attention pattern so that the model learns to compress instruction prefixes into a small number of virtual tokens. Chevalier *et al.* [10] extended this with *AutoCompressor*, which fine-tunes a language model to process long documents segment by segment and produce summary vectors. Ge *et al.* [11] formalised the architecture as the *In-context Autoencoder (ICAE)*, which trains a LoRA-adapted encoder to produce a fixed number of memory tokens from an input context. Wang *et al.* [31] present the In-Context Former, a faster cross-attention variant. Rae *et al.* [32] introduced the Compressive Transformer architecture, which is the closest single-context analogue of the cliff bound this thesis derives.

What is missing in prior work. The learned-compression papers require training, which introduces both a training-data asymmetry between compressors (which makes cross-compressor comparison less clean) and an operational deployment cost (every new compressor needs its own fine-tune). This thesis deliberately stays training-free; the LLMingua-2 wrapper, the instruction-aware filter, the Phi-3-Mini extractive span selector, and the truncation baseline are all inference-only, which makes the cliff comparison turn on the algorithm-of-action rather than on training-budget asymmetry. The learned-compression literature could be re-evaluated under the multi-fragment cliff methodology of this thesis as a follow-on study; the thesis does not do so directly.

2.4.4. *Extractive Small-LLM Span Selection*

Recent work uses small instruction-tuned LLMs as extractive compressors that produce verbatim spans of the source rather than abstractive summaries. The candidate small LLMs are the Phi-3 family [33] and other $\sim 3\text{B}$ -parameter instruction-tuned models. Extractive compression is privacy-friendly because the compressed output is by construction a subset of the source tokens; this is one of the reasons it appears in this thesis as a compressor option.

What is missing in prior work. The published guidance on how to enforce the extractive constraint reliably is thin; small LLMs frequently insert function-word filler even when told not to. The Phi-3-Mini extractive compressor of this thesis ships a post-hoc verifier with a 15% novel-token tolerance and an LLMingua-2 fallback that together salvage outputs whose extractive fraction is above 50%; the architecture is documented in detail in Section 3.3 of Chapter 3 and is one of the engineering contributions of the thesis.

2.4.5. Comparison Table

Table 1: The compressor families this thesis surveys, along the design dimensions that matter for multi-fragment evaluation. The third column is the property H1 of Chapter 4 turns on: no prior compressor was evaluated on a controlled multi-fragment coordination sweep before this thesis.

Family	Training-free?	Task-aware?	Eval'd on multi-fragment?	Deployed in this thesis?
LLMLingua family [6–8]	yes (inference-only)	question-conditioned	no	LLMLingua-2
Selective Context, RECOMP	yes	yes	no	—
Gist Tokens / AutoCompressor / ICAE [9–11]	no (training required)	varies	no	—
Extractive small-LLM (Phi-3 family) [33]	yes (inference-only)	yes	no	Phi-3-Mini ext.
Instruction-aware filter	yes	yes	no	yes (this thesis)
Truncation (prefix-keep)	yes	no	no	yes (baseline)
This thesis (multi-fragment cliff sweep)	yes	yes	yes	four compressors

2.5. Multi-Agent LLM Systems and Agentic Memory

The multi-agent LLM-systems literature has grown rapidly since 2023. The most-cited frameworks are AutoGen [16] (an open-source framework for multi-agent conversation; appeared at the ICLR 2024 LLM-Agents Workshop rather than as a main-conference paper), MetaGPT [34] (ICLR 2024 main), CAMEL [35] (NeurIPS 2023), Reflexion [36] (NeurIPS 2023), MemGPT [37] (preprint, not yet peer-reviewed at the time of writing of this thesis), and Collaborative Memory [38] (Park et al., 2025 preprint). Each framework has its own implicit answer to the question of how shared context flows between agents and what an individual agent needs to see. None of the frameworks formally bounds the information loss introduced when that shared context is compressed.

The agentic-memory line of work [37, 39–43] addresses the related question of how to store and retrieve agent state across long horizons. The closest published academic analogue to this thesis’s memory bus is MemIndex [42], which the three-layer memory-bus architecture of Chapter 3 explicitly extends. The MemGPT-style “LLM as operating system” framing [37] is related but is concerned with single-agent long-horizon memory rather than multi-agent shared state. Mem0 [40] is a production-oriented agentic-memory service.

What is missing in prior work. The multi-agent and agentic-memory frameworks treat compression as an implicit deployment optimisation; none measures the *multi-fragment* coordination-quality cost of compression structurally or bounds the loss quantitatively. This thesis fills the multi-fragment measurement gap directly (the cliff sweep, Section 4.3, and

the inference-disclosure measurement, Section 4.8); it does *not* fill the multi-round-agent measurement gap, which the AutoGen backend in `src/m6/orchestrator/` is positioned to address as future work (Section 5.6). The memory bus of Chapter 3 is designed as the reference target for both, but the empirical evidence in this thesis is multi-fragment, not multi-round.

2.6. Retrieval-Augmented Generation and Long-Context Benchmarks

Retrieval-augmented generation as a paradigm dates to Lewis *et al.* [44]. The landmark works since include RAPTOR [45] (recursive abstractive processing for tree-organised retrieval, ICLR 2024), GraphRAG [46] (a Microsoft preprint at the time of writing, not yet at a peer-reviewed venue), HippoRAG [47, 48] (NeurIPS 2024 and follow-on), and Self-RAG [49] (ICLR 2024). The pipeline catalogue this thesis evaluates (P1, P2, P3) places compression on the retrieval critical path explicitly; the closest published analogue to the post-retrieval-compression pipeline (P2) is the LongLLMLingua configuration [7], and the closest published analogue to the joint relevance-conditional pipeline (P3) is the adaptive context-compression framework (ACC-RAG) of Guo *et al.* [50], which varies the compression rate by input complexity.

The long-context-benchmark literature ([17, 18, 24, 51–53]) measures how effectively a model uses a long context. The “lost in the middle” result [24] is the canonical observation that information placed in the middle of a long context is recovered less reliably than information at the boundaries. RULER [18] extends the needle-in-a-haystack methodology to a broader set of long-context tasks. LongBench [17] is the standard multi-task benchmark. None of these benchmarks measures the multi-fragment-coordination task structure of this thesis; they measure single-agent long-context comprehension. They are complementary to the cliff measurement rather than competitive with it: long-context benchmarks measure the capacity of the model, while the cliff measurement measures the compression-quality effect on a task structure the long-context benchmarks do not include.

The multi-hop question-answering benchmarks HotpotQA [28] (EMNLP 2018), 2WikiMultiHopQA [54] (COLING 2020), and MultiHopRAG [55] (Findings of EMNLP 2024) are the real-data benchmarks on which Corollary 2 of Chapter 4 validates the information-density-scaling claim. The transfer is qualitative: the absolute cliff positions differ between synthetic and real benchmarks because the information density θ_{info} differs, but the cliff structure (existence, sharpness as a function of θ_{info}) transfers.

What is missing in prior work. The RAG and long-context benchmark literatures measure orthogonal properties to this thesis. The cliff measurement complements them rather than competing: RAG benchmarks measure retrieval quality, long-context benchmarks measure model capacity, and the cliff measurement measures the multi-fragment compression-quality effect that neither captures. The H3 catalogue of three RAG-plus-compression pipelines under matched cost regimes fills the compression-on-the-retrieval-critical-path measurement gap.

2.7. Privacy in Agentic Memory and Compressed Retrieval

The privacy-aware-retrieval literature has two strands. The first protects the retrieval index itself, on the assumption that direct access to the index would leak protected facts. Examples include PrivacyRAG [19] and SecureRAG [20]. The second strand protects the model weights from

training-data extraction, which is a property of the training pipeline rather than of the deployed retrieval-and-compression service.

The inference-disclosure metric this thesis contributes (C4) is distinct from both: it measures leakage *through the compressor itself*, on the assumption that the retrieval index is access-controlled but the compressed summaries are exposed to the planner and to any downstream reader the planner shares them with. The threat model is intuitive: an adversary who can query the planner can ask questions about the protected fact, and the disclosure rate is the fraction of those questions on which the planner’s answer recovers the protected fact given that it has the compressed summary as evidence. The H4 measurement operationalises this in a controlled setup; the memory-bus policy layer of Chapter 3 operationalises the resulting disclosure budget as a deployment-time enforcement mechanism.

The CFedRAG framework [56] is a related piece of work on coordination of federated retrieval-augmented generation; it addresses the retrieval-index-level privacy question rather than the compressor-level disclosure question that this thesis takes up.

What is missing in prior work. No published academic work, to the author’s knowledge, measures inference disclosure through the compressor at the summary level. Privacy-aware retrieval protects the index; this thesis protects the summary. The two are complementary, and a future system would presumably combine them.

2.8. The Gap This Thesis Closes

Four sentences, each one mapping to one of the four contributions of Chapter 1.

There is no shared multi-fragment coordination benchmark probing the cliff in a controlled, reproducible setup. C1 is the benchmark.

There is no published operational bound that predicts cliff position from compressor-side and task-side measurements within a defined regime. The compounding-error model and its cross-architecture frontier validation are the bound.

There is no honest catalogue of RAG-plus-compression pipeline placement under matched cost regimes, with the cost model explicit in EUR per workflow. The C3 catalogue is the catalogue, and its honest finding is that the assumed sign-flip does not appear.

There is no published academic measurement of inference disclosure through the compressor at the summary level. The H4 measurement, integrated with the memory bus’s policy layer, is the measurement.

The closest published industry analogue to C2 is Anthropic’s multi-agent system report [13], labelled honestly as industry observation rather than controlled measurement. The closest published academic analogues are the long-context benchmark suite and the LLMLingua line of work, which measure orthogonal properties on the single-agent axis. The thesis sits at the intersection of these three threads, and contributes the controlled cross-compressor multi-fragment evaluation that the prior work does not provide.

Table 2 positions this thesis against the nearest prior systems on the single axis that separates it from them: whether the work measures the effect of compression on *multi-fragment* task success, as opposed to single-fragment question-answering, retrieval quality, or aggregate token economics.

Table 2: Positioning of this thesis against the closest prior systems. The discriminating column is the last one: no prior system measures how compression affects success on a task whose answer requires combining information across multiple independently compressed fragments.

Prior system	What it measures	What it does not measure
LLMLingua-2 [8]	Single-agent QA-F1 retention at a target ratio	Whether surviving content supports multi-fragment coordination
LongLLMLingua [7]	Post-retrieval compression for long-context QA	Cliff position; multi-fragment success; cost in EUR
RECOMP [30]	Trained extractive/abstractive compression for RAG QA	Training-free cross-compressor comparison; coordination cliff
MemIndex [42]	Distributed agent-memory storage and retrieval	Information loss introduced when shared context is compressed
Anthropic multi-agent report [13]	Token usage vs. performance variance on one production workflow	Controlled, cross-compressor, cross-architecture characterisation
This thesis	Compression vs. multi-fragment coordination success across 4 compressors, 3 families, 3+ planners	—

3. SYSTEM DESIGN AND IMPLEMENTATION

This chapter presents the artefacts the empirical evaluation rests on: the reference memory bus, the compressor framework, the multi-fragment *CI* coordination benchmark, the coordination-success and critical-token-recall metrics that the cliff analysis turns on, the precomputed compression cache that decouples compression from evaluation, and the single-command reproducibility envelope. The intent is to give a reader enough detail to reproduce the system from the open-source repository without consulting the code, and to make the design choices explicit so that the experimental results in Chapter 4 can be read against the constraints those choices imposed.

A scope reminder, stated once explicitly so that every claim in this chapter and the next is bounded by it: the *multi-fragment LLM workflows* that this thesis evaluates are not multi-agent simulations. The planner is a deterministic regex-based information-extraction solver for the cliff sweep (Hypotheses H1 and H2) and a single LLM call with all (compressed) fragments visible for the model-scaling sweep (Corollary 1, frontier validation, real-data transfer). The memory bus is *designed for* a future multi-agent integration with an AutoGen back-end available in the source tree, but its multi-agent operation is outside the empirical scope of this thesis. This boundary is documented in the project repository.

3.1. The Memory-Bus Reference Implementation

3.1.1. Three-Layer Architecture

The reference implementation organises the memory bus into three layers that match the access pattern of a fragment-passing workflow:

1. An **access layer**, exposing `read`, `write`, `subscribe`, and `audit` endpoints via FastAPI and enforcing per-request access-control checks through a policy middleware.
2. A **compression layer**, providing a pluggable Compressor protocol behind which six concrete variants live: an *identity* no-compression control, the *LLMLingua-2* [8] token classifier, an *instruction-aware filter* based on TF-IDF pruning plus a cross-encoder re-ranker, a *Phi-3-Mini extractive* span selector implemented over a local Ollama endpoint, a *truncation* prefix-keep baseline, and *CAAC* — a cliff-aware adaptive wrapper that backs off compression when per-fragment recall would violate the compounding-error model’s safety threshold (Section 3.3.5, full discussion in Chapter 5).
3. A **storage layer**, providing three protocols — `AuditLog`, `Scratchpad`, and `VectorStore` — implemented for the evaluation by SQLite in WAL mode with an append-only audit table, an in-memory time-to-live dictionary, and FAISS-CPU with an HNSW index.

Figure 1 summarises the three layers. Each component dependency in the source tree is one-directional: the compression layer does not import from the access layer, and the storage layer does not import from either. This layering means the storage backends can be replaced without touching the compression or access layers.

The three-layer structure is a direct extension of MemIndex [42], the FCG group’s baseline architecture for distributed agent memory. The two material differences are that compression is




Figure 1: Three-layer architecture of the memory-bus reference implementation. Each layer depends only on the layer below it, and each storage component is hidden behind a `Protocol` so that the storage backends can be replaced without touching the access or compression layers. CAAC is drawn inside the compression layer for convenience but is a wrapper: it composes any of the other compressors with a recall-side back-off rule (Section 3.3.5).

promoted to a first-class layer rather than treated as a downstream optimisation, and that per-slot tags are part of the data model rather than out-of-band metadata.

3.1.2. Data Flow for a Write

A single write request flows through the bus in the following order. The client sends `POST /v1/write` with a JSON body containing the fragment text, its tag vector, and an optional task hint. The policy middleware parses the requester principal from the `X-M6-Principal` header and writes it to `request.state.principal` as a `Starlette State()` attribute. The service first checks that the principal’s access-control mask is a superset of the fragment tag’s mask and that the principal’s classification is at least the fragment’s. If the check passes, it calls the active compressor’s `compress(fragment, target_ratio)` method to produce a `CompressedSlot`, places that slot in the in-memory scratchpad keyed on the slot identifier, and — if the compressor exposes an embedding — adds it to the FAISS index. Finally it appends a row to the audit log whose `event_type` is `WRITE`, `result` is `OK`, `prev_hash` is the previous row’s chain hash, and `chain_hash` is `SHA-256(prev_hash || payload_hash)`. The response carries the slot identifier, the chain hash hex-encoded, and the audit row identifier.

If the policy check fails, the bus appends a `DENY` row with the synthesised slot identifier `deny-<uuid8>`, with a 60-second per-*(subject, fragment-id)* deduplication so that an adversarial client hammering denied writes cannot inflate the audit log.

3.1.3. Data Flow for a Read

A read request flows symmetrically. The client sends `GET /v1/read/{slot_id}`, the policy middleware attaches the principal, the service looks up the slot in the scratchpad, checks the

principal–tag relationship, and either returns the slot together with the audit row identifier or raises a policy-denied or slot-not-found error. The slot-not-found path is itself audited with a 60-second per-*(subject, slot-id)* deduplication: the first read miss for a given subject and slot is appended to the audit log as a READ/ERROR row, and subsequent misses within the window short-circuit silently so that an adversarial reader cannot inflate the audit log either. Appendix 7 contains a verbatim `curl` trace exercising the write, read, and audit endpoints end-to-end against the reference service.

3.1.4. Tags, Classifications, and the Audit Log

Every fragment and every compressed slot carries a tag vector recording its provenance and access-control state. The tag vector has four fields: a `uint64` access-control mask, a five-tier classification level (`PUBLIC` < `INTERNAL` < `CONFIDENTIAL` < `RESTRICTED` < `SECRET`) drawn from enterprise data-classification practice (e.g., ISO/IEC 27001 organisational schemes); the broader risk-management framing follows the NIST AI Risk Management Framework [57], but the five-tier lattice itself is enterprise convention and not prescribed by NIST AI RMF (which defines Govern/Map/Measure/Manage functions rather than a labelled classification lattice). A tuple of source identifiers recording provenance, and a tuple of parent-slot identifiers recording which previously-compressed slots this slot inherits from. When slots are merged, the union’s access-control mask is the bitwise OR of the parents’ masks and the union’s classification is the maximum; the policy predicate is symmetric for reads and writes.

The audit log is an append-only SQLite table whose columns are listed in Table 3. An `INSTEAD OF UPDATE` and an `INSTEAD OF DELETE` trigger refuse any mutation of an existing row, and a unit test drops the trigger, edits a `payload_hash` at the byte level, and confirms that the `verify()` chain-walker reports failure.

Table 3: Schema of the SQLite `audit_log` table. Triggers prevent `UPDATE` and `DELETE` on this table at the database level.

Column	Type	Notes
<code>rowid</code>	INTEGER PK	monotonic, auto-increment
<code>slot_id</code>	TEXT NOT NULL	matches the slot identifier
<code>event_type</code>	TEXT CHECK	one of WRITE / READ / SUBSCRIBE / DENY / COMPRESS / EVICT
<code>requester_acl</code>	INTEGER	64-bit ACL mask, stored signed
<code>requester_class</code>	INTEGER	classification level 0–4
<code>result</code>	TEXT CHECK	OK / DENIED / ERROR
<code>prev_hash</code>	BLOB(32)	SHA-256 of the previous row’s <code>chain_hash</code>
<code>payload_hash</code>	BLOB(32)	SHA-256 of the event’s payload bytes
<code>chain_hash</code>	BLOB(32)	SHA-256(<code>prev_hash</code> <code>payload_hash</code>)
<code>created_at</code>	TEXT NOT NULL	ISO-8601 UTC timestamp

A second SQLite table records every paid model call: provider, model, input tokens, output tokens, EUR cost, and wallclock. The cost ledger shares the audit-log database for transactional consistency but is queried separately. Pricing constants are drawn from the financial-analysis document [58]: the frontier-cloud reference is approximately EUR 2,76 per million input tokens and EUR 13,80 per million output tokens, and the on-premises amortised marginal

cost is EUR 0,05 per million tokens. The ledger feeds the EUR-per-workflow cost model that parameterises the RAG pipeline catalogue (Section 3.2).

3.2. Retrieval-Augmented Generation Pipeline Catalogue

The C3 contribution evaluates three retrieval-augmented generation pipeline architectures that differ in where compression sits relative to retrieval. The three are described here as artefacts; the empirical comparison between them under matched cost regimes is the H3 result of Chapter 4, Section 4.7.

P1: compress, then retrieve. The corpus is compressed up-front and the FAISS index is built over the compressed representation. Retrieval and synthesis both operate on compressed text. P1 trades quality for storage efficiency: the index is small but information dropped at compression time is unrecoverable. The pre-processing baseline against which P1 is informally compared is Anthropic’s contextual-retrieval work [59], which prepends LLM-generated chunk-level explanatory context before embedding.

P2: retrieve, then compress. The corpus is indexed in full; the compressor runs as a post-retrieval node-processor over the top- k retrieved fragments before they are synthesised. This is the classical LongLLMLingua configuration [7]. P2 trades storage cost for retrieval recall: the index is the full corpus, but the synthesis context can be compressed aggressively because the retriever has already identified the relevant chunks.

P3: joint, conditional on relevance. The corpus is indexed in full. At query time, each retrieved fragment is routed into one of three branches based on its retrieval relevance score s : fragments with $s \geq \theta_{\text{high}}$ pass through verbatim; fragments with $\theta_{\text{low}} \leq s < \theta_{\text{high}}$ are compressed; fragments with $s < \theta_{\text{low}}$ are discarded. The default thresholds $\theta_{\text{high}} = 0,75$ and $\theta_{\text{low}} = 0,45$ on BAAI/bge-large-en-v1.5-cosine similarities are the analytical starting points. P3’s relevance-conditional routing is in the spirit of adaptive context-compression schemes that vary the compression decision by input relevance, such as the ACC-RAG framework of [50]; the three-way verbatim/compress/discard routing implemented here is this thesis’s own design.

The three pipelines share the same FAISS-CPU + HNSW infrastructure with BAAI/bge-large-en-v1.5 as the embedder. The cost ledger introduced above feeds the EUR-per-workflow cost model that parameterises the H3 evaluation regimes (storage-bounded with FAISS index size ≤ 100 MB, accuracy-bounded with retrieval recall@10 $\leq 0,85$).

3.3. Compressor Framework

Every compressor variant conforms to a single Python Protocol with four methods — `compress(fragment, target_ratio)` returning a `CompressedSlot`; `decompress(slot)` returning a best-effort text reconstruction; `embed(slot)` returning an optional embedding for vector-store indexing; and `model_card()` returning a structured card containing the compressor identifier, family, base model, default target ratio, and provenance hashes. The thesis ships six concrete implementations. The identity compressor is the no-compression control condition required by every experiment as a mitigation against runtime-bug confounds.

The remaining five compressors fall into four families: a token-level hard-prompt filter (LLMLingua-2); an instruction-aware heuristic filter; an extractive small-LLM span selector (Phi-3-Mini extractive); a trivial prefix-truncation baseline; and a cliff-aware adaptive wrapper (CAAC). *All five are training-free*, which is a deliberate design choice: the thesis isolates the compression mechanism from any training-data effect and lets the cross-compressor comparison turn on the algorithm rather than on training-budget asymmetry.

To fix the counting once: there are **six** Compressor implementations in total (identity + five active); the **five** active ones are all training-free; and **four** of them — LLMLingua-2, the instruction-aware filter, Phi-3-Mini extractive, and truncation — are the compressors *swept* in the H1/H2 cliff experiments, since CAAC is a wrapper around an inner compressor and is evaluated separately (Chapter 5). Wherever this manuscript says “four compressors” it means the swept four; “five” means the active five; “six” includes the identity control.

3.3.1. *LLMLingua-2 (Token-Level Classifier)*

The LLMLingua-2 wrapper is a thin adapter around the `llmlingua` package [8]. Construction takes a target compression ratio; the package is asked to retain $1/r$ of the tokens, scored by a small XLM-RoBERTa classifier distilled from a more capable teacher. The force-tokens list (sentence boundary marks) is validated against the wrapped tokenizer’s vocabulary at initialisation and unknown tokens are dropped with an explicit log line. The wrapped XLM-RoBERTa model is supported by PyTorch on both CPU and Apple’s Metal Performance Shaders backend, so the compressor runs natively on Apple Silicon without modification.

The empirical token-recall curve $q(r)$ produced by LLMLingua-2 under critical-token-recall measurement (Section 3.4.5) is smooth and monotonic across the sweep ratios used in Chapter 4, which is the property the compounding-error model presumes.

3.3.2. *Instruction-Aware Filter*

The instruction-aware filter is a pure-Python heuristic baseline. It operates in two stages. The first is a sentence-level filter that ranks sentences by their relevance to the task hint using a cross-encoder reranker (`BAAI/bge-reranker-base`) when available and a lexical-overlap fallback when not. The second is a token-level trim that drops English stop-words and short tokens (≤ 2 characters) until the target compression ratio is reached. The stop-word list is encoded explicitly so that the trim is deterministic given the task hint, the fragment, and the target ratio — a subtle bug found during the development of this thesis showed that dropping stop-words by Python set iteration order introduced run-to-run variance, which is why the list is explicit.

The filter’s reranker preserves task-keyword tokens that simpler token-level methods drop, which is why its $q(r)$ curve on the multi-step retrieval family (Section 3.4, family-c) decays gracefully rather than collapsing: chain-reference tokens like “entry” and “FINAL” survive aggressive compression.

3.3.3. *Phi-3-Mini Extractive Span Selector*

The Phi-3-Mini extractive compressor is the only thesis compressor that uses an LLM at inference time. It runs the Phi-3-Mini-3.8B instruction-tuned model [33] via a local Ollama endpoint [60] with a structured prompt that asks for verbatim contiguous spans from the source. The post-hoc verifier `Phi3ExtractiveCompressor._verify_extractive` checks the output: every token in the output must appear in the source. Two loosening of the strict verifier were necessary in practice. First, a 15% novel-token tolerance: small language models frequently insert function words (“the”, “is”, “and”) even when explicitly asked not to. Second, a post-hoc `_strip_novel_tokens` pass that removes any token not in the source after the verifier passes; this guarantees zero novel tokens in the final output even if the model produced some during generation. Outputs that fail even after stripping fall back to LLMLingua-2 at the same target ratio.

A practically important property of Phi-3-Mini extractive is its *compression ceiling*. Span-level selection plus the stripping pass plus the LLMLingua-2 fallback combine to bound the achievable compression ratio at approximately 2.5 $\times$, regardless of the requested target. This is documented as a limitation in Chapter 5 and is the reason all evaluation reports *achieved* compression ratio (the ratio actually delivered) alongside the requested target ratio.

3.3.4. *Truncation (Prefix-Keep Baseline)*

The truncation compressor keeps the first $1/r$ fraction of tokens by whitespace split. It is deterministic by construction and achieves the requested ratio exactly. Its purpose is the lower-bound argument: every more sophisticated compressor must beat truncation on the coordination metric to justify its cost.

Truncation has a peculiar property: its single-agent question-answering F1 is competitive with the more sophisticated compressors at low to moderate compression ratios because the prefix tokens are naturally representative. Its coordination success collapses earlier than the others, because the tokens removed at the tail include the task-critical content that the planner needs to combine. This dissociation is part of the H1 finding: a token-overlap quality metric over-reports the utility of even the most naive compressor.

3.3.5. *Cliff-Aware Adaptive Compression (CAAC)*

CAAC is a wrapper, not a fifth compressor family. Given an inner compressor (LLMLingua-2 in the headline experiments; any of the others in the ablation) and a per-family recall threshold θ_q , CAAC runs the inner compressor at the requested target ratio, measures the resulting critical-token-recall, and if recall falls below $\theta_q^{1/N}$ binary-searches downward for the largest ratio that keeps recall above the threshold. A `min_ratio` floor of 1.5 $\times$ prevents CAAC from abdicating to no-compression on hard fragments. The full algorithmic and empirical discussion lives in Chapter 5 together with the constructive framing that CAAC is the operational realisation of the compounding-error bound. Here it suffices to note that CAAC consumes the same Compressor protocol as the other variants, which is what makes its drop-in evaluation in Chapter 4 straightforward.

3.4. The C1 Coordination Benchmark

3.4.1. Why a Synthetic Benchmark

The empirical centrepiece of this thesis — the *coordination cliff* τ^* — is a position on a curve. Detecting and characterising a cliff requires sweeping compression ratios densely from $1\times$ to $16\times$, holding everything else constant, and reading off where coordination success crosses a threshold. Real-world benchmarks (HotpotQA, MultiHopRAG) confound the cliff with task difficulty, retrieval distribution, dataset bias, and per-instance variance. The thesis takes the position that the cliff is best characterised on a synthetic generator first, where information density can be placed by construction, and then externally validated on real benchmarks. The synthetic-first posture is responsible for the clean Chapter 4 results on H1, H2, and Corollary 1; the real-benchmark validation of Corollary 2 confirms that the structure transfers.

The C1 generator produces 150 instances across three workload families and is reproducible from a single seed. The configuration hash and the resolved configuration are recorded in the manifest alongside the generated workloads so that any output CSV can be traced back to the exact generator inputs that produced it.

3.4.2. Three Workload Families

Family-a — cross-document fact aggregation. Each instance provides eight source documents, one per system from the Vignette 3.7 institutional inventory (Publication DB, Contract DB, Moodle, Patio, Peppi, CRM, Tatu SAP, SAP Travel [12]). Each document carries one recorded-hours and one approved-budget number drawn from a seed-fixed distribution. The planner is asked to aggregate the two quantities across the eight documents. Ground truth is the arithmetic sum, deterministic from the seed; the scorer is a regex-based extraction pipeline that matches *hours: X* and *budget: EUR Y* patterns. Family-a is the densest of the three: every multi-digit number is task-critical, and $\theta_{\text{info}} \approx 0,97$ (Section 4.9.1). The family-a cliff is sharp; it is the family on which the compounding-error model is most directly testable.

Family-b — constraint-satisfaction planning. Each instance specifies N sub-tasks with per-task loads and K workers with per-worker capacities. The planner must produce an assignment of sub-tasks to workers that respects per-worker capacity and assigns every sub-task. The generator constructs a feasible solution by a capacity-respecting greedy algorithm that bumps capacities only when strictly necessary, guaranteeing that every workload has at least one feasible assignment. The scorer is a *feasibility checker* rather than an exact-match: any assignment that satisfies the constraints scores 1, so the planner is rewarded for finding *any* solution rather than memorising one. Family-b exposes a property of small LLM planners that is independent of compression: constraint tracking is hard for $\leq 8\text{B}$ -parameter models even without compression, so the family-b baseline success p_0 is low for small planners and the cliff is not detectable in that regime — a *floor effect* that Chapter 4 treats explicitly.

Family-c — multi-step retrieval. Each instance presents a linear chain of fragments where each fragment carries either a pointer to the next (*entry X*) or, at the leaf, the answer (*FINAL-XXXX*). Chain lengths range from 2 to 4; each fragment is padded with 4 to 8

distractor sentences so that the retrieval step is non-trivial. The scorer reads the planner’s output and matches against the leaf answer. Family-c is the most distributed of the three: many tokens are distractors, information density $\theta_{\text{info}} \approx 0,62$ (the canonical AUC-based estimator on the family-c $\times$ LLMLingua-2 sweep; not to be confused with the recall-threshold parameter $\theta_q^{(c)} = 0,590$ reported in Chapter 4 Section 4.4.4, which is a different quantity computed by a different method — see CONTEXT.md and Section 4.9.1), and the cliff is gradual — it is the family on which the instruction-aware filter’s reranker, which preserves chain-reference keywords, can avoid a detectable cliff entirely.

Each family contributes 50 instances. The three families together span a deliberate range of information densities so that the cliff mechanism (a recall threshold θ_q) and the cliff-position shift mechanism (information density θ_{info}) can be separated; this separation is what Chapter 4 formalises as Corollary 2.

3.4.3. Tag Distributions for the Privacy Axis

The C1 generator supports three synthetic tag distributions — *uniform*, *skewed*, and *hierarchical* — used by H4 (inference disclosure) to probe the governance-stringency dimension. The hierarchical distribution sets higher classification levels to imply strict supersets of the bits at lower levels, which matches the real-world structure of enterprise tag hierarchies (e.g., ISO/IEC 27001-style schemes; the broader risk-management framing follows NIST AI RMF [57], though the five-tier lattice itself is industry convention rather than a NIST AI RMF artefact) and lets the H4 analysis disentangle classification level from access-control mask.

3.4.4. Coordination Success and Supporting Metrics

The coordination scorer is a pure function of the workload and the planner’s output; it does not call any LLM. It returns one primary metric — a per-instance binary *coordination success* — and four supporting metrics: the F1 over extracted answer tokens versus ground truth, the achieved compression ratio ($|source|/|compressed|$), generic token-recall (the fraction of source tokens preserved in the compressed text), and *critical-token-recall* (CTR; Section 3.4.5). Aggregation across workloads $\times$ seeds is by arithmetic mean; 95% confidence intervals are computed by workload-level bootstrap, never by per-row bootstrap which double-counts the seed dimension.

The reason the scoring pipeline never calls an LLM is the same reason H1 and H2 use a deterministic regex parser as the planner: isolating the compression-quality signal from LLM run-to-run variance. The trade-off is that the cliff measured here is an upper bound on what an LLM-based scorer would measure; the trade is in favour of clean attribution rather than ecological validity, and Chapter 5 discusses the limitation.

3.4.5. The Critical-Token-Recall Metric

A standard token-recall measurement $q_{\text{token}}(r) = |T_{\text{compressed}} \cap T_{\text{source}}|/|T_{\text{source}}|$ counts every preserved token equally. This turned out, during the development of this thesis, to be a poor proxy for task-relevant information preservation: the Phi-3-Mini extractive compressor at low

ratios preserves common function words like “the” and “and” but drops some critical numeric tokens. Its token-recall is high; its coordination success is low.

The critical-token-recall (CTR) metric is the family-specific restriction:

$$q_{\text{CTR}}(r) = \frac{|T_{\text{compressed}}^{\text{crit}} \cap T_{\text{source}}^{\text{crit}}|}{|T_{\text{source}}^{\text{crit}}|},$$

where T^{crit} is the set of *task-critical* tokens for the family:

- Family-a: multi-digit numeric tokens (at least two digits, to exclude single-digit noise);
- Family-b: all numeric tokens (load and capacity values);
- Family-c: chain-reference tokens (*entry-X* patterns) and the literal *FINAL* marker.

CTR is what the compounding-error model uses as the per-round retention rate $q(r)$; CAAC uses CTR as its back-off signal; and the predicted-versus-empirical τ^* figure in Chapter 4 uses CTR to derive θ_q . Generic token-recall is reported for backwards compatibility with prior analyses but is not used in the verdict pipeline.

3.5. Compression Cache

Compression is the slowest step in the evaluation. LLMingua-2 at a single ratio over 150 workloads takes minutes; Phi-3-Mini extractive over the same range takes hours because each fragment is a ~ 11 second Ollama call including verification and possible retry. The cliff sweep requires every (compressor, ratio, workload, task-hint) cell — approximately 30,000 cells in the canonical H1/H2 protocol — and a naive nested-loop implementation that recompresses on every evaluation pass would dominate the wallclock.

The thesis ships a precomputed compression cache, `CompressionCache` and `CachedCompressor` in `src/m6/compressors/cache.py`, with cache keys (compressor, r , fragment_id, hash(task_hint)). The cache is a JSON store for portability; the task-hint hash is the first twelve characters of an MD5 so that the keys are short enough to remain readable. The cache is populated once by the `precompute_cache.py` script on the GPU server and all subsequent runs consume the JSON. This decoupling is what makes the evaluation pipeline runnable end-to-end on a laptop without the GPU: H1, H2, the frontier API sweep, and the CAAC ablation all consume the cache rather than the live compressors.

The compression cache is also a methodological device. It guarantees that every evaluation in Chapter 4 sees the exact same compressed text for a given (compressor, ratio, workload) cell, so that any difference between runs is attributable to the evaluation protocol rather than to compressor stochasticity. The cache hit rate over the canonical sweep is 100% by construction.

3.6. Inference Backends

The unified `InferenceBackend` protocol exposes a single `complete(prompt, max_tokens, temperature, stop)` method together with an optional `embed(text)` method. Three concrete backends are used in the evaluation. The first is a local Ollama endpoint [60] that hosts Qwen-2.5-Instruct at the three local sizes (1.5B, 3.8B, 8B) used in the model-scaling sweep, Phi-3-Mini-3.8B for the extractive compressor, and Llama-3.1-8B-Instruct as the local

protected-fact reader. The second is the OpenAI-compatible Featherless API used for the frontier validation (Qwen-2.5-72B-Instruct, DeepSeek V4 Pro). The third is the OpenAI API for the GPT-class frontier arm (excluded from the calibrated regime per Chapter 5, retained on disk as a noncanonical diagnostic). All API calls are routed through the cost ledger so that the EUR-per-workflow figures in Chapter 4 are auditable.

3.7. Scope of Evaluation: the Multi-Fragment Proxy

The empirical protocol uses two planners. H1 and H2 use a deterministic regex-based information-extraction solver that parses the compressed fragments for the task-critical patterns of Section 3.4.2; its purpose is to isolate the compression-quality variance from LLM run-to-run variance. Corollary 1 (the model-scaling sweep), the frontier validation, and the real-data transfer arms use a single LLM call with all (compressed) fragments visible in a flat prompt; their purpose is to demonstrate that the cliff structure measured by the deterministic solver appears identically when the planner is a non-trivial language model.

An AutoGen v0.4 multi-round agent backend exists [16] in `src/m6/orchestrator/` and integrates with the memory bus described in Section 3.1. It is not used in the empirical evaluation. The decision is documented in the project repository: the run-to-run variance introduced by multi-round agent negotiation dominates the compression signal at the C1 ratios this thesis sweeps, and the value of the AutoGen runner is to demonstrate that the memory-bus design is multi-agent-ready, not to be the planner whose cliff is measured. The boundary is that the memory bus is *designed for* multi-agent integration; the experiments *evaluate* on a multi-fragment proxy.

Coordination success, as defined in Section 3.4, refers to the structural task property that the planner must combine information drawn from multiple fragments. It does not refer to the multi-round communication quality between agents; that property is what the AutoGen backend would enable measuring in future work, and is named in Chapter 5 as part of the limitations and the future-work roadmap.

3.8. Compute Envelope and Reproducibility

The empirical work runs on two machines. The first is an Apple M4 Pro 48 GB workstation, which hosts the local Ollama planners for the cliff sweep development and the deterministic-solver pipeline, and which is sufficient for every figure and verdict in Chapters Chapter 4 and Chapter 5 when the compression cache is in place. The second is an RTX 5090 32 GB workstation (Tailscale-accessible WSL2 host) used for the compression precomputation, the model-scaling sweep at 8B, the frontier API sweep, and the HotpotQA / MultiHopRAG transfer arm. Total wallclock for the canonical evaluation pipeline (cache precompute on GPU + cliff sweep + scaling + RAG + disclosure + real-data transfer) is approximately 30 hours of GPU time and ~ 12 hours of laptop time end-to-end.

The reproducibility package is the complete repository: a `docker-compose.yml` that brings the memory-bus reference service up locally, a GitHub release tag for the manuscript and code, model and data cards under `docs/`, and a `README.md` that exposes one-command reproduction for every chapter headline figure. The filesystem layout under `results/` is fixed: one CSV per hypothesis run plus a `manifest.yaml` that records the resolved configuration, the configuration

hash, the git revision, the Python version, the platform string, and the start time. Appendix 7 describes the package contents in detail and lists every command needed to reproduce the headline figures.

4. EXPERIMENT DESIGN AND PROTOCOL

This chapter presents the empirical work that this thesis turns on: the cliff sweep on which H1, H2, and the compounding-error model are based; the model-scaling sweep reframed as Corollary 1; the frontier validation that promotes the model-independence claim from small-model evidence to a cross-architecture result inside a defined calibrated regime; the retrieval-augmented-generation pipeline catalogue (H3); the summary-level inference-disclosure measurement (H4); and the information-density scaling result on real benchmarks reframed as Corollary 2. The chapter is organised around the falsifiable claims rather than the chronology of the runs; canonical results directories under `results/` are referenced for each section so that every number is traceable to an immutable artefact on disk.

4.1. Experimental Protocol

4.1.1. *The Cliff Sweep*

The cliff sweep underlies H1, H2, and the calibration step of the compounding-error model. Four compressors — LLMingua-2, Phi-3-Mini extractive, the instruction-aware filter, and a truncation baseline — are evaluated at ten compression ratios $\{1, 2, 3, 4, 5, 6, 8, 10, 12, 16\} \times$ on the 150 C1 workloads across the three families. Five seeds per cell give a total of approximately 30,000 evaluation cells per run. The canonical results directory is `results/h1_h2_v2/`; the earlier three-compressor variant `results/h1_h2_final/` is retained on disk for backwards-compatibility checks but is not the source of the headline numbers below. The truncation baseline is the lower-bound argument: any compressor that does not beat truncation on the coordination metric is paying training or inference cost for no coordination-quality return.

Every cell records the coordination-success binary, an F1 over extracted answer tokens versus ground truth (this is what H1’s “QA accuracy” refers to), the achieved compression ratio, the generic token-recall, and the critical-token-recall described in Section 3.4.5 of Chapter 3. The compression cache (Section 3.5) guarantees that every cell sees the exact same compressed text for a given (compressor, ratio, workload) tuple, so within-cell variance is exclusively planner-side variance, not compressor stochasticity.

4.1.2. *Statistical Protocol*

The statistical protocol is fixed in advance for every hypothesis. Effect-size summaries are computed at the workload level, never the cell level: per-workload means are computed first, and the 95% bootstrap confidence intervals are constructed by resampling workloads with replacement [61], never by resampling individual seed–cell pairs. The bootstrap method is the bias-corrected and accelerated (BCa) interval (per the construction in [61]) with $n_{\text{boot}} = 1000$ resamples, except where noted otherwise (the θ_q resampling pipeline of Section 4.4.4 also uses BCa with $n_{\text{boot}} = 1000$). The two-cell comparison test is the paired Wilcoxon signed-rank test [62], chosen over the Mann–Whitney U test [63] because the same workloads appear in both conditions, which violates the independence assumption of the two-sample test. Multiple-comparisons correction is Holm’s sequentially-rejective procedure [64] applied *within each*

hypothesis's family: the 12 (compressor, family) cells of the cliff sweep are Holm-corrected jointly for H2; the three condition pairs of H4 are Holm-corrected jointly for the disclosure signal and again jointly for the disclosure reduction; the three pipelines of H3 are Holm-corrected jointly per regime. Family sizes are deliberately kept narrow because the hypotheses are independent claims and joint correction across them would be over-conservative.

The H1 verdict statistic is the *workload-level* Spearman correlation ($n = 150$ workloads, one mean Δ -pair per workload, for every compressor; Phi-3-Mini extractive's $\sim 2.5\times$ achievable-compression ceiling collapses the upper ratio range and therefore reduces its *pooled* n from $150 \times 9 = 1350$ to $150 \times 5 = 750$, but does not reduce the workload-level n , Section 3.3.3). A pooled correlation over $n = 150 \times 9$ (workload, ratio) measurements would over-state significance by roughly an order of magnitude, because the nine per-workload ratios are repeated measures rather than independent draws; that pooled n and its asymptotic p -value are therefore retained in Table 4 only as a non-inferential comparison column. The two-tailed asymptotic p for $H_0 : \rho = 0$ is reported at the workload level; the verdict rests on the BCa bootstrap CI (over workloads) excluding 0,6.

Cliff position τ^* is extracted by fitting both a piecewise- linear model and a logistic

$$f(r) = p_0 + \frac{p_\infty - p_0}{1 + e^{-k(r-\tau^*)}}$$

to each cell's coordination-success-versus-ratio curve and selecting the model with the lower root-mean-squared error. The piecewise fit is constrained to an interior 10% margin from the sweep boundaries. This prevents the optimiser from parking τ^* at $r_{\max} = 16$, a degenerate fit the unconstrained optimiser would otherwise select. 95% bootstrap CIs on τ^* are computed by resampling the underlying workload means before the curve fit.

4.2. H1: Question-Answering Accuracy Does Not Transferably Predict Coordination Success

H1. Single-agent question-answering accuracy under compression is not a transferable predictor of multi-fragment coordination success: workload-level Spearman $\rho(\Delta_{qa}, \Delta_{coord}) < 0,6$ for every compressor, with 95% bootstrap CIs excluding 0,6 from above. **Verdict: SUPPORTED.** The finding is most honestly stated as *QA-F1 change does not positively predict coordination change for any compressor at the workload level*, plus a methodological warning that pooled single-number correlations are distorted. At the workload level (the inference-valid unit), the instruction-aware filter is strongly negative ($\rho = -0,82$), Phi-3-Mini extractive moderately positive ($\rho = +0,32$), and LLMingua-2 and truncation collapse to near-zero ($\rho = +0,03$ and $+0,05$, CIs spanning zero) — even though all four look moderately correlated ($\rho \approx \pm 0,4-0,6$) when ratios are pooled. We deliberately avoid the word “decorrelate” for the filter ($|\rho| = 0,82$ is a strong correlation), but for LLMingua-2 and truncation near-zero is exactly what the workload-level data show. A per-family decomposition (Table 5) explains both effects: within most (compressor, family) cells coordination collapses uniformly across workloads, leaving no rank variance (“degenerate”); the filter’s strong overall negative value is a *between-family* artifact (within family-a it is only $-0,13$, n.s.); and the single significant within-family relationship is *negative* (LLMingua-2/family-c, $-0,46$). The pooled positive correlations for LLMingua-2 and truncation are thus inflated by pseudo-replication and between-family pooling and must not be read as evidence that QA-F1 tracks coordination.

4.2.1. Result

Table 4 reports the Spearman correlation between the change in token-overlap F1 (the “QA accuracy” metric the LongLLMingua line of work reports [7, 8]) and the change in coordination success under compression, at both the workload level (the inference-valid unit) and pooled-over-ratios (for comparison only). The signature finding is twofold. First, the workload-level correlations span from $-0,82$ (filter) to $+0,32$ (Phi-3-Mini extractive) and none reaches 0,6, with LLMingua-2 and truncation sitting essentially at zero ($+0,03$, $+0,05$). Second, the workload-level and pooled columns disagree sharply: pooling the nine ratios per workload inflates LLMingua-2 and truncation to $\approx +0,38$ — a pseudo-replication artifact, since the nine ratios on one workload are not independent. The filter’s strongly negative workload-level value is the easiest to misread and is why this section opens with the verdict: the filter’s re-ranker preserves task-keyword tokens that the token-overlap metric does not reward, so as the F1 metric goes down the coordination metric can go *up*. But the per-family decomposition (Table 5) shows even this is a between-family effect — within any single family the filter’s relationship is weak or undefined. The robust H1 statement is the negative one: token-overlap F1 change does not positively, consistently predict coordination change for any compressor at the workload level.

The contrast between the two ρ columns is itself a finding: pooling the nine ratios per workload as if independent inflates LLMingua-2 and truncation from a workload-level ≈ 0 to a pooled $\approx +0,38$, and shifts the filter from $-0,82$ to $-0,59$. The workload-level column is the inference-valid one; the verdict (all CIs below 0,6) holds in both, but only the workload-level magnitudes should be interpreted.

Table 4: H1 verdict table. Spearman correlation between Δ token-overlap F1 and Δ coordination success under compression. The workload-level ρ (one mean Δ -pair per workload, $n = 150$) is the verdict statistic. The pooled column ($n = 1350 = 150$ workloads $\times$ 9 ratios) is shown for comparison only and is **not** used for inference: the nine per-workload ratio measurements are repeated measures, not independent draws, so the pooled n and its asymptotic p over-state significance by roughly an order of magnitude. The verdict (all CIs exclude 0,6) holds in both columns. CI is BCa-bootstrap over workloads, 2000 resamples. Phi-3-Mini extractive saturates at $\sim 2,5\times$ achieved compression so its pooled n is 750 (150×5); see Section 3.3.3 and Appendix B.

Compressor	ρ (workload, n)	95% CI (workload)	ρ (pooled, n)
Instruction-aware filter	−0,818 (150)	[−0,852, −0,760]	−0,593 (1350)
LLMLingua-2	+0,026 (150)	[−0,135, +0,180]	+0,381 (1350)
Phi-3-Mini extractive	+0,315 (150)	[+0,177, +0,442]	+0,193 (750)
Truncation (baseline)	+0,051 (150)	[−0,092, +0,198]	+0,384 (1350)
All four: all workload-level CIs exclude 0,6 from above SUPPORTED			

Table 5: H1 per-family decomposition (workload-level Spearman, $n = 50$ per family; computed from `results/h1_h2_v2/sweep_results.csv`). This is the load-bearing diagnostic against a Simpson’s-paradox reading of the pooled sign. “Degenerate” marks cells where the per-workload coordination change has effectively zero variance (every workload cliffs identically), so no rank correlation is defined. The only statistically significant within-family correlation is *negative* (LLMLingua-2 on family-c).

Compressor	Family-a	Family-b	Family-c
Instruction-aware filter	−0,13 (n.s., $p = 0,35$)	degenerate	degenerate
LLMLingua-2	degenerate	degenerate	−0,46 ($p = 7 \times 10^{-4}$)
Truncation	degenerate	degenerate	−0,05 (n.s., $p = 0,72$)
Phi-3-Mini extractive	+0,16 (n.s., $p = 0,27$)	+0,24 (n.s., $p = 0,09$)	degenerate

The decomposition is the decisive evidence that the pooled compressor-level correlations are not a uniform mechanism. Within a single family at the workload level, coordination success usually collapses *uniformly* across workloads (every workload cliffs at the same ratio), leaving no rank variance — eight of the twelve cells are degenerate for this reason. Of the four cells that retain variance, three are non-significant and the **only** statistically significant within-family relationship is *negative* (LLMLingua-2 on family-c, −0,46). **Stated plainly: across the twelve (compressor, family) cells there are zero statistically-significant positive within-family correlations.** The pooled-column “positives” for LLMLingua-2 and truncation ($\rho \approx +0,38$) therefore do not correspond to any within-family signal a practitioner could exploit; they are pseudo-replication artefacts of treating the nine ratios per workload as independent draws. Crucially, the strong overall filter correlation (−0,82, Table 4) does **not** reproduce within any single family (family-a is only −0,13 and n.s.; b and c are degenerate): it is a between-family pooling effect, not a within-task law. The practical reading is the conservative one — QA-F1 change does not positively and consistently predict coordination change anywhere we can measure it within a task family, which is exactly why a single-agent F1 ranking cannot be trusted to transfer to a multi-fragment deployment.

4.2.2. Interpretation

The H1 finding is a methodological claim, not just an empirical one. The LongLLMLingua and LLMLingua-2 papers report question-answering F1 retention at $4\times$ compression [7, 8]; what this thesis shows is that F1 retention at $4\times$ does not predict whether a planner can solve a multi-fragment coordination task at the same operating point. The two metrics are not the same metric on the same axis with different noise; they *disagree on the ranking of compressors*. For the multi-fragment workflows the FCG use cases [12] are built around, the F1 ranking would lead a practitioner to choose the wrong compressor.

The negative-correlation behaviour of the filter compressor is the mechanism behind H1. The TF-IDF and cross-encoder re-ranker stages of the filter are explicitly task-aware: they preserve tokens that are relevant to the task hint and drop tokens that are not. The F1 metric, by contrast, scores token-overlap against the *ground-truth answer*, not the task hint. So the filter drops F1-positive tokens (the answer string itself is sometimes re-ranked out) while preserving the chain-of-reasoning tokens the planner needs to derive the answer. The result is that the same compression operation that the F1 metric scores as a quality drop is, by the coordination metric, a quality preservation. The single-agent benchmark is measuring the wrong thing.

4.2.3. Comparison to Prior Work

The closest published analogue is the family of long-context question-answering benchmarks (LongBench [17], RULER [18]) which measure single-agent F1 on compressed inputs. The H1 finding does not contradict those benchmarks — they are correct on what they measure — but shows that they are insufficient for multi-agent or multi-fragment deployment decisions. Anthropic’s industry observation that token usage explains approximately 80% of performance variance in their multi-agent research workflow [13] is consistent with H1: a metric that does not see the structural multi-fragment property of the planner’s task will give a systematically over-optimistic compressor ranking.

4.3. H2: A Sharp Coordination Cliff Exists

H2. On each (compressor, family) cell of the cliff sweep, the coordination-success curve has a sharp cliff τ^* with a relative drop of at least 30% and a paired-Wilcoxon p -value of at most 0,05 after Holm correction across cells. **Verdict: SUPPORTED** on 11 of the 12 cells; the only exception is the instruction-aware filter on family-c, where the re-ranker preserves chain-reference tokens and no cliff is detectable below $r = 16$.

4.3.1. Result

Figure 2 is the chapter’s most-quoted figure: the coordination-success curve for LLMLingua-2 on family-a, on the coarse sweep grid, falling from coordination $\approx 1,0$ at $r = 1$ to coordination $\approx 0,0$ by $r = 6$. On this grid the logistic fit places the cliff at the quantised midpoint $\tau^* \approx 2,5\times$; the fine-grid re-resolution below shows the true deterministic-solver cliff sits lower, near $\tau^* \approx 1,1$ (and the LLM-planner cliff on the same cell near 2,7, Section 4.5). What the figure establishes

is the *shape*: the drop is concentrated in a narrow span of compression ratios, which is the signature the compounding-error model is built around.

τ^* -register table. Four numerical values for the family-a $\times$ LLMingua-2 cliff appear in the manuscript; they correspond to different measurement contexts and are not in conflict. Table 6 consolidates them so the chapter’s later sections can reference one number per context without re-explaining the register.

Table 6: τ^* register for the family-a $\times$ LLMingua-2 cell. The deterministic coarse-grid value 2,5 is a grid-quantised artefact; the fine-grid re-resolution $\approx 1,1$ is the true deterministic-solver cliff; the LLM-planner value $\approx 2,48\text{--}2,79$ is the deployment-relevant quantity (Section 4.5); and $\tau^* = 2,6997$ is a superseded auxiliary fit retained only in the audit-trail JSON.

Context	τ^*	Source / use
Deterministic, coarse-grid sweep	2,5	Table 7 (sweep midpoint, quantised)
Deterministic, fine-grid re-resolution	$\approx 1,1$	results/h1_h2_finegrid/
LLM planner, local 8B + frontier 72B	$\approx 2,48\text{--}2,79$	Section 4.5, deployment-relevant
Auxiliary fit on h1_h2_final	2,6997	superseded; retained in audit trail only

The deployment-relevant value is the LLM-planner range; the deterministic-solver values are conservative lower bounds, not deployment numbers.




Figure 2: The coordination cliff. LLMingua-2 on family-a, five seeds per ratio. Coordination success collapses from $\sim 1,0$ at $r = 1$ to ~ 0 by $r = 6$, with the cliff midpoint at $\tau^* \approx 2,7\times$. The shaded band is the workload-level 95% bootstrap CI. This is the canonical reference image for the H2 finding; see Figure 3 for the twelve-cell grid.

Table 7 reports the empirical τ^* , the relative drop, and the Holm-corrected p -value for each (compressor, family) cell. Eleven of the twelve cells reach the H2 threshold of a 30% relative

drop with $p < 0,05$ Holm-corrected. The one cell that does not is *filter* $\times$ *c*: the instruction-aware filter on the multi-step retrieval family. The reason is mechanical and consistent with the H1 finding — the filter’s re-ranker preserves the chain-reference tokens (*entry*, *FINAL*) that family-c’s scorer depends on, so the family-c task does not degrade across the sweep. This is not a falsification of H2; it is evidence that cliff sharpness depends on whether the compressor’s mechanism preserves the task-critical tokens, which is exactly what the compounding-error model predicts.

Table 7: H2 verdict table: empirical cliff position τ^* , observed-range drop, slope-extrapolated drop Δ , and Holm-corrected paired-Wilcoxon p for each (compressor, family) cell from `results/h1_h2_v2/`. The **observed drop** column is the intuitive verdict metric: (baseline coordination – post-cliff coordination)/baseline, bounded in $[0, 1]$, and the H2 criterion (drop $\geq 0,30$) is read in these units. The **slope drop** column is the piecewise-linear-extrapolated ratio from the fit’s slope-extrapolated boundary values (it can exceed 1 because the left extrapolation can sit slightly above the observed baseline of 1,0 and the right slightly below 0); it is retained as a secondary fit diagnostic only. Every significant cell clears the 0,30 criterion in observed units — the smallest is Phi-3-Mini ext. $\times$ a at exactly 0,30. Cells marked † have τ^* grid-quantised at the midpoint of the sweep gap $\{1, 2\}$: on these eight (compressor, family) pairs the deterministic solver scores coordination as a binary $1,0 \rightarrow 0,0$ step between $r = 1$ and $r = 2$, and the logistic fit lands at the midpoint with no further resolving power. The cliffs exist (Holm-corrected $p < 10^{-11}$ in all eight) but the precise positions $\leq 2,5$ are unresolved at the sweep’s ratio grid. Family-c filter cell does not exhibit a detectable cliff because the re-ranker preserves chain-reference tokens; this is documented behaviour rather than a falsification.

Cell	τ^*	observed drop	slope drop
LLMLingua-2 $\times$ a	2,5 †	1.00	1.61
LLMLingua-2 $\times$ b	2,5 †	1.00	1.61
LLMLingua-2 $\times$ c	6.7	0.66	1.23
Filter $\times$ a	2,5 †	1.00	1.60
Filter $\times$ b	2,5 †	1.00	1.61
Filter $\times$ c	—	0.00	0.00
Phi-3-Mini ext. $\times$ a	2,5 †	0.30	0.42
Phi-3-Mini ext. $\times$ b	2,5 †	0.72	0.97
Phi-3-Mini ext. $\times$ c	2,5 †	0.94	1.37
Truncation $\times$ a	2,5 †	1.00	1.61
Truncation $\times$ b	2,5 †	1.00	1.61
Truncation $\times$ c	4.95	0.71	1.36

Verdict: 11/12 cells significant with observed drop $\geq 0,30$ (3 of those 11 are Phi-3 cells whose τ^* positions are not num

Fine-grid re-resolution (results/h1_h2_finegrid). To test whether the $\tau^* = 2,5$ value was a measurement or an artefact, the dense-family cells were re-swept on a refined ratio grid adding $\{1,25, 1,5, 1,75, 2,5\}$ below $r = 2$. The quantised 2,5 is an artefact: the true deterministic-solver cliffs sit well below it. LLMLingua-2 $\times$ a collapses $1,0 \rightarrow 0,0$ between $r = 1$ and $r = 1,25$ ($\tau^* \approx 1,1$); Filter $\times$ a runs $1,0, 0,92, 0,46, 0,02$ at $r = 1,5, 1,75, 2,0$ ($\tau^* \approx 1,7$); Truncation $\times$ a runs $1,0, 0,92, 0,76, 0,0$ ($\tau^* \approx 1,85$). The distributed family-c cliffs are unchanged (LLMLingua-2 $\times$ c $\approx 6,7$, Truncation $\times$ c $\approx 5,0$), confirming they were never quantised. Two honest consequences follow. First, the canonical synthetic reference $\tau^* = 2,5$ used for the

frontier comparison (§4.6) is itself a coarse-grid artefact; the genuine *deterministic* family-a cliff is near 1,1. Second — and more substantively — the deterministic regex solver cliffs at $\approx 1,1$ while the LLM planners (local 8B and frontier Qwen-2.5-72B) cliff at $\approx 2,7$ on the same compressor and family: the brittle parser fails the instant a critical token is dropped, whereas a language model tolerates substantially more compression by reconstructing partial information. This solver-versus-LLM divergence is why Corollary 1 is anchored to the *LLM-planner* cliff rather than the deterministic-solver value (§4.5–§4.6); the deterministic cliff is a conservative *lower* bound on the operating point an LLM planner can sustain.

The twelve-cell grid in Figure 3 reads the result at one glance: cliffs are sharp on the dense fact-aggregation family and shallow on the distributed multi-step retrieval family, the ordering between (LLMLingua-2 / filter / Phi-3) is consistent within a family, and truncation cliffs at the lowest ratio of any compressor — which is what the lower-bound argument predicts.




Figure 3: Cliff curves for the four compressors $\times$ three families. Sharp cliffs are visible on family-a for every compressor; family-c degrades gradually for LLMLingua-2 and Phi-3-Mini extractive, and does not show a detectable cliff for the instruction-aware filter. Truncation cliffs earliest in every family, providing the lower-bound baseline.

4.3.2. Interpretation

The within-family ordering of τ^* corresponds to the ordering of the four compressors’ critical-token-recall curves $q(r)$, measured independently of the coordination metric. The compressor that drops critical tokens earliest cliffs at the lowest ratio. This is the H2 finding pointing forward: a cliff exists, its position is determined by the compressor’s mechanism of action on task-critical tokens, and the mechanism is captured by a single statistic — $q(r)$.

One comparability caveat applies to the Phi-3-Mini extractive cells. Phi-3-Mini extractive saturates at $\sim 2,5\times$ *achieved* compression regardless of the requested *target* ratio (Section 3.3.3), so above target $2,5\times$ its x-axis (target ratio) decouples from what the compressor actually delivers. Its tabulated τ^* values are therefore on a different effective x-axis from the token-level compressors, and the phi3 cliffs are better read against *achieved* ratio (Figure 3 is plotted against achieved ratio for this reason). The phi3 cells are retained in the 11/12 count because the coordination collapse is real and significant, but their cliff *positions* should not be compared numerically against the other compressors’ target-ratio positions. The next section turns the $q(r)$ observation into a quantitative bound.

4.4. The Compounding-Error Model

The cliff curves of Section 4.3 share a structural property: each has a sharp transition between a high-coordination plateau and a low-coordination floor. The compounding-error model is a short derivation whose *primary* contribution is explaining **why** that transition is sharp — a threshold on critical-token survival produces a step in success — and whose *secondary* output is a **first-order** estimate of the transition position τ^* from the compressor-side measurement $q(r)$ and a per-family recall-threshold parameter θ_q . It is not, and is not claimed to be, a tight position predictor; Section 4.4.4 quantifies the residual error directly.

4.4.1. Setup and Derivation

Let T_i^{crit} be the set of task-critical tokens in the i -th workload’s source fragments (multi-digit numerics for family-a, all numerics for family-b, chain-reference tokens for family-c; see Section 3.4.5 of Chapter 3). Let X_i be the number of these tokens that survive compression, $M_i = |T_i^{\text{crit}}|$ the total, and $q(r) = E[X_i/M_i]$ the per-round token-retention function of the compressor at ratio r , measured by the critical-token-recall metric. Let $\theta_q \in [0, 1]$ be the *cliff-recall threshold*: the minimum fraction of task-critical tokens that the planner needs to succeed. The threshold-success model is

$$\text{success}_i = \mathbf{1} \left[\frac{X_i}{M_i} \geq \theta_q \right],$$

and the bound the model produces is

$$E \left[\frac{X_i}{M_i} \right] = q(r), \tag{1}$$

$$P(\text{success} \mid r) \rightarrow \mathbf{1}[q(r) \geq \theta_q] \quad \text{in expectation}, \tag{2}$$

so the cliff position τ^* solves $q(\tau^*) = \theta_q$. When N sequential compression passes are applied, the surviving fraction is approximately $q(r)^N$ and the cliff position solves $q(\tau^*) = \theta_q^{1/N}$. **In every experiment in this thesis** $N = 1$, so the cliff equation simplifies to $q(\tau^*) = \theta_q$. The multi-pass formulation is recorded for future work.

The derivation is a paragraph; it is not a formal proof, and the artefact is named the *compounding-error model* rather than *Theorem 1* throughout the manuscript. The change is deliberate: the empirical match rate (Section 4.4.4) is 33% within $\pm 25\%$ relative error and 67% within $\pm 50\%$, which is the empirical-bound character of a model with quantified residuals rather than the formal character of a theorem with a proof. A finite- M Chernoff-bound refinement that sharpens the in-expectation step above into an explicit $\exp(-2M(q - \theta_q)^2)$ concentration bound exists in `src/m6/theory/cliff_prediction.py` as `chernoff_success_curve()`; the decision to keep the finite- M refinement out of the manuscript is deliberate, since the empirical match-rate residuals are dominated by A2 specification error rather than sampling error and the in-expectation form is the right register for the model’s actual accuracy. The codebase identifier `theorem_validation.json` is preserved for backwards compatibility with on-disk artefacts; the manuscript term is *the compounding-error model* (the mapping is documented in the project glossary).

4.4.2. Assumptions A1--A4

The bound rests on four assumptions, each of which is satisfied approximately rather than strictly. Naming them explicitly is what makes the calibrated regime predicate in Section 4.4.3 operational rather than mystical.

A1 — Round independence. The per-round retention is independent across compression passes. This is trivially satisfied at $N = 1$. A1 also has a *within-pass* reading — per-token retention is independent across token positions — which holds approximately for the token-level compressors evaluated here (LLMLingua-2, the instruction-aware filter, and the truncation baseline). It is violated by Phi-3-Mini extractive, which selects whole spans by LLM judgement, so adjacent token-survival events are correlated within a span. The phi3-extractive cells are retained in the Section 4.4.4 validation as a deliberate robustness check of the bound *outside* its formal regime rather than as a within-regime prediction; the corresponding match-rate cells should be read accordingly.

A2 — Binary token importance. Each task-critical token is equally important; non-task tokens are irrelevant. This is the strongest of the four assumptions; H4 (Section 4.8) measures graded importance and shows that A2 is a first-order simplification. The predicted-versus-empirical gap of Section 4.4.4 is the empirical price the model pays for A2.

A3 — Threshold success. Coordination success is binary at a single recall threshold θ_q . The cliff sharpness measured in Section 4.3 is what licenses A3 *operationally*; we acknowledge that this justification is circular in the strict sense — A3 is motivated by the sharp empirical cliff that A3 in turn explains. Breaking the circularity requires an independent A3 probe (directly varying the surviving fraction of task-relevant tokens rather than mediating that variation through compression), which is named as a future-work direction in Section 5.6. **A stronger, family-a-specific form of this circularity must be stated explicitly.** On family-a the planner is the regex solver, whose success criterion is “did enough task-critical numeric tokens survive for the parser to recover the sum,” and CTR (the model’s $q(r)$) is *the surviving fraction of those same tokens*. The cliff equation $q(\tau^*) = \theta_q$ on family-a therefore relates two near-identical measurements, so family-a validation primarily confirms internal consistency rather than testing a prediction. The non-tautological tests of the position estimate are the families where solver-success and CTR are mechanically decoupled (family-c, and the LLM-planner arms of Sections 4.5–4.6); the description of family-a as the family on which the model is most directly testable should be read in this light — most directly *computable*, not most independently *validating*.

A4 — Per-round retention measured by critical-token-recall. The $q(r)$ that enters the bound is the family-specific CTR rather than the generic token-recall. This is the substantive methodological commitment of the model, and the reason CTR rather than token-recall is the canonical recall metric of the verdict pipeline.

4.4.3. The Calibrated Regime

A planner is *in the calibrated regime* if both of the following hold:

1. the baseline coordination success p_0 at $r = 1$ is at least θ_q (no floor effect); and

2. the planner does not recover from sub-threshold information via extended reasoning beyond what the priors-only baseline supplies.

The comparison “ $p_0 \geq \theta_q$ ” in condition 1 deserves an explicit unit reading: p_0 is a success probability in $[0, 1]$ and θ_q is a fraction of task-critical tokens in $[0, 1]$, so they are comparable as numbers but they measure different things. Operationally the condition reads “*the planner’s no-compression baseline accuracy meets or exceeds the threshold-success model’s minimum-recall floor*” — equivalently, p_0 is above the success-detection floor that the threshold-success model would otherwise satisfy vacuously. This re-reading is what makes the comparison meaningful as a regime predicate; mathematically the inequality remains a comparison of two numbers in $[0, 1]$, not a unit-checked physical statement.

The second condition is operationalised by the priors-only baseline measurement of H4 (Section 4.8): a planner is in-regime if its accuracy at $q \rightarrow 0$ is no greater than its priors-only baseline plus a small slack. This formalisation matters because the model’s quantitative predictions only apply inside the regime; outside the regime — specifically for extended-reasoning planners that recover information via chain-of-thought — the cliff position can shift substantially, as the GPT-oss 120B diagnostic in Section 4.6 shows.

4.4.4. Predicted versus Empirical τ^*

For each (compressor, family) cell, θ_q is estimated from a training subset of the cliff sweep — specifically the recall at which coordination success first drops below 0.5 of the cell’s p_0 — via the `derive_theta()` function in `src/m6/theory/cliff_prediction.py`. Under critical-token-recall the canonical per-family estimates on the `h1_h2_v2` sweep are $\theta_q^{(a)} = 0.632$, $\theta_q^{(b)} = 0.838$, and $\theta_q^{(c)} = 0.590$. A workload-level bootstrap CI is constructed on θ_q by resampling workloads (`scripts/bootstrap_theta_q.py`, $n_{\text{boot}} = 1000$, BCa interval). The predicted τ^* is then obtained by inverting $q(r) = \theta_q$ on the empirical critical-token-recall curve. Resampling θ_q through this inverse map gives a *predicted- τ^* band* that quantifies the *sampling* uncertainty in θ_q — but not the model’s specification error. The two should not be confused: the band reports “how much would the prediction wobble if we resampled workloads”, not “what is the model’s actual accuracy”. The empirical τ^* falls *outside* this bootstrap band on 11/11 testable cells — the band is too tight relative to the model’s specification error. The match-rate paragraph below reports the actual accuracy, which is what the chapter’s verdict claims should be calibrated against.

Figure 4 overlays the predicted- τ^* band on the empirical cliff curve for the family-c LLMingua-2 cell. Family-c is shown rather than the more spectacular family-a because the cliff is gradual enough on family-c that the predicted-versus-empirical *gap* is visible — the family-a cell drops from 1.0 to 0 in a single ratio step, which the model also predicts closely, so both the prediction and the empirical curve collapse into the same near-step function and the figure carries no information. The family-c panel shows that the compounding-error model is *conservative* for family-c: it predicts the cliff at $r \approx 2.85$ while the empirical cliff sits at $r \approx 6.7$, a 58% relative under-prediction. The empirical τ^* sits *outside* the bootstrap band, as on every other cell — which is the empirical evidence that the band reflects θ_q -sampling uncertainty rather than the model’s full error budget. The gap is the second-order term that Corollary 2 explains (Section 4.9): family-c has lower information density than family-a, so the threshold-success approximation A2 (binary token importance) breaks earliest on family-c.




Figure 4: Predicted-versus-empirical τ^* for the family-c $\times$ LLMLingua-2 cell. The shaded band on the right panel is the 95% workload-level bootstrap CI on the predicted τ^* derived from CTR-based θ_q ; the dashed curve is the empirical cliff. The band reflects *sampling* uncertainty in θ_q only — the empirical τ^* falls outside the band here and on all other testable cells, which is the model’s specification error rather than a sampling artefact (see prose above). The model is conservative on family-c ($\Delta\tau^* \approx +0,64$), which Corollary 2 explains via the family-c information-density estimate $\theta_{\text{info}} \approx 0,62$ being lower than family-a’s $\theta_{\text{info}} \approx 0,97$. The numerical proximity of $\theta_{\text{info}}^{(c)} \approx 0,62$ and $\theta_q^{(c)} = 0,590$ is coincidental — the two are computed by different methods on different quantities and are not interchangeable; see Section 4.9.1.

Aggregating the predicted-versus-empirical comparison across all twelve (compressor, family) cells gives an empirical match rate of **33%** (4/12 cells) at the $|\tau_{\text{empirical}}^* - \tau_{\text{predicted}}^*|/\tau_{\text{empirical}}^* \leq 25\%$ criterion and **67%** (8/12 cells) at $\pm 50\%$. The binomial Wilson 95% confidence interval on the 4/12 strict match rate is approximately [13%, 61%] — the small denominator gives a wide interval that the prose below should not be read past. Reporting the full distribution of relative errors alongside the binary match rate: the median is 35,3%, the interquartile range is [19,2%, 45,0%] (10 cells with finite relative error), and the maximum is 57,4%. Two cells produce infinite predicted τ^* because $q(r)$ does not cross θ_q on the measured range: filter/c, where the no-cliff prediction is matched by a no-cliff empirical observation (Table 7 records no detectable cliff for filter/c); and phi3-extractive/c, where the no-cliff prediction is *not* matched — Table 7 records empirical $\tau^* = 2,5$ with $p < 10^{-11}$ for that cell, so phi3-extractive/c is counted as a MISS in the match-rate denominator alongside the other phi3-extractive cells. The four cells that pass the strict criterion (lingua2/b, filter/a, truncation/a, truncation/b) span three compressors, which is consistent with the bound being structural rather than compressor-specific. The phi3-extractive cells are flagged in Section 4.4.2 as outside A1’s formal regime (span-level selection induces correlated token-survival events); they are retained in the aggregate above as a deliberate robustness check on the bound’s behaviour outside its formal regime rather than as within-regime predictions. The headline empirical claim this paragraph underwrites is the *cliff-position dependence structure* — τ^* is controlled by the (compressor, task) pair — which Corollary 1 turns into the model-independence claim, validated below.

The match rates above are *in-sample*: the same sweep that derives θ_q via `derive_theta()` is used to score predicted-versus-empirical τ^* . (Note also that these match rates — not the bootstrap-CI band above — are what quantifies the model’s actual accuracy. The band reports θ_q -sampling uncertainty only and is too tight to contain the empirical τ^* on any cell.) A leave-one-family-out (LOO) cross-validation that derives θ_q on $n-1$ families and predicts τ^* on the held-out family — breaking the in-sample circularity — returns **25%** (3/12) at $\pm 25\%$

and **75%** (9/12) at $\pm 50\%$, with per-family LOO θ_q at $a = 0,71$, $b = 0,61$, $c = 0,73$. The LOO match rate is lower at the strict tolerance, as expected for an in-sample-trained θ_q ; the gap (33% in-sample vs 25% LOO at $\pm 25\%$) is the size of the in-sample optimism in the headline number. The primary uncertainty measure the chapter reports is the bootstrap CI on θ_q propagated through $q(r) = \theta_q$ to the predicted- τ^* band of Figure 4. The on-disk artefacts are `results/h1_h2_v2/theorem1_validation_per_family.json` (in-sample) and `results/h1_h2_v2/theorem1_validation_loo.json` (LOO).

4.4.5. *Direct A3 Probe: a Non-Circular Test of the Threshold Mechanism*

The match-rate evaluation of Section 4.4.4 is limited by the family-a circularity flagged in A3 (Section 4.4.2): on family-a, both “success” and “surviving critical-token fraction” are near-identical measurements of the same underlying quantity, so $q(\tau^*) = \theta_q$ on family-a is a tautological identity rather than a prediction. The *direct* A3 probe (`results/a3_probe/`) is the one test the chapter contains that breaks this circularity. The probe replaces the compressor with a hand-curated deletion procedure that removes a controlled fraction of the per-task critical tokens directly, then runs the LLM planner on the remainder and measures success as a function of the surviving-token fraction k/M . By construction the surviving fraction is the experimental knob, and success is read out *after* the knob is set — the two are not the same measurement.

The probe is run on two families: family-a (dense numeric aggregation) and family-c (multi-step retrieval). The fitted logistic on family-a returns $k \approx 15$ (sharp threshold) and $r_0 \approx 0,842$ with a transition band (0,2–0,8) of width $\approx 0,15$; family-c returns $k \approx 5,9$ (graded transition) and $r_0 \approx 0,538$ with band width $\approx 0,54$. The qualitative reading is the load-bearing one: **the dense family exhibits the step shape A3 posits, the distributed family does not**. A3 holds on dense aggregation tasks and degrades to a graded-success model on distributed retrieval tasks — which is also where the cliff itself is visibly less sharp (Phi-3 $\times$ c, LLMLingua-2 $\times$ c in Table 7).

There is a quantitative discrepancy worth surfacing: the probe’s family-a threshold ($r_0 \approx 0,842$) sits well *above* the CTR-derived $\theta_q \approx 0,632$ that the compounding-error model uses (see Section 4.4.4 per-family θ_q). The interpretation: compressors preferentially retain answer-adjacent tokens, so the surviving-token fraction they deliver at the cliff is biased upward relative to a uniform deletion. The model’s use of $\theta_q \approx 0,63$ on family-a therefore reflects a compressor-side selection effect, while the $r_0 \approx 0,84$ is the underlying planner-side threshold. The two numbers are not in conflict; they measure different stages of the pipeline. The model’s first-order accuracy (Section 4.4.4) is consistent with this gap being absorbed into the empirical θ_q fit.

The A3 probe is the strongest direct evidence the chapter provides for the threshold-success assumption. It is reported here in the results chapter rather than the future-work section because the qualitative finding (step on dense, graded on distributed) is the empirical justification for both the compounding-error model’s domain of applicability and its first-order character outside that domain.

4.5. Corollary 1: Ceiling--Cliff Separation

Corollary 1 (ceiling--cliff separation). The planner’s parameter count m determines its baseline coordination success $p_0(m)$, while the cliff position τ^* is determined by the compressor’s $q(r)$ and the task’s θ_q , not by the planner. When $p_0(m) < \theta_q$ a floor effect prevents detection of the cliff; when $p_0(m) \geq \theta_q$ we detect **no shift** in the cliff position with m within the calibrated regime (we report “no detected shift” rather than proven invariance). (The “ $p_0(m) \geq \theta_q$ ” comparison is between a success probability and a fraction of task-critical tokens; the operational reading and unit-bridge gloss is in Section 4.4.3.) **Verdict: SUPPORTED as a within-class, cross-scale, cross-architecture invariance of the LLM-planner cliff; consistent on the local arm; not established at the strict tolerance for the deterministic-solver cliff** (the fine-grid re-resolution of Section 4.3 shows the synthetic $\tau^* = 2,5$ is a deterministic-solver artefact, so the load-bearing comparison is LLM-vs-LLM: Llama-3.1-8B and Qwen-2.5-72B both cliff at $\approx 2,5$ – $2,8$ on family-a $\times$ LLMlingua-2 across a $9\times$ scale-up and a change of architecture — see the re-anchoring paragraph below; the original $n = 10$ Qwen run is 7,3% from the coarse-grid 2,5, now read as an approximate coincidence rather than a validation against that number; DeepSeek V4 Pro’s point estimate $\tau^* = 2,15$ is only weakly identified, bootstrap CI [1,76, 7,14], median 5,7); the local cross-architecture sweep at {Qwen-2.5-1.5B, Phi-3-Mini-3.8B, Llama-3.1-8B} shows a τ^* spread of 24% on family-c, which sits *below* the validator’s default 50% tolerance but *above* the strict 20% tolerance matching the H2 spec — the local sweep is *necessary but not sufficient* and the corollary’s binary verdict rests on the cross-architecture frontier evidence. The family-a small-model cells correctly exhibit the predicted floor effect and are excluded; the family-b cells are below threshold for all three models because constraint tracking is hard for ≤ 8 B planners independent of compression.

This corollary is the reframing of the original H5 hypothesis, which predicted strict monotonicity of τ^* across planner scales. The empirical result on family-c across the three local planners — **Qwen-2.5-1.5B-Instruct, Phi-3-Mini-3.8B-Instruct, and Llama-3.1-8B-Instruct, three different architecture families at the small/mid scale** — is $\tau_{\text{Qwen-1.5B}}^* = 3,69$, $\tau_{\text{Phi-3-3.8B}}^* = 4,68$, $\tau_{\text{Llama-3.1-8B}}^* = 4,01$ — a spread of approximately **24%** (23,91% exactly) within the bootstrap CI of each cell, which is what an invariance result looks like rather than a monotonicity result. (The local arm is therefore a *three-architecture* invariance test, not a within-family parameter-count sweep: only the 1.5B planner is Qwen; Phi-3-Mini (3.8B) and Llama-3.1 (8B) are different architecture families.) The original H5 predicate (gap $\geq 1,5$ ratio units) is not satisfied because the spread is well below 1,5; the reframed Corollary 1 predicate (invariance within the calibrated regime) is supported by the local sweep at a $\pm 50\%$ tolerance but *fails the strict $\pm 20\%$ tolerance* that matches the H2 falsifiable-claim spec. The local-model sweep is therefore *necessary but not sufficient* for Corollary 1; the load-bearing evidence is the LLM-cliff invariance re-anchored below (the re-anchoring paragraph; Section 4.6) rather than the local arm. The canonical results directories are `results/h5_final/` (local sweep) and `results/frontier_qwen72b_e2/` (the $n = 30$ Qwen re-run), with `results/frontier_deepseekv4/` retained as weak corroboration. A reader opening `results/h5_final/model_independence_20pct.json` will see `corollary1_supported: false`; that flag is *exactly* the strict-20%-tolerance local-arm result (`n_testable_families: 1`, `n_invariant_families: 0`, the other two floored), which this section already reports as failing the strict bar — it is honest about the local

arm, not a refutation of the reframed claim, whose support comes from the cross-architecture LLM-cliff comparison the validator does not evaluate.

Re-anchoring to the LLM-planner cliff (the cleanest form of the result). The fine-grid re-resolution of Section 4.3 forces a correction that, followed through, *strengthens* this corollary. The canonical synthetic reference $\tau^* = 2,5$ is a coarse-grid artefact of the **deterministic regex solver**, whose true family-a $\times$ LLMLingua-2 cliff is near 1,1; comparing a frontier LLM against that number compares two different kinds of planner. The honest comparison holds the *planner type* fixed and varies its scale and architecture. Measured as the 0,5-crossing of the coordination curve on the family-a $\times$ LLMLingua-2 cell, the **LLM-planner cliff is essentially invariant**: Llama-3.1-8B (local) cliffs at $\tau^* \approx 2,48$ and Qwen-2.5-72B (frontier, re-run at the larger $n = 30$ workloads $\times$ 5 seeds, `results/frontier_qwen72b_e2`) cliffs at $\tau^* \approx 2,79$ — a $\sim 12\%$ spread across a $9\times$ parameter scale-up *and* a change of architecture family, both at $p_0 = 1,0$ (in-regime). The deterministic solver on the same cell cliffs at $\approx 1,1$. So planner *type* moves the cliff (brittle parser 1,1 vs LLM $\approx 2,5$ – $2,8$), but scale and architecture *within the LLM class* do not. Read this way, Corollary 1 rests on a within-class, cross-scale, cross-architecture comparison rather than on a single frontier cell against an artefactual reference; the previously-cited match to “2,5” survives as an approximate coincidence between the LLM cliff and the coarse-grid deterministic value, now stated as such. (The DeepSeek V4 Pro $n = 30$ re-run was abandoned to API rate-limiting after its clean $p_0 = 1,0$ baseline; the existing $n = 10$ DeepSeek result, $\tau^* = 2,15$ with a wide CI, is retained as weak corroboration.)

Figure 5 reports the area-under-the-coordination-curve as a function of planner scale across all three families. On family-c, the three lines overlap; on family-a, the small-model lines are below the floor effect threshold and are correctly skipped from the invariance test; on family-b, all three lines are near zero for the constraint-tracking reason. The ceiling-versus-cliff structure is most cleanly visible on the joined plot of $p_0(m)$ and $\tau^*(m)$ on family-c (Figure 6): p_0 increases monotonically with m , τ^* does not.




Figure 5: Coordination-curve AUC by planner scale and family. Family-c shows three near-overlapping lines, consistent with ceiling–cliff separation. Family-a small-model cells are in the floor regime ($p_0 < \theta_q$); family-b is below threshold for all sizes due to constraint-tracking difficulty.

The honest statement of the reframing is in the title of this section. The original H5 predicate did not hold; the reframed Corollary 1 predicate, which is the structurally interesting claim, did.




Figure 6: Family-c ceiling versus cliff: baseline coordination $p_0(m)$ (solid line) increases monotonically with planner scale m , while the cliff position $\tau^*(m)$ (dashed line) is flat within the bootstrap CI. This is the simplest empirical demonstration of Corollary 1.

The reframing was made on the basis of the cliff machinery’s prediction and was registered as a design decision before the frontier validation that promotes it from a small-model result to a cross-architecture one.

4.6. Frontier Validation Within the Calibrated Regime

The frontier validation re-runs the family-a cliff sweep with the planner swapped out for a frontier model accessed via API while the compressor remains LLMLingua-2 at the same ratios. The question is the model-independence claim of Corollary 1: does the cliff position τ^* stay where the local-model cliff was, or does it shift with the planner? The three planners tested are Qwen-2.5-72B-Instruct, DeepSeek V4 Pro, and GPT-oss 120B. The first two are standard non-reasoning frontier models; the third is an extended-reasoning model and is the calibrated-regime boundary case.

Sample-size caveat (binding). The original frontier sweep runs only 10 family-a workloads $\times$ 3 seeds per cell (an extended Qwen-72B re-run later doubles this to 30×3). At this scale the bootstrap CI on τ^* is wide for every frontier model, and “CI contains the synthetic reference” is weaker evidence than a same-scale local-arm comparison would be. Every per-model number reported below should be read against this sample-size constraint; Section 5.3 names it as the binding limitation of the frontier arm and Section 5.6 names a re-run with more seeds (and a TOST equivalence test against the $\pm 20\%$ tolerance) as the direct upgrade.

4.6.1. Setup and Statistical Protocol

The original frontier sweep uses **10 family-a workloads at 6 sweep ratios with 3 seeds per cell, for a total of 180 cells per frontier model** (each canonical frontier CSV contains 180 rows; the larger $n = 30$ Qwen re-run of Section 4.5 is in `results/frontier_qwen72b_e2/`). Planner inference is via the OpenAI-compatible Featherless endpoint for Qwen-2.5-72B and DeepSeek V4 Pro and via the OpenAI Responses API for GPT-oss 120B. τ^* is fit by the same paired-fit-and-select procedure as the local sweep, and the bootstrap CI is constructed by resampling workloads ($n_{\text{boot}} = 500$). **The canonical reference for this section is the LLM-planner cliff $\tau_{\text{LLM}}^* \approx 2,7$ on the family-a $\times$ LLMLingua-2 cell** — the 0,5-crossing of the coordination curve for an in-regime language-model planner (local Llama-3.1-8B at $\approx 2,48$; see Section 4.5). The two earlier reference values are retained only for traceability and are **superseded**: the coarse-grid deterministic-solver logistic fit $\tau^* = 2,5$ (`results/h1_h2_v2/verdicts.json`) is, per the fine-grid re-resolution of Section 4.3, an artefact of the brittle regex solver (whose true cliff is near 1,1), and the auxiliary fit $\tau_{\text{aux}}^* = 2,6997$ on the older `h1_h2_final` curve appears in the verdicts JSON only because the frontier runner predates `h1_h2_v2`. Comparing a frontier *language model* against the deterministic-solver 2,5 compares two different kinds of planner; the load-bearing comparison here is therefore LLM-vs-LLM, against $\tau_{\text{LLM}}^* \approx 2,7$. The canonical results directories are `results/frontier_qwen72b_e2/` (the $n = 30$ Qwen re-run) and `results/frontier_qwen72b/` (the original $n = 10$ run) for Qwen, `results/frontier_deepseekv4/` for DeepSeek, and `results/frontier_gptoss120b/` (retained as a non-canonical diagnostic) for GPT-oss.

The small frontier sample (10 workloads $\times$ 3 seeds) is the binding statistical constraint on this section: bootstrap CIs are wide for all three models, and “CI contains the canonical reference” is a weaker test than the headline single-cell point-estimate comparison would imply. Section 5.3 names this as a sample-size limitation.

4.6.2. Qwen-2.5-72B-Instruct: In-Regime, within Bootstrap CI

The canonical $n = 30$ Qwen-2.5-72B re-run (`results/frontier_qwen72b_e2`) places the cliff at the 0,5-crossing $\tau^* \approx 2,79$; the original $n = 10$ sweep returned $\tau^* = 2,68$ (bootstrap CI [1,92, 7,23]). Against the canonical LLM-planner reference $\tau_{\text{LLM}}^* \approx 2,7$ (Llama-3.1-8B at $\approx 2,48$, Section 4.5), the Qwen-72B cliff is within $\sim 12\%$ — a same-band agreement across a $9\times$ parameter scale-up *and* a change of architecture family, with both planners at $p_0 = 1,0$ (squarely in-regime). This is the load-bearing evidence for Corollary 1: across a $9\times$ scale-up and a change of architecture family, *no cliff-position shift is detected within the LLM-planner class* (the point estimates cliff together; the formal $\pm 20\%$ equivalence test these cells do not pass is reported in Section 4.6.3, and marks the limit of what the present sample establishes). We do **not** lean on the within-Qwen-family “ $48\times$ from local 1.5B to frontier 72B” comparison, because the local 1.5B Qwen has a family-a floor effect that prevents measuring its cliff on that cell at all; the cross-architecture Llama-8B-vs-Qwen-72B comparison is the clean one. The baseline coordination $p_0(\text{Qwen-72B}) = 1,00$ is at the ceiling, which places the model squarely inside the calibrated regime: there is no floor effect to argue around. (For traceability: against the superseded coarse-grid deterministic value 2,5 the $n = 10$ point estimate is $\sim 7\%$ off, and against the auxiliary fit 2,6997 it is 0,8% off; Section 4.3 explains why both of those reference

numbers are artefacts of the regex solver rather than cliff positions a language model would exhibit.)

4.6.3. *DeepSeek V4 Pro: within Tolerance but Weakly Identified*

The DeepSeek V4 Pro sweep returns a τ^* point estimate of 2,15 — about 20% below the canonical LLM-planner reference $\tau_{\text{LLM}}^* \approx 2,7$ (and 13,8% below the superseded coarse-grid 2,5) — but with a bootstrap CI of [1,76, 7,14] (width 5,4 ratio units) and, tellingly, a bootstrap *median* of 5,7 that sits far from the 2,15 point estimate. That gap between point estimate and bootstrap median is the diagnostic: the cliff-fit is unstable under resampling (the curve admits both an early-cliff and a late-cliff fit on different resamples), so the position is only weakly identified. The honest framing is therefore: DeepSeek’s point estimate agrees with the reference within tolerance, but the agreement is not robust — CI-containment of the reference is weak evidence when the CI spans most of the swept range, and is not an equivalence test.

To replace that affirming-the-null reasoning with a proper test, the cliff position was re-bootstrapped from the retained per-row `coord_success` data (500 workload resamples, the same interior-clamped piecewise fitter the canonical pipeline uses; `results/tost_corollary1.json`, `results/tost_deepseek.json`) and a percentile-bootstrap two-one-sided-tests (TOST) procedure was run against the $\pm 20\%$ equivalence band [2,16, 3,24] around the reference $\tau_{\text{LLM}}^* \approx 2,70$. **Neither frontier cell demonstrates equivalence at $\pm 20\%$.** For Qwen-2.5-72B the 90% bootstrap interval on τ^* is [1,96, 7,23] (only 22% of resamples land inside the band); for DeepSeek V4 Pro it is [1,77, 7,13] (only 6% inside). The point estimates match the reference well (0,8% and 20% relative error), but the fits are too weakly identified for the data to *rule out* a difference larger than 20%. The correct conclusion is therefore *no cliff-position shift is detected* across scale and architecture — not that invariance is positively demonstrated. Establishing equivalence formally would require more seeds per frontier cell to tighten the bootstrap interval; this is named in Section 5.3. The thesis accordingly takes the conservative reading that the frontier evidence *is consistent with* Corollary 1 — the point estimates cliff together — while the equivalence test it does not pass marks the limit of what the current sample can establish.

4.6.4. *GPT-Oss 120B: Out-Of-Regime Diagnostic*

The GPT-oss 120B sweep returns $\tau^* = 6,62$ — 145% above the synthetic reference, and far outside any reasonable bootstrap CI on θ_q . The v1 sweep with baseline $p_0 = 0,53$ admitted a floor-effect interpretation (the regime predicate of Section 4.4.3 fails on its ceiling condition: $p_0 < \theta_q$ under the unit-bridge reading), but the v2 sweep with baseline $p_0 = 1,00$ does not. The model is unambiguously at the ceiling and is unambiguously cliffing at a substantially higher ratio than Corollary 1 predicts. This is the boundary case the calibrated regime was formalised around — and the failure is on the regime’s *second* condition (extended reasoning, Section 4.4.3), not the first.

The hypothesised mechanism is the second condition of the calibrated regime predicate (Section 4.4.3): GPT-oss is an extended-reasoning planner, and at sub-threshold token recall it can recover the task-critical information by chain-of-thought reasoning, violating assumption A3. The standard non-reasoning planners cannot. The 145% deviation is not

hidden: results/frontier_gptoss120b/ is preserved on disk as a non-canonical diagnostic that records the scoping, and the result is reported here as a *positive contribution* — a structural diagnostic of where the compounding-error model breaks. Section 5.3 names the extended-reasoning regime as the largest future-work direction the thesis opens.

4.6.5. Multi-Model Figure and the Cross-Architecture Finding

Figure 7 overlays the three frontier cliff curves with the synthetic reference band. The Qwen and DeepSeek curves overlap the reference band; the GPT-oss curve diverges visibly. Figure 8 is the per-model τ^* scatter with the calibrated-regime boundary marked explicitly.




Figure 7: Frontier-model cliff curves overlaid on the synthetic reference band. Qwen-72B and DeepSeek V4 Pro overlap the reference; GPT-oss 120B diverges. The GPT-oss curve is included explicitly to record the out-of-regime diagnostic rather than to support the Corollary 1 verdict.

4.6.6. Significance

Putting the three results together: we detect **no shift** in the LLM-planner cliff position across **three local architecture families** (Qwen-2.5-1.5B, Phi-3-Mini-3.8B, Llama-3.1-8B; 24% spread on family-c, within the 50% but not the strict 20% tolerance), across a **9× cross-architecture scale-up** to the frontier (local Llama-3.1-8B at $\tau^* \approx 2,48 \rightarrow$ frontier Qwen-2.5-72B at $\approx 2,79$ on family-a $\times$ LLMLingua-2, the one well-identified frontier cell), and — more weakly — across **two frontier vendors** (Qwen-72B $\rightarrow$ DeepSeek V4 Pro, the latter only weakly identified) — all within the calibrated regime — while the position *does* shift for the extended-reasoning regime (GPT-oss 120B). The defensible structural claim is therefore “no shift detected on the resolvable cells, carried by one well-identified frontier cell, and a visible break at the regime boundary,” not a proven invariance across all scales. This is still a substantive result —




Figure 8: Per-model τ^* against the synthetic reference. Qwen-72B sits at the reference; DeepSeek V4 Pro’s CI brackets the reference; GPT-oss 120B sits at the right tail with the *out-of-regime* annotation. The dashed envelope is the $\pm 25\%$ tolerance band against the synthetic reference.

the cliff is a property of the (compressor, task) pair, not the planner, on everything we could resolve — but it is stated as a tested-and-not-refuted consistency rather than a theorem.

For practitioners: τ^* can be measured once on a small local model and deployed at the same ratio on a frontier model without re-measuring, *provided* the frontier model is in the calibrated regime — a working hypothesis supported on the cells we could resolve, not a guarantee. For researchers: the evidence here is that the cliff behaves as a property of the (compressor, task) pair rather than of the planner, and that single-agent benchmarks that vary the model and hold the compressor fixed cannot reveal this structure. We state the cross-architecture consistency as the most promising contribution of this thesis to the compression literature *while being explicit about its evidential weight*: it rests on one well-identified frontier cell (Qwen-2.5-72B), is consistent on the local three-architecture arm, weakly corroborated by DeepSeek V4 Pro, and not yet established at the strict tolerance. Corollary 1 promotes the original local-model-only observation to a cross-architecture consistency result; confirming it as invariance requires the larger frontier sample named in Section 5.6.

4.7. H3: RAG Pipeline Placement --- Compress-First Dominance

H3 (original predicate). The optimal pipeline between compress-first (P1) and retrieve-first (P2) sign-flips between a storage-bounded and an accuracy-bounded regime. **Verdict: NOT SUPPORTED** for the predicted sign-flip. **Reframed finding:** P1 (compress-first) is robustly preferred over P2 (retrieve-first) in both regimes; P3 (joint relevance-conditional routing) leads both. The reframed result is itself a substantive finding because it does not reproduce, on the single retriever/embedder stack tested here, the post-retrieval-compression preference of the LongLLMLingua [7] line of work (a single-stack non-reproduction, not a general falsification).

4.7.1. Setup

The three pipelines (P1, P2, P3) are the catalogue described in Section 3.2 of Chapter 3. Each is run under a storage-bounded regime (FAISS index ≤ 100 MB) and an accuracy-bounded regime (retrieval recall@10 $\leq 0,85$), measured on the C1 family-a generator with both F1 score and the EUR-per-workflow cost model [58]. The canonical results directory is `results/h3_final/`.

4.7.2. Result

Table 8 reports the per-regime mean F1 and the per-query EUR cost for each pipeline. The headline observation is that the ranking $P3 > P1 > P2$ holds in *both* regimes. P1 is above P2 by 3,2 percentage points (pp) F1 in the storage-bounded regime and by 2,0 pp in the accuracy-bounded regime, with 95% paired-bootstrap CIs ($[1,84, 4,72]$ pp and $[1,15, 2,85]$ pp respectively) that comfortably exclude zero. The H3 sign-flip predicate is therefore *not satisfied*: the optimal does not change between regimes, and P1 is the better of the two compressed-corpus pipelines in both. P3 (joint relevance-conditional routing) leads both regimes by a substantial margin.

Table 8: H3 verdict table. Mean F1, per-query EUR cost, and *achieved* compression ratio ($\times$) by regime and pipeline from `results/h3_final/` and the cost-instrumented re-run `results/h3_eprice/` ($n = 150$ workloads per cell). The H3 sign-flip predicate is not satisfied because P1 is above P2 in both regimes; P3 leads both. The reframed finding is the robust compress-first preference ($P1 > P2$). The achieved-ratio column is the key to reading P3: only P1 actually compresses ($7.49\times / 2.11\times$), while P2 and P3 route fragments verbatim ($1,00\times$), so the cross-pipeline F1 gap is not a matched-compression comparison.

Pipeline	Storage-bounded			Accuracy-bounded		
	F1	EUR/q	$\times$	F1	EUR/q	$\times$
P1 (compress→retrieve)	0,396	$2,72e^{-5}$	7,49	0,648	$3,29e^{-5}$	2,11
P2 (retrieve→compress)	0,364	$2,82e^{-5}$	1,00	0,628	$3,00e^{-5}$	1,00
P3 (joint, conditional)	0,820	$3,07e^{-5}$	1,00	0,895	$3,10e^{-5}$	1,00
p_{Holm} (P1 vs P2, paired bootstrap)	2×10^{-4}			2×10^{-4}		

4.7.3. Interpretation: Why the Sign-Flip Did Not Appear

The original H3 predicate rested on the LongLLMLingua-style intuition that post-retrieval compression (P2) is more efficient than pre-retrieval compression (P1) because the retriever has already filtered the corpus down to the relevant chunks before the compressor runs. The interpretation depends on two assumptions that the empirical work invalidates: first, that retrieval over an uncompressed corpus is cheap (it is not — the FAISS index over the uncompressed corpus grows linearly in corpus size and dominates the storage budget); second, that compression after retrieval can be more aggressive than compression before (the empirical compression ratios at iso-F1 are comparable). What remains is the symmetric effect: pre-indexing compression amortises the compressor cost across queries while post-retrieval compression pays it per query.

The compress-first preference over retrieve-first is a substantive finding because it does not reproduce the assumed post-retrieval-compression default on the stack we tested, and the P3-higher-F1 result points to a promising routing variant. For practitioners, the recommendation is deliberately hedged on cost: **compress-first (P1) is the safe default over retrieve-first (P2) on a comparable stack; P3 (joint relevance-conditional routing) is worth evaluating where the F1 gain matters** — its F1 leads compress-first by 42 pp (storage) and 25 pp (accuracy) — but note that P3 attains this F1 by routing fragments *verbatim* (achieved ratio 1,00×, Table 8) rather than by better compression, and the cost model is corpus-dominated and does not meter compression compute (Section 4.7.4), so the F1 gain cannot yet be claimed as cost-free and must be re-measured with a faithful per-pipeline cost model before deployment; post-retrieval compression (P2) is only preferable when the corpus is small enough that storage cost is irrelevant and retrieval recall is the binding constraint.

4.7.4. Limitations and the Cost Model Caveat

The cost model uses the requested compression ratio rather than the achieved compression ratio because the achieved ratio depends on the compressor and the request flows through the pipeline as target. For Phi-3-Mini extractive the ceiling effect makes this an asymmetric error in favour of P2 (which pays the compressor per query and is therefore most sensitive to the ceiling), so the P1 result is the conservative one.

A second, stronger caveat is sharpened by a cost-instrumented re-run that records the *achieved* compression ratio per pipeline (`results/h3_eprice`). The `eur_per_query` column is *not* constant — it is computed per pipeline from a per-pipeline token count (P1: full corpus plus compressed corpus; P2: corpus plus retrieved tokens; P3: corpus plus 0,7× retrieved) — but the pipelines differ by only 4,1–4,9% in EUR within a regime, because a shared corpus-indexing term dominates `input_tokens` and is added identically to every pipeline. This small, corpus-dominated spread is easily mistaken for a single constant; the substantive defect, however, is not that the EUR values are near-identical but that the cost model is *corpus dominated* and, more importantly, **does not meter compression compute at all** — only the resulting token counts — so a pipeline that runs an expensive compressor pass is not charged for it. The achieved-ratio column makes the consequence concrete: only compress-first (P1) actually compresses (7,49× storage, 2,11× accuracy), while both retrieve-first (P2) and the joint router (P3) record an achieved ratio of 1,000 — they send the retrieved fragments to synthesis essentially *verbatim*. Together these explain the headline cleanly: **P3’s large F1 advantage comes substantially from not compressing**, at a compression and compute cost the current corpus-dominated model does

not faithfully charge for. As a consequence the thesis makes **no cost-parity and no Pareto claim** for the P3-vs-P1 comparison; the cross-pipeline F1 comparison is *not* a matched-compression comparison, and P3 should be read as “verbatim routing preserves F1 at a compression and cost the current model does not faithfully charge for” rather than as a free-lunch improvement. The remedy is a faithful per-pipeline cost model that meters the compressor pass and amortises a persistent index across queries, rather than re-charging the corpus on every query. Section 5.6 names this as the most natural follow-on study, alongside the single-retriever, single-embedder limitation (Section 5.3).

4.8. H4: Summary-Level Inference Disclosure

H4. The summary-level inference-disclosure metric (i) distinguishes a baseline planner from a priors-only planner (signal: positive paired effect, $p < 0,05$), and (ii) is reduced by compression at $4\times$ versus $1\times$ (reduction: positive paired effect, $p < 0,05$). **Verdict: SUPPORTED on the unbiased benchmark for three of four compressors (truncation, filter, LLMLingua-2), under two caveats below.** Phi-3-Mini extractive shows no significant reduction (0,4 pp, $p = 0,91$) and is the structural exception. **Caveat 1 (reader bias):** the Llama-3.1-8B reader exhibits a strong YES-bias (rate $\approx 0,03$ on priors-only with a YES ground truth, $\approx 1,00$ on a NO ground truth); the pooled- rate verdict is balanced by ground-truth-class distribution rather than by reader-side accuracy, and the operational reading is “*compression preserves enough information to flip a no-biased reader to confident YES on a protected-fact question*”. Per-class breakdown and the conservative-privacy framing are in Section 4.8.7. **Caveat 2 (construct validity):** the disclosure-reduction ranking is *monotonic in how much text-information the compressor destroys* — truncation (most destructive) tops the ranking at 24,6 pp, then filter (20,4 pp), then LLMLingua-2 (18,9 pp), then Phi-3-Mini extractive (most semantically preserving) at 0,4 pp not-significant. This pattern is consistent with the metric measuring general information destruction rather than privacy-specific filtering; an **oracle field-redaction control** (results/h4_oracle) that removes *only* the answer-bearing numeric values — destroying $\approx 0\%$ of the text — matches truncation’s reduction (23,4 pp vs 24,6 pp), confirming that the metric *does* detect targeted removal when it occurs and that blanket compression is the wrong privacy lever when a targeted one is available. The H4 *finding* is robust but the *privacy interpretation* is deflationary: aggressive compression and trivial field redaction reduce disclosure by indistinguishable amounts, and field redaction does it at near-zero coordination cost. Full discussion in Section 4.8.5.

4.8.1. Threat Model

The H4 measurement is built around a concrete adversary. Stating the adversary’s capabilities explicitly is what makes the metric’s verdict meaningful as a privacy claim rather than just an information-retention measurement.

Adversary’s goal. Recover a *protected fact* about a CONFIDENTIAL fragment held by the memory bus. A protected fact is defined operationally as a binary (yes/no) proposition over the fragment’s content whose answer the C1 generator records as ground truth (e.g., “did agent A approve more than EUR N ?”, “did the meeting attendee list include person P?”).

The canonical protected-fact list for each fragment is in `data/processed/c1-v0.1/`; the H4 evaluation iterates over the full list.

Adversary’s capabilities. The adversary (i) can query the downstream reader with arbitrary yes/no questions about the fragment; (ii) knows the compressor identity and ratio (is not blinded by the policy layer); (iii) does *not* have direct access to the uncompressed source — the only signal flowing to the adversary is what the compressed summary carries through the reader. In the H4 evaluation, this is operationalised by giving the reader a fixed public preamble plus the compressed fragment (in `BASELINE / COMPRESSED-4×`) or only the preamble (in `PRIORS`).

What the adversary does not have. Direct read access to the bus (the policy layer blocks that for `CONFIDENTIAL` slots), training-time access to the reader (the reader is a frozen Llama-3.1-8B), or signal-amplification through repeated queries with feedback (each yes/no question is asked once per fragment).

The metric as operational measurement. Given this adversary, the inference-disclosure metric (Section 4.8.3) measures the *upper bound* on what the adversary can recover at a chosen (compressor, ratio) operating point. The privacy claim the memory bus’s policy layer enforces is: if the operator selects an operating point whose disclosure is below the governance budget, the adversary above cannot recover protected facts at a higher rate than that budget. The budget is enforceable because the disclosure rate is empirically measured rather than asserted.

4.8.2. *Why This Measurement Matters*

The memory bus stores summarised representations of source fragments. When a fragment is classified `CONFIDENTIAL` or higher, the goal of the policy layer is to ensure that the summary does not leak the protected facts that motivated the classification. The H4 metric quantifies, per compressor and per ratio, the rate at which the Section 4.8.1 adversary can recover protected facts from the compressed summary. This converts the qualitative goal (“compression should not leak”) into a quantitative operating-point selection (“at this ratio, this compressor leaks $x\%$ of protected facts to the Section 4.8.1 adversary”). The metric is what makes the memory bus’s privacy claim auditable against a concrete threat.

4.8.3. *Metric Definition and the Unbiased Benchmark*

For each `CONFIDENTIAL` fragment, the C1 generator emits a list of (yes/no question, ground-truth answer) pairs whose ground-truth answer is a protected fact (as defined in Section 4.8.1) about the fragment. The canonical benchmark variant balances the YES/NO distribution at the question level and uses a single comparator phrasing for the yes/no question template, eliminating a surface-pattern bias in the original H4 benchmark generator (`fact_aggregation.py:119–156`, fixed 2026-05-29). Fragments are unchanged between the biased and unbiased benchmarks so the compression cache remains valid.

The reader is the local Llama-3.1-8B-Instruct model accessed via the Ollama backend. Three conditions are evaluated:

PRIORS. Reader sees only the workload’s public preamble. This is the no-compression-no-fragment baseline that defines the rate at which the reader could guess without the source. (This baseline is also what operationalises the *second* condition of the Section 4.4.3 calibrated-regime predicate — a planner is in-regime if its accuracy at $q \rightarrow 0$ is no greater than its priors-only rate plus a small slack.)

BASELINE. Reader sees the uncompressed source fragments. This is the disclosure rate without any privacy protection.

COMPRESSED-4×. Reader sees the source fragments compressed at 4× by one of the four compressors.

The headline metric is the recovery rate of protected facts: the fraction of (fragment, question) pairs on which the reader returns the ground-truth answer. The H4 *signal* is whether **BASELINE** sits above **PRIORS** (the metric must measure something real); the H4 *reduction* is whether **COMPRESSED-4×** sits below **BASELINE** (compression must reduce disclosure). Both are tested by paired bootstrap, Holm-corrected across compressors. The canonical results directory is `results/h4_unbiased_v2/` (the 4-compressor run including truncation).

4.8.4. Result

Table 9 reports the per-compressor signal and reduction. The signal is positive and p -significant for every compressor (the uncompressed baseline beats priors-only by 28,6 pp on average), and the reduction is positive and p -significant for the three token-destructive compressors (truncation −24,6 pp, instruction-aware filter −20,4 pp, LLMingua-2 −18,9 pp) and *not significant* for the extractive copier (Phi-3 −0,4 pp, $p_{\text{Holm}} = 0,91$). The ordering is monotonic in how aggressively each compressor destroys source tokens — the construct-validity reading is developed in Section 4.8.5.

Table 9: H4 verdict table from `results/h4_unbiased_v2/` (4 compressors including truncation; canonical from 2026-05-30 GPU rerun). The signal (priors $\rightarrow$ baseline) is positive and significant for every compressor; the reduction (baseline $\rightarrow$ compressed-4×) is positive and significant for three of four compressors and *not significant* for Phi-3-Mini extractive ($p = 0,91$). Compressors are listed in order of how aggressively they destroy source tokens: truncation drops the tail of the text entirely; the token-level filter and LLMingua-2 drop individual tokens chosen by ranking / classifier; Phi-3 extractive copies verbatim spans of the source. Reader: local Llama-3.1-8B-Instruct. See Section 4.8.7 for the reader-bias note and Section 4.8.5 for the construct-validity finding on this ordering.

Compressor	priors	baseline	compr-4×	signal Δ	reduction
Truncation (most destructive)	0.50	0.78	0.54	+0,286 ($p_{\text{Holm}} = 8 \times 10^{-4}$)	−0,246 (p_{Holm})
Instruction-aware filter	0.50	0.78	0.58	+0,286 ($p_{\text{Holm}} = 8 \times 10^{-4}$)	−0,204 (p_{Holm})
LLMingua-2	0.50	0.78	0.59	+0,286 ($p_{\text{Holm}} = 8 \times 10^{-4}$)	−0,189 (p_{Holm})
Phi-3-Mini extractive (most preserving)	0.50	0.78	0.78	+0,286 ($p_{\text{Holm}} = 8 \times 10^{-4}$)	−0,004 (p_{Holm})
Verdict: Signal significant for all 4; reduction significant for 3 of 4 compressors. SUPPORTED (with construct-validity)					

Figure 9 visualises the disclosure-versus-coordination trade-off: each point is a (compressor, ratio) cell, with disclosure on one axis and coordination success on the other. The instruction-aware filter and LLMingua-2 occupy a low-disclosure, moderate-coordination region at 4×

Phi-3 extractive sits in a high-disclosure, high-coordination region; truncation is the extreme of the same line — lowest disclosure, lowest coordination.

4.8.5. *Construct-Validity Finding: Disclosure-Reduction Tracks Information Destruction, Not Privacy-Specific Filtering*

The four compressors order monotonically in Table 9: **truncation** (most aggressive destruction, drops the tail of the text entirely) shows the *largest* disclosure reduction at 24,6 pp; the **instruction-aware filter** and **LLMLingua-2** (token-level filtering, intermediate destruction) follow at 20,4 and 18,9 pp respectively; **Phi-3-Mini extractive** (most semantically preserving — the compressed output is by construction a verbatim subset of the source tokens) shows *no significant reduction* (0,4 pp, $p = 0,91$).

This pattern is exactly what would be predicted if the H4 metric reflected general information destruction rather than a privacy-specific information-filtering mechanism. The thesis therefore interprets H4 as evidence that compression reduces inference disclosure roughly in proportion to how much text-information it destroys, and *not* as evidence that the smarter compressors (filter, LLMLingua-2) are doing privacy-preserving filtering in any sense beyond destroying tokens that happen to include protected facts.

A discriminating control was run to test this interpretation from the opposite end (results/h4_oracle, scripts/h4_oracle_redaction.py): an **oracle redaction** that removes *only* the answer-bearing numeric values (recorded hours and approved budget) the protected-fact questions probe, leaving every other token intact. If disclosure could only be reduced through bulk destruction, this near-zero edit would barely move the metric. Instead it reduces recovery from a balanced baseline of 0,79 to 0,56 — a 23,4 pp reduction, essentially at the priors floor 0,51 — while destroying $\approx 0\%$ of the text (the redacted fragments are in fact $\sim 7\%$ longer, since the [REDACTED] marker is longer than the digits it replaces). That 23,4 pp matches truncation’s 24,6 pp, which required destroying the entire tail of the fragment. Two things follow. First, the H4 metric *does* detect privacy-specific removal when it occurs — the construct is not blind to targeted redaction; the blanket compressors simply do not do it, reducing disclosure only as a side effect of destruction. Second, the memory bus gains a concrete, non-lossy privacy mechanism: field-redacting protected values achieves blanket-compression-level disclosure reduction at near-zero coordination cost. The empirical implication for the memory bus’s policy layer (Section 4.8.8) is therefore stronger than before: disclosure-budget enforcement works as designed, *and* targeted redaction is an available lever with a far better privacy-per-utility profile than aggressive compression. The remaining open control — a full privacy-gain-versus-coordination-loss frontier across compressors — is named in Section 5.6.

4.8.6. *Interpretation*

The reduction signal scales with how aggressively the compressor operates on tokens, not with the compression ratio per se. The filter and LLMLingua-2 both drop tokens at the lexical level; the fragments that survive are filtered for relevance, which disproportionately removes the protected-fact tokens. Phi-3 extractive preserves contiguous spans verbatim; if the protected fact is in a span the verifier selects, it is in the output. This makes the disclosure metric *useful* as a




Figure 9: The privacy-versus-coordination operating-point landscape. At 4× compression, the aggressive token-level compressors (filter, LLMingua-2) trade a few coordination-success points for substantial disclosure reduction; Phi-3 extractive preserves more coordination at the cost of more leakage. Each point is a (compressor, ratio) cell; the axes are pooled across families.

compressor-selection tool: it ranks compressors by privacy at iso-ratio, not just by compression efficiency.

The H4 result lets the memory bus’s policy layer convert a governance constraint (“CONFIDENTIAL fragments should leak protected facts at no more than $X\%$ to a downstream reader”) into a compressor-and-ratio selection. The integration is described in the next subsection.

4.8.7. The Reader’s YES-Bias Caveat

The Llama-3.1-8B reader exhibits a strong YES-bias: when the ground-truth answer is YES, the reader’s recovery rate is approximately 0,03 on the priors-only condition and 0,58 on the baseline; when the ground-truth answer is NO, both rates are approximately 1,00. The pooled rates of 0,50 (priors) and 0,78 (baseline) reflect a balanced ground-truth distribution, not a reader that is equally accurate in both directions. The correct framing of H4 is that the metric measures “*compression preserves enough information to flip a no-biased reader to confident YES on a protected-fact question*”. This is an asymmetric leakage signal but it is a valid one: a reader that defaults to NO is the conservative case for privacy auditing because it is the case in which the absence of the protected fact most cleanly translates to a privacy preservation.

The caveat is documented here rather than in the verdict block because the pooled rates are what the verdict pipeline tests and the verdict is stable under the caveat. A future-work direction

The policy check is effectively free (sub-microsecond, $\sim 5\text{--}7$ million checks per second across the two hosts), so policy enforcement is not a throughput bottleneck. End-to-end write and read run at $\sim 18,200$ and $\sim 19,900$ operations per second on Host A ($\sim 5,750$ and $\sim 6,400$ on Host B) with sub-millisecond tail latency, and the audit hash-chain verifies at $1,68\ \mu\text{s}$ per row (Host A) and $2,19\ \mu\text{s}$ per row (Host B) with the chain intact over the 5,500-row log. These figures characterise the reference implementation’s single-threaded core; a concurrent-HTTP-path throughput sweep and an audit-verify scaling curve remain future work (Section 5.6 of Chapter 5).

4.9. Corollary 2: Information-Density Scaling on Real Benchmarks

Corollary 2 (information-density scaling) — empirical observation, $n = 3$. Across the three benchmarks tested (C1 family-a, MultiHopRAG, HotpotQA), the *shape* of the cliff — sharpness and pre-cliff plateau — varies systematically with the AUC-based information-density estimate θ_{info} (defined formally in Section 4.9.1, distinct from θ_q). Dense tasks ($\theta_{\text{info}} \rightarrow 1$) exhibit a sharp cliff with a high coordination-success plateau; distributed tasks ($\theta_{\text{info}} \rightarrow 0$) degrade gradually from a lower baseline. **Important scope:** the corollary is reported as an empirical observation on three benchmarks, not as a derived consequence of the compounding-error model. (i) $n = 3$ is a small sample — the trend is consistent but not a scaling law. (ii) θ_{info} is computed from the same coordination curve whose shape it describes, so “ θ_{info} predicts shape” is a circular description: an external probe of task density (e.g., conditional entropy of the answer given the input) is required to break the circularity and is named as future work in Section 5.6. (iii) The corollary makes a claim about cliff *shape*, not about the absolute cliff position τ^* — HotpotQA cliffs at $\tau^* \approx 2,7$ (similar to dense C1 family-a) yet has $\theta_{\text{info}} \approx 0,37$ (very distributed); the baseline coordination on HotpotQA is only 0,59 and the degradation shape is gradual rather than precipitous. **Verdict: SUPPORTED as an empirical observation on $n = 3$ benchmarks** by the gap in θ_{info} between C1 family-a ($\approx 0,97$, dense), MultiHopRAG ($\approx 0,48$, distributed), and HotpotQA ($\approx 0,37$, highly distributed), and by the qualitatively distinct degradation shapes visible in Figure 10 and Figure 3. The original H6 predicate (synthetic-vs-real τ^* within $\pm 15\%$) is honestly NOT SUPPORTED — MultiHopRAG cliffs at 11,3, far from the C1 reference — and that is what motivated the reframing from *position transfers* to *shape transfers*.

4.9.1. θ_{info} versus θ_q

A subtle point that the corollary reframing makes explicit: the information-density θ_{info} that Corollary 2 talks about is *not* the cliff-recall threshold θ_q that Corollary 1 and the compounding-error model talk about. The two quantities are derived from different measurements and answer different questions. θ_q is the recall threshold per family at which coordination success drops to 0,5 of p_0 , estimated by `derive_theta()`; θ_{info} is the per-task information-density estimate

$$\theta_{\text{info}} = 1 - \text{AUC} \left[\frac{\text{coord-success}(r)}{p_0} \right]_{r=1}^{r_{\text{max}}},$$

estimated by `estimate_task_theta()`, which captures how *concentrated* the task’s information is in compression-sensitive tokens. The two are easily conflated; the project glossary distinguishes them explicitly.

4.9.2. Result on Three Benchmarks

Table 11 reports θ_{info} and the empirical τ^* for the three benchmarks. The information density spans from 0,97 on the dense C1 family-a to 0,37 on HotpotQA, and the cliff position correspondingly spans from $\tau^* \approx 2,7$ on family-a to $\tau^* \approx 11,3$ on MultiHopRAG — a $4\times$ shift in the cliff position that the single- θ_q compounding-error model cannot account for and that the information-density scaling claim does.

Table 11: Corollary 2 verdict table. θ_{info} and empirical τ^* across the three benchmarks. The cliff shifts late as the information density drops, with the gap between dense and distributed tasks far exceeding the 0,1 threshold of the corollary predicate. The C1 family-a row uses the canonical h1_h2_v2 fit $\tau^* = 2,5$ (Section 4.6); the earlier register $\tau_{\text{aux}}^* = 2,6997$ on h1_h2_final is superseded and is retained only as a parenthetical in the footnote. The $\sim 8\%$ shift between the two registers does not change the qualitative ordering or the corollary verdict.

Task	θ_{info}	τ^*	θ_{info} source CSV
C1 family-a (dense, numeric)	0,97	2,5	h1_h2_v2/ (a, lingua2)
MultiHopRAG (distributed QA)	0,48	11,3	h6_final/results.csv
HotpotQA (very distributed)	0,37	2,7	hotpotqa_sweep/

Gap (dense vs MultiHopRAG) = 0,48; (dense vs HotpotQA) = 0,59.

All three θ_{info} values are computed by `estimate_task_theta()` (AUC-based) from the per-task CSVs above and consolidated in `results/corollary2_theta_info.json`. The τ^* for C1 family-a is the canonical h1_h2_v2 fit 2,5 (the earlier 2,7 was the superseded h1_h2_final auxiliary fit; the fine-grid deterministic-solver cliff is lower still, $\approx 1,1$, Section 4.3).

The HotpotQA result requires brief comment: τ^* on HotpotQA is approximately 2,7, similar to C1 family-a’s, but the baseline coordination success on HotpotQA is only 0,59, so the absolute coordination success at the cliff is much lower ($\approx 0,30$). The corollary claim is about the *shape* of the degradation (gradual versus sharp), captured by θ_{info} via the AUC computation, not about the absolute position alone. Figure 10 shows the HotpotQA cliff with the gradual-degradation pattern visible.

4.9.3. Interpretation

The corollary makes the compounding-error model’s per-family θ_q structure plausible as the first-order term of a finer description. θ_{info} tells you how much of the task’s success curve is concentrated in compression-sensitive tokens; θ_q tells you the recall threshold at which success collapses given that concentration. The two together describe the cliff position and shape; either alone is a first-order summary. For the thesis, θ_{info} is what makes the cross-benchmark transfer of the cliff structure auditable: a practitioner who needs to deploy compression on a new task




Figure 10: HotpotQA cliff curve under LLMLingua-2 compression. The baseline coordination success is 0,59, lower than family-a’s 1,00; the cliff position by the $0,5 \cdot p_0$ rule is $\tau^* \approx 2,7$, similar to family-a’s; but the *shape* of the degradation is gradual rather than sharp, as captured by the lower θ_{info} of 0,37.

can estimate θ_{info} on a small calibration set and predict whether the task is dense (cliff early, sharp) or distributed (cliff late, gradual) before sweeping the full curve.

The dropped H6 original predicate (synthetic-versus-real τ^* within $\pm 15\%$) is honestly NOT SUPPORTED at the table-level numbers; the reframed Corollary 2 predicate (gap in $\theta_{\text{info}} \geq 0,1$ between dense and distributed benchmarks) is comfortably SUPPORTED. The synthetic benchmark is still the right place to characterise the cliff mechanism cleanly, and the real benchmarks are the right place to validate the structure transfers.

4.10. Reproducibility

The results of this chapter span three regimes of reproducibility, and conflating them would over- or under-state what a reader can re-derive. This section states the guarantee for each regime once; the one-command recipe, the per-arm wallclock budget, and the release manifest live in Appendix D. The structure follows the NeurIPS reproducibility-checklist convention [65] of separating *code-and-data availability* from *computational reproducibility*, but reports the latter as a narrative keyed to Table 12 rather than as a filled-in checklist.

Tier 1 — byte-deterministic local. The cliff sweep (H1, H2), the RAG-pipeline evaluation (H3), the bootstrap confidence intervals, and the figure regeneration are pure Python over a fixed compression cache: a regex information-extraction solver, numpy resampling, and deterministic compressors (truncation, the LLMLingua-2 token classifier at `argmax`, the TF-IDF-plus-cross-encoder filter). All randomness flows from a single master seed through `Seeds.derive`

Table 12: The three reproducibility tiers of this thesis. “Byte-deterministic” means a rerun produces a CSV that differs zero against the published one given the same seed; “in distribution” means verdict-level aggregates (per-cell means, τ^* fits) reproduce within the published bootstrap CI but individual rows may differ; “not reproducible” means a rerun produces closely-similar but not identical numbers, and the published CSV is itself the canonical record.

Arm	Byte-deterministic?	Canonical artefact
H1/H2/H3 cliff sweep, scoring, bootstrap CIs, figure regeneration	Yes , given seed + cache	results/h1_h2_v2/, results/h3_final/
H4 reader, H5 planner, Phi-3 extractive compressor	In distribution only	results/h4_unbiased_v2/, results/h5_final/
Frontier API arm (Qwen-2.5-72B, DeepSeek V4 Pro, GPT-oss 120B)	No	results/frontier_{qwen72b,deepseekv4}/

in `src/m6/utils/seed.py`, which seeds `numpy`, Python’s `random`, `PYTHONHASHSEED`, and `torch`. Rerunning `make repro-cliff` (or `repro-rag`, `repro-figs`) from the cache produces CSVs and PDFs that differ zero against the published artefacts; the bootstrap CIs are exact to the last decimal because the resampler is seeded and n_{boot} is fixed.

Tier 2 — distributionally deterministic local LLM. The H4 protected-fact reader, the H5 model-scaling planner, and the Phi-3-Mini extractive compressor invoke local language models through Ollama. Greedy decoding (temperature = 0 for the H4 reader, 0,1 for the H5 planner) with a fixed per-call seed is deterministic per call up to the reduction order of the CUDA/Metal kernels, which is not guaranteed identical across hardware or backend builds. A rerun on the same host therefore reproduces individual generations to within a few tokens, and the verdict-level aggregates — per-cell coordination success means and the τ^* fits — reproduce to the published values within the bootstrap CI. The published CSVs were produced on the RTX 5090 host with the model digests recorded in `docs/MODEL_CARDS.md` (verifiable with the ID column of `ollama list`); pinning those digests is what makes the tier reproducible *in distribution* rather than byte-for-byte.

Tier 3 — frontier API, not reproducible. The frontier arm (Section 4.6) calls vendor-hosted models over an API. Three mechanisms break byte-reproducibility: server-side request batching changes per-token logits, the seed parameter is best-effort rather than contractual, and the vendor may silently update the served weights. Rerunning `make repro-frontier` therefore produces closely-similar but not identical numbers; the published `results/frontier_*/results.csv` files *are* the canonical record, and every analysis downstream of them (τ^* fit, bootstrap CI, the TOST of Section 4.6.3) is itself Tier 1 — deterministic given those CSVs. A rerun whose τ^* CI no longer overlaps the published CI is the documented trigger for suspecting a vendor-side model update; the GPT-oss 120B diagnostic additionally carries a `STATUS_NONCANONICAL.txt` marker because it is scoped out of the calibrated regime (Section 4.6.4).

Traceability. Every quantitative claim in this thesis is traceable to one named directory under `results/` through the `CANONICAL_NUMBERS.md` registry in the repository root, which maps each reported figure to its source artefact and the key or column that yields it. The open-source release bundles the manuscript source, the code and configuration, the canonical CSVs and JSONs, the model and data cards (`docs/MODEL_CARDS.md`, `docs/DATA_CARDS.md`), the pinned `requirements.lock.txt`, and the `docker-compose.yml` that brings the memory-bus reference service up locally.

4.11. Summary of Verdicts

Table 13 consolidates the verdict block of this chapter for easy reference. Each row links to the canonical results directory under `results/` that produced the numbers.

Table 13: Consolidated verdict table. Original predicate verdicts are reported alongside the reframed corollary verdicts where the reframing was made before the chapter was written.

ID	Claim	Verdict	Canonical dir
H1	QA-F1 does not transferably predict coordination success	SUPPORTED	<code>h1_h2_v2/</code>
H2	Sharp coordination cliff exists	SUPPORTED (11/12)	<code>h1_h2_v2/</code>
H3	RAG sign-flip across regimes	NOT SUPPORTED (reframed as compress-first dominance)	<code>h3_final/</code>
H4	Compression reduces inference disclosure	SUPPORTED on unbiased	<code>h4_unbiased_v2/</code>
H5 → Cor. 1	Cliff-position invariance within calibrated regime	SUPPORTED as <i>no shift detected</i> (frontier point estimates match; $\pm 20\%$ TOST not passed)	<code>h5_final/</code> , <code>frontier_qwen72b_e2/</code> , <code>frontier_deepseekv4/</code>
H6 → Cor. 2	θ_{info} scales the cliff	SUPPORTED (empirical observation, $n = 3$ benchmarks)	<code>h6_final/</code> , <code>hotpotqa_sweep/</code>

Chapter 5 draws the threads together: how the findings interact, what CAAC contributes as a constructive realisation of the compounding-error bound, what the limitations of the work are, what its significance to practice and to the field is, and where the largest remaining unknowns sit.

5. DISCUSSION AND SUMMARY

This chapter draws together the empirical work of Chapter 4 and the system design of Chapter 3: what the findings say jointly, what CAAC adds as a constructive realisation of the compounding-error bound, where the work is limited, what its significance is to practitioners and to the field, how it relates to industry observations on multi-agent token economics, and where the most natural follow-on studies sit.

5.1. Synthesis

What stands and what doesn't. Of the six pre-registered hypotheses, **H1 (QA-F1 mis-ranks compressors for multi-fragment use)** holds as stated and is the manuscript's most-citable result; **H2 (a sharp cliff exists)** holds on 11 of 12 (compressor, family) cells, with the remaining cell (filter $\times$ family-c) accounted for mechanically in Section 4.3. The remaining four hypotheses were reframed: **H3** (RAG sign-flip between regimes) is *not supported* on the predicate as written — a narrower single-stack compress-first-over-retrieve-first ordering survives; **H5** (planner-scale monotonicity) is reframed as *Corollary 1 (ceiling-cliff separation)* and is carried by one well-identified frontier cell (Llama-3.1-8B vs Qwen-2.5-72B, family-a $\times$ LLMLingua-2) rather than by the original parameter-scale sweep; **H6** (cross-benchmark τ^* transfer) is reframed as *Corollary 2 (information-density scaling)* on three benchmarks; **H4** (compression reduces inference disclosure) is supported on three of four compressors but the privacy interpretation is deflationary — the reduction is monotonic in information destruction and matched at ~ 23 pp by trivial oracle field-redaction at near-zero destruction. The compounding-error model is a *first-order* position estimate (median $\sim 35\%$ relative error) and an *explanation* of why the cliff is sharp; the threshold mechanism it posits is independently confirmed on the dense family by the direct A3 probe (Section 4.4.5). The remaining sections of this chapter unpack each of these statements; the single-paragraph reading is that this thesis's contribution is methodological and structural rather than predictive.

The thesis answers the question posed in Chapter 1 with five quantitative results that fit together as a single story about context compression in multi-fragment LLM workflows.

Single-agent question-answering accuracy does not transferably predict multi-fragment coordination success — across the four compressors the workload-level Spearman correlation between the two ranges from $-0,82$ (filter) to $+0,32$ (Phi-3-Mini extractive), with LLMLingua-2 and truncation at essentially zero ($+0,03$, $+0,05$); all four CIs exclude the moderate-correlation threshold of $0,6$ (Section 4.2). The pooled-over-ratios correlations look larger (up to $+0,38$) but are pseudo-replicated, and the per-family decomposition shows no positive within-family relationship for any compressor. A token-overlap metric does not predict the harder structural property a planner needs, and the single-agent benchmarks that dominate the compression literature are insufficient for the multi-fragment deployment decisions FCG and TalentAdore-style workflows depend on.

A sharp coordination cliff τ^* exists on the C1 benchmark for every (compressor, family) cell at which the compressor's mechanism drops the task-critical tokens (Section 4.3); the only cell that does not exhibit a detectable cliff is the instruction-aware filter on the multi-step-retrieval family, where the cross-encoder re-ranker preserves the chain-reference tokens that the task scorer turns

on. The cliff is sharp enough that an empirical-bound model built around a threshold-success assumption (Section 4.4) recovers the position to within a useful first-order approximation.

The cliff position shows **no detected shift** across the three local architecture families tested (Qwen-2.5-1.5B, Phi-3-Mini-3.8B, Llama-3.1-8B) within a precisely-defined calibrated regime (Section 4.5) — though the local arm’s 24% spread fails the strict $\pm 20\%$ tolerance, so this arm is consistent-with but not sufficient-for the claim. The load-bearing evidence is the invariance of the LLM-planner cliff itself: the fine-grid re-resolution (Section 4.3) shows the synthetic reference $\tau^* = 2,5$ is a coarse-grid artefact of the deterministic solver, so the honest comparison is LLM-vs-LLM. On family-a $\times$ LLMingua-2, Llama-3.1-8B cliffs at $\tau^* \approx 2,48$ and Qwen-2.5-72B at $\tau^* \approx 2,79$ ($n = 30$) — a $\sim 12\%$ spread across a $9\times$ parameter scale-up and a change of architecture family, both at $p_0 = 1,0$, while the deterministic solver on the same cell cliffs at $\approx 1,1$. (The original $n = 10$ Qwen run, $\tau^* = 2,68$, is 7,3% from the coarse-grid 2,5 — now read as an approximate coincidence rather than a validation; DeepSeek V4 Pro’s point estimate $\tau^* = 2,15$ sits inside a bootstrap CI so wide — $[1,76, 7,14]$, median 5,7 — that its position is only weakly identified, so it corroborates without robustly confirming, Section 4.6.) The defensible claim is therefore that planner *type* moves the cliff but scale and architecture *within the LLM class* do not. The cliff is *not* invariant to reasoning regime: extended-reasoning planners such as GPT-oss 120B cliff at substantially higher ratios because they can recover from sub-threshold information by chain-of-thought, violating the threshold-success assumption A3. This is the empirical evidence that the cliff is a property of the (compressor, task) pair within a defined regime — the structural finding that the thesis contributes to the compression literature.

Compress-first is preferred over post-retrieval compression in both storage- and accuracy-bounded RAG regimes under the single retriever/embedder stack we evaluate (BAAI/bge-large-en-v1.5 + FAISS-CPU; Section 4.7). The H3 hypothesis as written (sign-flip between regimes with ≥ 5 pp effect) is not supported, but the reframed finding — that the simpler, more amortisable architecture is also the better one under the stack we tested — does not reproduce the post-retrieval-compression preference of the LongLLMLingua line of work on this stack and is the substantive C3 contribution. (Generalisation to other retrievers/embedders is an explicit limitation, Section 5.3.)

Compression measurably reduces inference disclosure of protected facts to a downstream reader (Section 4.8), and the reduction is *monotonically proportional to how much text-information the compressor destroys*. The truncation baseline (most destructive — drops the tail of the text entirely) reduces disclosure by 24,6 pp; the token-level compressors filter and LLMingua-2 reduce it by 20,4 pp and 18,9 pp respectively; the extractive copier (Phi-3-Mini extractive — which preserves verbatim spans of the source) shows no significant reduction (0,4 pp, $p = 0,91$). This construct-validity ordering (Section 4.8.5) means the H4 metric tracks information destruction rather than privacy-specific filtering; the metric is still useful for the memory bus’s policy layer (Section 3.1 of Chapter 3) — more aggressive compression yields lower disclosure rates as required — but the *privacy* interpretation requires a target threat model that values destruction-induced reduction, not just compressor-specific reduction.

The cliff position varies with task information density across benchmarks (Section 4.9): C1 family-a ($\theta_{\text{info}} \approx 0,97$) cliffs early and sharply, MultiHopRAG ($\theta_{\text{info}} \approx 0,48$) cliffs late and gradually, HotpotQA ($\theta_{\text{info}} \approx 0,37$) degrades gradually from a lower baseline. This is the empirical evidence that θ_{info} is a useful second-order term beyond θ_q for predicting cliff structure on a new task; it is also the evidence that the synthetic-benchmark methodology of this thesis transfers to real benchmarks at the structural level even when the absolute cliff positions differ.

Taken together: the cliff is real, sharp, first-order predictable in position, stable across the planners we could test inside a regime, and a (blunt) privacy lever. A practitioner who needs to deploy compression on a multi-fragment workflow can measure θ_q and θ_{info} once on a small calibration set, predict the cliff position from the compounding-error model within a known tolerance, select an operating point inside the cliff’s safe zone, and audit the deployment’s privacy budget against the disclosure-rate table. The next section describes the algorithmic mechanism that turns this measurement loop into a runtime back-off.

5.2. Cliff-Aware Adaptive Compression: a Constructive Realisation

CAAC — Cliff-Aware Adaptive Compression (`src/m6/compressors/caac.py`) — is the algorithmic counterpart of the compounding-error model of Section 4.4. **CAAC’s contribution is a measurement-bounded safety floor:** given a per-family θ_q from one calibration pass, CAAC produces a per-fragment operating point whose critical-token-recall is provably above the cliff threshold, at the cost of accepting a back-off in achieved compression ratio on fragments where the inner compressor would have crossed the cliff. The wrapper is therefore an *operating-point selector*, not a *Pareto-dominating wrapper*: by construction it trades compression for a recall guarantee, so it cannot strict-dominate the fixed-ratio compressor that makes the opposite trade. The empirical strict-Pareto rate against the fixed-ratio baseline is 0/7 at every cliff ratio of the canonical sweep — exactly the structural property the model predicts. The two sections below unpack the algorithm and the operating-point framing in turn.

5.2.1. Algorithm

CAAC wraps any inner compressor that satisfies the Compressor protocol. Given a target ratio r , a per-family recall threshold θ_q from `derive_theta()`, and the number of compression passes N (in every CAAC experiment in this thesis $N = 1$ so $q_{\min} = \theta_q$), CAAC executes the following per fragment:

1. Run the inner compressor at the requested target ratio r , producing a compressed fragment and measuring its critical-token- recall q .
2. If $q \geq q_{\min}$, return the compressed fragment as is.
3. Otherwise, binary-search downward over ratios in $[r_{\min}, r]$ until the largest ratio $r' \leq r$ is found that satisfies $q(r') \geq q_{\min}$. Return the result of the inner compressor at r' .
4. Floor: $r_{\min}$ is fixed at 1.5× to prevent abdication to no-compression on hard fragments.

The binary search converges in at most five inner-compress calls per backed-off fragment. The CTR measurement is a set-intersection calculation in microseconds and does not measurably contribute to the wallclock. The wrapper is therefore drop-in: any existing deployment using LLMingua-2 or the instruction-aware filter or Phi-3-Mini extractive can swap to its CAAC-wrapped version without changing the memory-bus access layer or the storage layer.

5.2.2. Operating-Point Framing

The strict-Pareto criterion against the fixed-ratio compressor requires CAAC to match or exceed both the coordination success *and* the achieved compression ratio of the fixed-ratio baseline at every target ratio. CAAC by construction trades compression for coordination: it backs off the achieved ratio *below* the requested target whenever the recall constraint $q \geq q_{\min}$ would be violated. So at high target ratios where the fixed compressor cliffs to zero coordination, CAAC is backing off to $\sim 3\times$ achieved ratio while the fixed compressor delivers $\sim 12\times$ at zero coordination. Strict-Pareto says CAAC does not dominate the fixed compressor; the empirical 0/7 result reflects this faithfully.

The reframing we adopt as the substantive contribution is the operating-point one: CAAC’s selected ratio is *determined by the compounding-error bound*, not by tuning. For each family, θ_q is a measurement; CAAC’s binary-search procedure makes the ratio it operates at a predictable function of the inner compressor’s $q(r)$ curve and that measurement. With three families, CAAC produces three operating points, one per family, all on the safe side of the cliff. The fixed-ratio compressor produces one operating point per target ratio, with no guarantee about which side of the cliff it sits on. The two compressors populate complementary regions of the (coord-success, achieved-ratio) frontier: Figure 11 shows the two regions on a single plot, with the $q(r) = \theta_q$ curve drawn explicitly as the boundary that CAAC is constructed to respect.




Figure 11: CAAC and the fixed-ratio compressor on the (coord-success, achieved-ratio) plane. The shaded curve is the $q(r) = \theta_q$ contour from the compounding-error model. CAAC’s operating points sit on the high-coord side of the contour by construction; the fixed compressor’s points sit at the target ratio regardless. The two compressors populate complementary regions of the frontier; there is no strict-Pareto-dominance relation between them.

5.2.3. Weak-Dominance Result and the θ_q/N_q Ablation

Although strict-Pareto does not hold, *weak* dominance — the coordination-only comparison — does. Across all seven target ratios of the canonical sweep and across the two inner compressors evaluated, CAAC’s coordination success is greater than or equal to the fixed compressor’s 100% of the time. At the high target ratios where the fixed compressor cliffs to zero coordination, CAAC retains coordination at the price of compression. This is what the model predicts, and it is the correct framing of CAAC’s contribution: a guaranteed safety floor, conditional on the per-family recall threshold being correctly measured.

The θ_q/N_q ablation (Sections 5.2 `caac_theta` and `results/caac_N_*`) returns an informative null. Here N_q is CAAC’s internal *recall-exponent* knob — the exponent in the back-off threshold

$q_{\min} = \theta_q^{1/N_q}$ that controls how aggressively CAAC tightens its target — and is *not* the number of compression passes, which is fixed at one throughout this thesis. The coordination-success plateau of CAAC is invariant to the choice of θ_q in $\{0,6,0,7,0,8\}$ and the choice of N_q in $\{2,3,4,5\}$: the LLMLingua-2 plateau is 33% across all seven configs, the filter plateau is 50% across all seven configs, and only the achieved compression ratio responds, by backing off more aggressively as either parameter increases. The primary knob is the $r_{\min}$ floor, not θ_q or N_q . A sweep over $r_{\min} \in \{1,2,1,5,2,0,2,5\}$ is the natural follow-on that this thesis defers as future work.

5.2.4. What CAAC Contributes to the Thesis

CAAC is the demonstration that the compounding-error model is operationally usable, not just descriptively predictive. The algorithm is short, training-free, drop-in, and bounded in behaviour by a measurement. The contribution sits in Chapter 5 rather than as a headline contribution in Chapter 1 because the model is the substantive contribution and CAAC is its constructive realisation — but the design choice that CAAC enables is real: a practitioner who has measured θ_q on one family can deploy CAAC at any target ratio with the guarantee that the achieved operating point sits on the safe side of the cliff. This deployment-mode property is what makes the cliff measurement useful as something other than a publication finding.

5.3. Limitations

The limitations are catalogued explicitly so that a reader can calibrate the scope of the findings.

No multi-round agent coordination. The empirical evaluation uses a deterministic regex solver (H1, H2) or a single LLM call with all (compressed) fragments visible (Corollary 1, frontier, real-data transfer). The AutoGen v0.4 multi-round agent backend in `src/m6/orchestrator/` integrates with the memory bus but is not used in any reported result. Multi-round coordination introduces a per-round variance that the deterministic solver deliberately excludes; the cliff measured here is the solvability cliff under compression, not the multi-round communication-quality cliff. This scoping decision is recorded in the project repository.

Single compression pass. Every experiment uses $N = 1$. The compounding-error model admits arbitrary N (q compounds as q^N), but the $N > 1$ regime is not empirically validated. The θ_q/N_q ablation (Section 5.2) is an algorithmic sweep over CAAC’s internal recall-exponent knob N_q , not a sweep of the number of compression passes (which stays fixed at one). Multi-pass coordination is a natural follow-on study.

Family-a homogeneity. The 50 instances of family-a are all sum-of-eight-numbers variants. The cliff structure observed on family-a may not generalise to heterogeneous aggregation tasks (max, min, weighted average, conditional aggregation). The transfer to real benchmarks (Corollary 2) addresses the broader concern that synthetic tasks might not transfer to real ones, but the within-family homogeneity is a separate concern that future work could resolve by sweeping aggregation types within family-a.

Family-b LLM floor. The constraint-satisfaction family-b is hard for $\leq 8B$ planners independent of compression: baseline coordination success is near zero for the 1,5B (Qwen-2.5) and 3,8B (Phi-3-Mini) planners. The family-b cliff is not detectable in this regime, and Corollary 1’s invariance result is therefore tested on family-c only for the local sweep. A planner capable of solving family-b at the no-compression baseline (e.g., a frontier reasoning model) would enable a family-b validation of Corollary 1 that this thesis cannot produce.

Phi-3-Mini extractive ceiling. The Phi-3-Mini extractive compressor saturates at approximately $2,5\times$ achievable compression regardless of the requested target, because the verifier and the LLMingua-2 fallback combine to bound the achievable ratio. All evaluation reports the achieved ratio alongside the requested target so this saturation is visible, but it does mean that Phi-3-Mini extractive cannot participate in the high-compression-ratio comparisons; its curves flatten above $\sim 2,5\times$ rather than continuing to follow the family’s degradation pattern.

Extended-reasoning regime is out of scope. The calibrated regime predicate (Section 4.4.3) excludes planners that recover sub-threshold information by chain-of-thought reasoning. GPT-oss 120B and likely the GPT-4-class “thinking” models violate the predicate. The compounding-error model’s quantitative predictions do not apply to these planners; the GPT-oss 120B result is reported as an out-of-regime diagnostic in Section 4.6 rather than as a counterexample. Characterising the cliff in the extended-reasoning regime is the largest follow-on direction this thesis opens.

Compounding-error model is a first-order bound. The strict empirical match rate of the predicted-versus-empirical τ^* comparison is 33% within $\pm 25\%$ relative error and 67% within $\pm 50\%$ (in-sample; Wilson 95% CI on 4/12 is approximately [13%, 61%], reflecting the small denominator), or 25% within $\pm 25\%$ and 75% within $\pm 50\%$ under leave-one-family-out cross-validation; the relative-error distribution has median 35,3% and IQR [19,2%, 45,0%]. The model is useful as a first-order predictor and as a derivation that explains the cliff structure (specifically the threshold-success form $E[X/M] = q(r)$ with $P(\text{success} \mid r) \rightarrow \mathbf{1}[q(r) \geq \theta_q]$ in expectation; see Section 4.4); it is not a tight prediction. The graded-importance assumption A2 is the strongest of the four (Section 4.4.2); H4’s graded-disclosure measurements provide evidence that the binary-importance simplification is what the model gives up.

Predicted- τ^* bootstrap band is too tight. The bootstrap CI on θ_q that the Section 4.4.4 predicted- τ^* band is built from quantifies *sampling* uncertainty only — how much the prediction would wobble if we resampled workloads. It does not quantify the model’s *specification* error. The empirical τ^* falls *outside* this bootstrap band on 11/11 testable cells. The right way to read Figure 4 is “the band is tight because θ_q is well-estimated; the model is biased because A1–A4 are first-order approximations”; the actual accuracy of the model is what the in-sample/LOO match rates above measure.

Corollary 1’s local arm misses the strict tolerance. The local cross-architecture sweep at {Qwen-2.5-1.5B, Phi-3-Mini-3.8B, Llama-3.1-8B} on family-c shows a τ^* spread of approximately 24%, which is well below the `validate_model_independence` default 50% tolerance but *above* the strict $\pm 20\%$ tolerance matching the H2 falsifiable-claim spec. The local arm is therefore necessary but not sufficient for Corollary 1; the cross-architecture frontier evidence in Section 4.6 (Qwen-2.5-72B point estimate $\sim 7\%$ above the canonical synthetic

reference $\tau^* = 2,5$ and inside its bootstrap CI; DeepSeek V4 Pro bootstrap CI contains the reference) is what carries the corollary’s binary verdict across the strict bar. The on-disk record is `results/h5_final/model_independence_20pct.json`.

A3 (threshold success): circularity now broken by a direct probe. A3 was originally licensed only by the sharp empirical cliff of Section 4.3, which A3 in turn explains — a circular justification. That circularity is now broken by the independent A3 probe reported in Section 5.6 (`results/a3_probe`): hand-curated deletion of task-critical tokens *without* a compressor, scored by an LLM planner (Llama-3.1-8B) rather than the brittle regex solver, so a planner that *could* reconstruct partial information is what is tested. The probe supports A3 *directly* on the dense fact-aggregation family (threshold-like, logistic steepness $k \approx 15$, transition recall $r_0 \approx 0,84$) and shows it is only first-order on the distributed retrieval family (graded, $k \approx 5,9$). The residual limitation is therefore the narrower and honest statement that A3 is a threshold model for dense tasks and a first-order approximation for distributed ones, with a graded-success refinement as the open direction — the same information-density axis Corollary 2 measures. The cliff *exists* regardless; A3 is the tractable description of its shape, now validated on the axis where it is expected to vary.

A1 within-pass independence is approximate; Phi-3 extractive violates it. A1 (per-token retention independent across positions) holds approximately for token-level compressors (LLMLingua-2, instruction-aware filter, truncation baseline). It is violated by Phi-3-Mini extractive, which selects whole spans, inducing correlated token-survival events within a span. The phi3-extractive validation cells in Section 4.4.4 should be read as a robustness check of the bound *outside* its formal A1 regime rather than as a within-regime prediction.

Synthetic-task focus, real-task transfer at the structural level. The C1 benchmark is synthetic, by design: cliff characterisation requires controlled information density. The transfer validation on HotpotQA and MultiHopRAG shows that the cliff structure (existence, sharpness gradient with θ_{info}) transfers to real benchmarks, but the absolute cliff positions differ substantially. Generalisation to arbitrary task structures beyond the three task families this thesis examines is unverified and would require either a broader benchmark suite or a `theta_info` sweep over a corpus of real benchmarks.

Single-retriever, single-embedder RAG comparison. The H3 RAG pipeline catalogue uses FAISS-CPU with HNSW and BAAI/bge-large-en-v1.5 embeddings. Different retrievers or embedders may shift the cost balance of the storage- and accuracy-bounded regimes. The compress-first-over-retrieve-first finding is qualitatively robust, since it depends on the storage-cost asymmetry, which holds for any retriever. It is nonetheless demonstrated on only this one stack, so we report it as a single-stack non-reproduction of the LongLLMLingua preference rather than a general falsification.

H3 cost model is corpus-dominated and does not meter compression. The `eur_per_query` column in `results/h3_final/results.csv` is computed per pipeline (from per-pipeline `input_tokens`) and is *not* constant, but the three pipelines differ by only 4,1–4,9% within a regime because a shared corpus-indexing term dominates the token count, and compression *compute* is not metered at all (Section 4.7.4). The cost-instrumented re-run (`results/h3_eprice`) shows P2 and P3 route fragments verbatim (achieved ratio 1,00×)

while only P1 compresses, so the cross-pipeline F1 comparison is not matched on compression. As a result the thesis makes no cost-parity or Pareto claim for the P3-vs-P1 comparison, and retrieval-recall is additionally constant across regimes. A faithful per-pipeline cost model — one that meters the compressor pass and amortises a persistent index rather than re-charging the corpus each query — is required to attribute P3’s F1 advantage to routing rather than to uncharged cost (Section 5.6).

Frontier equivalence not established (underpowered). The cliff position was re-bootstrapped from the retained per-row `coord_success` data and a percentile-bootstrap TOST was run against the $\pm 20\%$ equivalence band (Section 4.6.3; `results/tost_corollary1.json`, `results/tost_deepseek.json`). *Neither* frontier cell passes the equivalence test: the 90% bootstrap intervals on τ^* are [1,96, 7,23] (Qwen-2.5-72B) and [1,77, 7,13] (DeepSeek V4 Pro), so although the point estimates match the reference well (Qwen-2.5-72B is $\sim 7\%$ from the canonical $\tau^* = 2,5$ and within 0,8% of the auxiliary `h1_h2_final` fit 2,6997; DeepSeek within 20%), the fits are too weakly identified to rule out a $> 20\%$ difference. Corollary 1’s frontier evidence is therefore correctly stated as “no cliff-position shift detected” rather than “invariance demonstrated”; tightening this to a passed equivalence test requires more seeds per frontier cell (more workloads to stabilise the piecewise fit under resampling), which is named as future work.

Reader-bias in H4 measurement. The Llama-3.1-8B reader under-predicts YES protected-fact answers. The pooled rates that the H4 verdict rests on are balanced by the ground-truth class distribution rather than by reader-side accuracy. The H4 metric measures “*compression preserves enough information to flip a no-biased reader to confident YES*”, which is an asymmetric leakage signal. A reader chosen specifically to be unbiased would distinguish the disclosure-reduction signal from the reader’s prior; this is named as future work in Section 5.6.

5.4. Significance

5.4.1. For Practitioners

Four practitioner-facing implications follow from the findings.

First, single-agent question-answering benchmarks (LongBench, RULER, the LLMLingua line of work’s own QA evaluations) *systematically over-report* compressor utility for multi-fragment deployments. A multi-fragment evaluation is necessary; this thesis ships one (C1) and a methodology (the cliff sweep) that other groups can apply on their own workloads.

Second, the cliff position can be predicted from per-round token recall plus a one-shot θ_q measurement, within a known first-order tolerance. A practitioner who needs to choose a compression ratio for a new multi-fragment workflow can run the θ_q measurement on a small calibration set and apply the compounding-error model to obtain a safe operating point without running the full cliff sweep.

Third, compress-first (P1) is the better RAG architecture between the two compressed-corpus pipelines in both regimes *on the single retriever/embedder stack tested* (a result that does not reproduce the LongLLMLingua post-retrieval-compression preference on this stack, but is not a general falsification). Joint relevance-conditional routing (P3) scores higher F1 still (42 pp storage, 25 pp accuracy). That advantage, however, comes from *not compressing*: Section 4.7

shows P3 routes fragments verbatim (achieved ratio 1,00×), and the corpus-dominated cost model does not meter the compressor pass. The F1 ranking $P3 \succ P1 \succ P2$ is therefore reported *without* a cost-parity claim; whether P3’s quality gain is cost-free is undetermined until the per-pipeline cost model is fixed. The defensible practitioner takeaway is narrower than a clean ranking: on this stack, prefer compress-first over retrieve-first, and treat P3 as a promising routing variant whose cost must be measured before deployment.

Fourth, compression is a privacy lever — but a *blunt* one, and the bluntness creates a policy conflict the memory bus does not yet resolve. The disclosure metric ranks compressors by leakage at iso-ratio, so a governance budget on protected-fact leakage can be enforced by routing CONFIDENTIAL fragments through the compressor (and ratio) whose disclosure rate meets the budget. The conflict is this: Section 4.8 shows the privacy gain is *destruction-driven* — leakage falls because tokens are destroyed — which is the **same lever**, pointed the other way, that the coordination cliff warns against. For a fragment that is simultaneously privacy-sensitive and coordination-critical, “compress more to leak less” and “compress less to stay above τ^* ” are directly opposed, and the disclosure-reduction and coordination-loss are not independent axes but two readings of one quantity (information destroyed). The reference bus exposes both signals (disclosure rate and CTR) but provides **no arbitration rule** for fragments where they conflict; choosing an operating point that satisfies a leakage budget *and* a coordination floor — when both are feasible at all — is named as future work (Section 5.6). Practitioners should treat the privacy lever as usable only for fragments that are *not* on the coordination-critical path, until such an arbitration policy exists.

5.4.2. For the Field

Three field-facing claims are substantive.

First, the cliff is a property of the (compressor, task) pair and of the planner *class*, not of planner scale or architecture within a class. The empirical evidence is that the LLM-planner cliff on family-a $\times$ LLMingua-2 is essentially invariant — Llama-3.1-8B at $\tau^* \approx 2,48$ and Qwen-2.5-72B at $\tau^* \approx 2,79$, a 9× scale-up across architecture families, both at $p_0 = 1,0$ — and sits well above the brittle deterministic-solver cliff ($\approx 1,1$) on the same cell. The implication for the compression literature is that future compressor evaluations should either characterise the calibrated regime they apply in or explicitly target the extended-reasoning regime as a separate case.

Second, extended-reasoning planners are a different regime that the existing compression literature has not characterised. The GPT-oss 120B out-of-regime result is a positive diagnostic of where the threshold-success assumption breaks; characterising this regime systematically is a publishable follow-on. The practical-deployment implication is that compression ratios calibrated on a non-reasoning planner cannot be transferred to a reasoning planner without revalidation.

Third, information density θ_{info} scales the cliff across benchmarks in a way that future benchmarks should control for. A multi-benchmark compression-quality comparison that does not measure or control θ_{info} will report cliff positions that conflate compressor mechanism with task structure; the cross-benchmark transfer result of Corollary 2 is the evidence that the two must be separated.

5.5. Comparison to Industry Observations

The closest published industry observation, Anthropic’s report that token usage explains approximately 80% of the performance variance in their multi-agent research workflow [13], and the follow-on context-engineering and context-editing reports [14, 15], are consistent with the findings of this thesis and underwrite the practitioner urgency. What they do not provide — and what the thesis contributes — is the structural characterisation. The Anthropic report is an aggregate observation on a single multi-agent workflow run by a single industry team; the cliff measurement in this thesis is a controlled cross-compressor, cross-family, cross-architecture sweep with a quantitative model that predicts the cliff position. The two complement each other: the industry observation is the macro reason a practitioner cares about the question; the thesis findings are the micro answer that lets them act on it.

5.6. Future Work

The most natural follow-on studies in order of expected impact.

Extended-reasoning regime. Characterise the cliff for planners that recover from sub-threshold information by chain-of-thought reasoning. The empirical setup is a re-run of the cliff sweep with GPT-oss-class planners, GPT-4-Turbo with thinking, and Claude Sonnet 4.6 with extended thinking enabled. The theoretical work is replacing A3 (threshold success) with a graded-success model.

Graded-success refinement of the A3 model. The direct A3 probe (now in the results chapter, Section 4.4.5) confirms the threshold mechanism on the dense aggregation family but shows a clearly *graded* success curve on the distributed retrieval family ($k \approx 5,9$, transition band $\approx 0,54$ wide). The remaining open theoretical direction is therefore a graded-success refinement of A3 that recovers the threshold form as the information-density-high limit. The same axis also underlies Corollary 2’s θ_{info} scaling (dense $\Rightarrow$ threshold, distributed $\Rightarrow$ graded), so the refinement is the natural common ancestor of A3 and Corollary 2.

Per-task θ_q estimation. Lift the model from per-family θ_q to per-task θ_q . The current per-family validation match rate is 33% strict in-sample (25% under leave-one-family-out cross-validation; see Section 4.4.4 for the gap); a per-task fit should improve this to the 60–70% range observed on the θ_q -already-known calibration cells, against the LOO baseline rather than the in-sample one. The empirical setup is straightforward; the methodological cost is fitting θ_q per task on a held-out calibration split.

Multi-pass compression at $N > 1$. Validate the q^N formulation of the model empirically. The empirical setup is a small sweep across two-pass and three-pass compression on the family-a cells; the analysis turns on whether the empirical cliff shifts predictably with N .

CAAC $r_{\min}$ sweep. The θ_q/N_q ablation showed that the CAAC operating-point plateau is invariant to those two parameters. The primary knob the ablation deferred is $r_{\min}$. A sweep over $r_{\min} \in \{1,2, 1,5, 2,0, 2,5\}$ produces the operating-point family that CAAC traces out as a function of the floor.

Multi-round AutoGen coordination. The AutoGen v0.4 multi-round backend is wired into the memory bus but excluded from the empirical evaluation because of run-to-run variance. A focused study that controls the round-count variance and isolates the per-round compression effect would extend the thesis findings to the multi-round coordination regime that the deterministic-solver scope of this thesis deliberately defers.

Faithful per-pipeline RAG cost model. The H3 cost ledger is corpus-dominated and does not meter compression *compute*: per-pipeline EUR differs by only 4,1–4,9% (Section 4.7.4) because a shared corpus term dominates, and P2/P3 route fragments verbatim (achieved ratio 1,00×). This prevents any cost-parity claim. A corrected model would meter each pipeline’s actual token economics separately — compression compute paid by P1 up-front, full-index synthesis tokens for P2/P3, verbatim-passthrough synthesis tokens for P3’s high-relevance branch — amortise a persistent index rather than re-charge the corpus each query, and re-evaluate the P3-vs-P1 comparison on a binding storage constraint (the current ≤ 100 MB cap is non-binding at 0,521 MB). Only then can the P3 F1 advantage be attributed to routing rather than to an unmetered cost.

Frontier equivalence: more seeds to tighten the CI. The TOST equivalence test has now been computed by re-bootstrapping τ^* from the retained per-row data (Section 4.6.3; `results/tost_corollary1.json`): *neither* frontier cell passes equivalence at $\pm 20\%$ (Qwen-2.5-72B 90% CI [1,96, 7,23], DeepSeek [1,77, 7,13]), so the defensible claim is “no shift detected”, not “equivalent”. The remaining work is to *tighten* the test: increasing the per-cell frontier sample (more workloads) would stabilise the piecewise fit under resampling — the DeepSeek bootstrap median 5,7 sitting far from its point estimate 2,15 is the signature of the instability — and could convert “no shift detected” into a passed equivalence test (or refute it).

Memory-bus systems benchmark. C4’s single-threaded core operations are now measured on two hosts — an Apple M4 Pro (arm64) and a Linux x86_64 workstation. The policy check runs at ~ 5 – 7 million ops/s, and end-to-end write/read at $\sim 18,200/\sim 19,900$ ops/s on the M4 Pro ($\sim 5,750/\sim 6,400$ on the Linux host), both with sub-millisecond tail latency. The audit hash-chain verifies at 1,68–2,19 μ s/row (Section 4.8.9 of Chapter 4; `scripts/bench_memory_bus.py`, `results/bus_bench/bench.json` and `results/bus_bench/bench_gpu.json`). The benchmark uses the identity compressor (isolating the bus machinery) and the in-process SQLite audit log, so the remaining step that would upgrade C4 from “characterised single-threaded” to “production characterised” is a concurrent-HTTP-path throughput sweep and an audit-verify scaling curve across chain lengths.

Standalone dissemination of the H1 result. The H1 finding — that single-agent question-answering accuracy does not transferably predict multi-fragment coordination success, so the QA benchmarks dominating the compression literature mis-rank compressors for multi-fragment use — is the most citable standalone contribution and is independent of the model and the bus. A focused workshop paper built around H1 (the per-family Spearman decomposition plus the cliff existence result) is the lowest-friction route to external validation of the thesis’s central methodological claim.

Broader task families. The three C1 families are deliberate probes of a single property each (density, constraint-tracking, chain-reference). A broader benchmark suite with

heterogeneous aggregation types, multi-tool reasoning, and multimodal fragments would test the generalisability of the cliff structure beyond the families this thesis evaluates.

An unbiased H4 reader — partially addressed by class re-weighting. The cheaper of the two options named here, re-weighting the H4 results by the reader-side prior, was carried out: the per-question results were re-balanced so the YES and NO ground-truth classes contribute equally, removing the reader’s YES-bias from the aggregate. The disclosure-reduction conclusion is **robust** to this de-biasing — class-balanced reductions (truncation 0,247, filter 0,197, LLMLingua-2 0,188, Phi-3-Mini extractive 0,005 n.s.) are within a percentage point of the pooled values, and the destruction-monotonic ordering is unchanged. The Section 4.8.7 caveat is thus downgraded from “the verdict may rest on a reader artefact” to “the verdict survives class-balancing.” The remaining open item is the larger-reader confirmation (e.g. Llama-3.3-70B), expected to be confirmatory rather than corrective.

5.7. Closing

The five empirical findings, the compounding-error model that ties them together, the memory-bus reference implementation that operationalises the privacy lever, and the CAAC algorithm that operationalises the cliff bound, jointly constitute the thesis’s contribution: a structural characterisation of context compression in multi-fragment LLM workflows, validated across compressors, families, planner scales, architectures, and benchmarks, with a predictive model, a deployment artefact, and an open-source reproducibility package. The single deliverable sentence is the one the abstract opens with: every token that does not change the answer is waste, and what this thesis contributes is a way of knowing which tokens those are, when removing them is safe, and what the cost is when removing them is not.

6. REFERENCES

- [1] A. S. Luccioni, Y. Jernite, and E. Strubell, “Power hungry processing: Watts driving the cost of AI deployment?”, in *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency (FAccT ’24)*, ACM, 2024, pp. 85–99. doi: 10.1145/3630106.3658542. [Online]. Available: <https://dl.acm.org/doi/10.1145/3630106.3658542>.
- [2] S. Samsi et al., “From words to watts: Benchmarking the energy costs of large language model inference”, in *2023 IEEE High Performance Extreme Computing Conference (HPEC)*, IEEE, 2023. doi: 10.1109/HPEC58863.2023.10363447.
- [3] R. Pope et al., “Efficiently scaling transformer inference”, in *Proceedings of Machine Learning and Systems (MLSys)*, vol. 5, 2023.
- [4] D. Patterson et al., *Carbon emissions and large neural network training*, arXiv preprint arXiv:2104.10350, 2021.
- [5] S. Altman, *Tens of millions of dollars on please-and-thank-you compute*, Industry reporting on social media (X/Twitter), April 2025, Cited as industry observation; not a peer-reviewed source., 2025.
- [6] H. Jiang, Q. Wu, C.-Y. Lin, Y. Yang, and L. Qiu, “LLMLingua: Compressing prompts for accelerated inference of large language models”, in *Proc. Conf. Empirical Methods in Natural Language Processing (EMNLP)*, 2023.
- [7] H. Jiang et al., “LongLLMLingua: Accelerating and enhancing LLMs in long context scenarios via prompt compression”, in *Proc. Annual Meeting of the Association for Computational Linguistics (ACL)*, 2024.
- [8] Z. Pan et al., “LLMLingua-2: Data distillation for efficient and faithful task-agnostic prompt compression”, in *Findings of the Association for Computational Linguistics (ACL Findings)*, 2024.
- [9] J. Mu, X. L. Li, and N. Goodman, “Learning to compress prompts with gist tokens”, in *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- [10] A. Chevalier, A. Wettig, A. Ajith, and D. Chen, “Adapting language models to compress contexts”, in *Proc. Conf. Empirical Methods in Natural Language Processing (EMNLP)*, AutoCompressor., 2023.
- [11] T. Ge, J. Hu, L. Wang, X. Wang, S.-Q. Chen, and F. Wei, “In-context autoencoder for context compression in a large language model”, in *Proc. International Conference on Learning Representations (ICLR)*, 2024.
- [12] Future Computing Group, “Use case: A week at the university in the AI service economy”, CSE Research Unit, University of Oulu, Tech. Rep., 2026, Internal document, March 2026. Vignette 3.7 (Period-end project reporting) is the focal scenario.
- [13] Anthropic. “How we built our multi-agent research system”, Anthropic Engineering Blog, Accessed: May 12, 2026. [Online]. Available: <https://www.anthropic.com/engineering/multi-agent-research-system>.

- [14] Anthropic. “Managing context on the Claude developer platform: Context editing and the memory tool”, Anthropic News, Accessed: May 12, 2026. [Online]. Available: <https://www.anthropic.com/news/context-editing-memory-tool>.
- [15] P. Rajasekaran, K. Dixon, J. Ryan, A. Hadfield, et al. “Effective context engineering for AI agents”, Anthropic Applied AI Engineering Blog, Accessed: May 12, 2026. [Online]. Available: <https://www.anthropic.com/engineering/context-engineering-for-ai-agents>.
- [16] Q. Wu et al., *AutoGen: Enabling next-gen LLM applications via multi-agent conversation*, arXiv preprint arXiv:2308.08155, 2023.
- [17] Y. Bai, S. Tu, J. Zhang, et al., *LongBench v2: Towards deeper understanding and reasoning on realistic long-context multitasks*, arXiv preprint arXiv:2412.15204, 2024.
- [18] C.-P. Hsieh et al., *RULER: What’s the real context size of your long-context language models?*, arXiv preprint arXiv:2404.06654, 2024.
- [19] Zhou et al., *Privacy-Aware retrieval-augmented generation*, arXiv preprint arXiv:2503.15548, 2025.
- [20] Bassit and Boddeti, *SecureRAG: Privacy-preserving retrieval-augmented generation*, arXiv preprint, 2025.
- [21] A. Vaswani et al., “Attention is all you need”, in *Advances in Neural Information Processing Systems 30 (NIPS 2017)*, 2017, pp. 5998–6008. [Online]. Available: <https://proceedings.neurips.cc/paper/2017/hash/3f5ee243547dee91fbd053c1c4a845aa-Abstract.html>.
- [22] J. Kaplan et al., “Scaling laws for neural language models”, *arXiv preprint arXiv:2001.08361*, 2020, arXiv:2001.08361, preprint.
- [23] J. Hoffmann et al., “Training compute-optimal large language models”, in *Advances in Neural Information Processing Systems 35 (NeurIPS 2022)*, 2022. [Online]. Available: https://proceedings.neurips.cc/paper_files/paper/2022/hash/c1e2fa ff6f588870935f114ebe04a3e5-Abstract-Conference.html.
- [24] N. F. Liu et al., “Lost in the middle: How language models use long contexts”, *Transactions of the Association for Computational Linguistics (TACL)*, vol. 12, pp. 157–173, 2024.
- [25] Z. Li, Y. Liu, Y. Su, and N. Collier, “Prompt compression for large language models: A survey”, in *Proc. Conf. North American Chapter of the Association for Computational Linguistics (NAACL)*, 2025.
- [26] Cheng et al., *xRAG: Extreme context compression for retrieval-augmented generation with one token*, arXiv preprint arXiv:2405.13792, 2024.
- [27] T. Kočiský et al., “The NarrativeQA reading comprehension challenge”, *Transactions of the Association for Computational Linguistics (TACL)*, vol. 6, pp. 317–328, 2018.
- [28] Z. Yang et al., “HotpotQA: A dataset for diverse, explainable multi-hop question answering”, in *Proc. Conf. Empirical Methods in Natural Language Processing (EMNLP)*, 2018.

- [29] Y. Li, B. Dong, F. Guerin, and C. Lin, “Compressing context to enhance inference efficiency of large language models”, in *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, Association for Computational Linguistics, 2023.
- [30] F. Xu, W. Shi, and E. Choi, “RECOMP: Improving retrieval-augmented LMs with compression and selective augmentation”, in *Proc. International Conference on Learning Representations (ICLR)*, 2024.
- [31] Wang et al., “In-context former: Lightning-fast compressing context for large language model”, in *Proc. Conf. Empirical Methods in Natural Language Processing (EMNLP)*, 2024.
- [32] J. W. Rae, A. Potapenko, S. M. Jayakumar, C. Hillier, and T. P. Lillicrap, “Compressive transformers for long-range sequence modelling”, in *Proc. International Conference on Learning Representations (ICLR)*, 2020.
- [33] M. Abdin et al., “Phi-3 technical report: A highly capable language model locally on your phone”, Microsoft, Tech. Rep., 2024, arXiv:2404.14219, technical report. [Online]. Available: <https://arxiv.org/abs/2404.14219>.
- [34] S. Hong et al., “MetaGPT: Meta programming for a multi-agent collaborative framework”, in *Proc. International Conference on Learning Representations (ICLR)*, 2024.
- [35] G. Li, H. A. A. K. Hammoud, H. Itani, D. Khizbullin, and B. Ghanem, *CAMEL: Communicative agents for “mind” exploration of large scale language model society*, arXiv preprint arXiv:2303.17760, 2023.
- [36] N. Shinn, F. Cassano, E. Berman, A. Gopinath, K. Narasimhan, and S. Yao, “Reflexion: Language agents with verbal reinforcement learning”, in *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- [37] C. Packer et al., *MemGPT: Towards LLMs as operating systems*, arXiv preprint arXiv:2310.08560, 2023.
- [38] Park et al., *Collaborative memory: Multi-user memory sharing in LLM agents with dynamic access control*, arXiv preprint arXiv:2505.18279, 2025.
- [39] Xu et al., “A-MEM: Agentic memory for LLM agents”, in *Advances in Neural Information Processing Systems (NeurIPS)*, 2025.
- [40] Chhikara et al., *Mem0: Building production-ready AI agents with scalable long-term memory*, arXiv preprint arXiv:2504.19413, 2025.
- [41] Rasmussen et al., *Zep: A temporal knowledge graph architecture for agent memory*, arXiv preprint arXiv:2501.13956, 2025.
- [42] Saleh et al., “MemIndex: Agentic event-based distributed memory management for large language model agents”, *ACM Transactions on Autonomous and Adaptive Systems (TAAS)*, 2025.
- [43] A. Saleh, R. Morabito, S. Dustdar, S. Tarkoma, S. Pirttikangas, and L. Lovén, “Towards message brokers for generative AI: Survey, challenges, and opportunities”, *ACM Computing Surveys*, vol. 58, no. 1, Article 20, 2025. doi: 10.1145/3742891.
- [44] P. Lewis et al., “Retrieval-augmented generation for knowledge-intensive NLP tasks”, in *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.

- [45] P. Sarthi, S. Abdullah, A. Tuli, S. Khanna, A. Goldie, and C. D. Manning, “RAPTOR: Recursive abstractive processing for tree-organized retrieval”, in *Proc. International Conference on Learning Representations (ICLR)*, 2024.
- [46] D. Edge et al., *From local to global: A GraphRAG approach to query-focused summarization*, arXiv preprint arXiv:2404.16130, 2024.
- [47] B. J. Gutiérrez, Y. Shu, Y. Gu, M. Yasunaga, and Y. Su, “HippoRAG: Neurobiologically inspired long-term memory for large language models”, in *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.
- [48] B. J. Gutiérrez, Y. Shu, W. Qi, S. Zhou, and Y. Su, “From RAG to memory: Non-parametric continual learning for large language models”, in *Proc. International Conference on Machine Learning (ICML)*, 2025.
- [49] A. Asai, Z. Wu, Y. Wang, A. Sil, and H. Hajishirzi, “Self-RAG: Learning to retrieve, generate, and critique through self-reflection”, in *Proc. International Conference on Learning Representations (ICLR)*, 2024.
- [50] S. Guo, S. Zhang, and Z. Ren, *Enhancing RAG efficiency with adaptive context compression*, arXiv preprint arXiv:2507.22931, 2025. DOI: 10.48550/arXiv.2507.22931.
- [51] G. Mialon, C. Fourrier, C. Swift, T. Wolf, Y. LeCun, and T. Scialom, *GAIA: A benchmark for general AI assistants*, arXiv preprint arXiv:2311.12983, 2023.
- [52] X. Liu, H. Yu, H. Zhang, Y. Xu, et al., “AgentBench: Evaluating LLMs as agents”, in *Proc. International Conference on Learning Representations (ICLR)*, 2024.
- [53] H. Yen et al., “HELMET: How to evaluate long-context language models effectively and thoroughly”, in *Proc. International Conference on Learning Representations (ICLR)*, 2025.
- [54] X. Ho, A.-K. Duong Nguyen, S. Sugawara, and A. Aizawa, “Constructing a multi-hop QA dataset for comprehensive evaluation of reasoning steps”, in *Proc. 28th International Conference on Computational Linguistics (COLING)*, 2020.
- [55] Y. Tang and Y. Yang, “MultiHopRAG: Benchmarking retrieval-augmented generation for multi-hop queries”, in *Findings of the Association for Computational Linguistics: EMNLP 2024*, 2024.
- [56] Addison et al., *C-FedRAG: A confidential federated retrieval-augmented generation system*, arXiv preprint, 2024.
- [57] National Institute of Standards and Technology, “AI risk management framework (AI RMF 1.0)”, U.S. Department of Commerce, Tech. Rep., 2023.
- [58] Future Computing Group, “Financial analysis: University AI service economy”, CSE Research Unit, University of Oulu, Tech. Rep., 2026, Internal document, March 2026. Five-year projection across the University of Oulu.
- [59] Anthropic. “Introducing contextual retrieval”, Anthropic News, Accessed: May 12, 2026. [Online]. Available: <https://www.anthropic.com/news/contextual-retrieval>.
- [60] Ollama. “Ollama: Get up and running with large language models locally”, Accessed: May 12, 2026. [Online]. Available: <https://ollama.com/>.

- [61] B. Efron and R. J. Tibshirani, *An Introduction to the Bootstrap*. New York: Chapman & Hall, 1993.
- [62] F. Wilcoxon, “Individual comparisons by ranking methods”, *Biometrics Bulletin*, vol. 1, no. 6, pp. 80–83, 1945.
- [63] H. B. Mann and D. R. Whitney, “On a test of whether one of two random variables is stochastically larger than the other”, *Annals of Mathematical Statistics*, vol. 18, no. 1, pp. 50–60, 1947.
- [64] S. Holm, “A simple sequentially rejective multiple test procedure”, *Scandinavian Journal of Statistics*, vol. 6, no. 2, pp. 65–70, 1979.
- [65] J. Pineau et al., “Improving reproducibility in machine learning research (a report from the NeurIPS 2019 reproducibility program)”, *Journal of Machine Learning Research*, vol. 22, no. 164, pp. 1–20, 2021.

7. APPENDICES

The appendices collect four artefacts that support the main text but would interrupt its narrative: the memory-bus HTTP contract (Appendix A), the canonical long-format results-dataframe schema used by every experiment runner (Appendix B), an example C1 workload instance from family (a) showing the source-fragment structure (Appendix C), and a one-command reproducibility recipe together with a verbatim `curl` trace exercising the memory-bus endpoints (Appendix D). The full reproducibility package, including the model cards and the data cards, is shipped with the open-source release referenced in the closing section of Chapter 5.

7.1. Appendix A. Memory-Bus HTTP Contract

The reference memory-bus service exposes four endpoints. The contract below summarises the request and response shapes; the authoritative OpenAPI schema lives at `docs/openapi.json` in the repository and is regenerated via `make bus-openapi`.

POST /v1/write Body: a Fragment JSON object (`fragment_id`, `text`, `tags`) and an optional `target_ratio`. Returns the assigned `slot_id`, the audit row identifier, and the hex-encoded `chain_hash`. HTTP 201 on success; 403 on policy denial.

GET /v1/read/{slot_id} Returns a CompressedSlot with its payload, tags, compressor identifier, nominal ratio, and creation timestamp. HTTP 200 on success; 403 on policy denial; 404 on slot not found.

POST /v1/subscribe Body: a SubscribeRequest with a query string, a `ttl_seconds` positive integer, and a top-*k* integer. Returns a server-sent-event stream emitting newly-matched `slot_ids` with their cosine scores until the TTL elapses.

GET /v1/audit/{slot_id} Returns the chronological list of AuditRow entries that mention the given slot.

Every request is processed by the PolicyMiddleware which parses the requester principal from the X-M6-Principal header in the development build and from an OAuth/JWT verifier in a production deployment. The principal carries a 64-bit access-control mask and a 5-tier classification level (PUBLIC < INTERNAL < CONFIDENTIAL < RESTRICTED < SECRET). A request is granted access to a slot when the principal's mask is a superset of the slot tag's mask and the principal's classification is at least the slot tag's classification.

7.2. Appendix B. Long-Format Results Schema

Every experiment runner writes its results to a CSV that conforms to the long-format schema implemented as `m6.evaluation.statistics.LongDF`. The columns are listed in Table 14.

Table 14: The canonical long-format dataframe schema used by every experiment runner. The same schema underlies every CSV under `results/`, which makes cross-hypothesis analysis a single `pandas.concat` away.

Column	Type	Notes
<code>experiment_id</code>	str	unique per run
<code>hypothesis</code>	str	one of <code>h1...h6</code> or <code>caac</code> , <code>frontier</code> , <code>hotpotqa</code>
<code>compressor</code>	str	e.g. <code>lingua2</code> , <code>filter</code> , <code>phi3-extractive</code> , <code>truncation</code> , <code>caac</code> , <code>identity</code>
<code>ratio</code>	float	nominal compression ratio
<code>actual_ratio</code>	float	measured ratio post-compression
<code>pipeline</code>	str	one of <code>none</code> , <code>P1</code> , <code>P2</code> , <code>P3</code>
<code>model</code>	str	e.g. <code>llama-3.1-8b-instruct</code>
<code>model_size</code>	str	<code>1.5b</code> , <code>3.8b</code> , <code>8b</code> , <code>72b</code>
<code>workload_family</code>	str	<code>a</code> , <code>b</code> , or <code>c</code>
<code>workload_id</code>	str	e.g. <code>c1-a-007</code>
<code>seed</code>	int	
<code>metric</code>	str	e.g. <code>coord_success</code> , <code>qa_f1</code> , <code>preservation_rate</code>
<code>value</code>	float	
<code>wallclock_ms</code>	int	
<code>eur_cost</code>	float	
<code>run_id</code>	str	
<code>git_sha</code>	str	
<code>config_hash</code>	str	
<code>created_at</code>	str	ISO-8601 UTC
<code>invalid</code>	bool	control-condition variance dominates?
<code>invalid_reason</code>	str	free-text explanation if invalid

7.3. Appendix C. Example C1 Workload (Family a)

The following JSON object illustrates a single instance of C1 family (a) (cross-document fact aggregation). Each fragment represents one institutional system; the planner is asked to aggregate hours and budget across all eight; the critic verifies the answer against the ground-truth aggregate.

```
{
  "workload_id": "c1-a-007",
  "family": "a",
  "seed": 1247,
  "tag_distribution": "skewed",
  "initial_prompt":
    "Aggregate the period-end report for project P-3219 across the
     following systems: publication_db, contract_db, moodle, patio,
     peppi, crm, tatu_sap, sap_travel. Report total hours and total
     budget.",
  "expected_answer": "hours=853;budget=148340",
}
```

```

    "n_agents": 8,
    "fragments": [
      {
        "fragment_id": "c1-a-007/publication_db",
        "text": "[publication_db] Project P-3219 - period 2025-Q3. ...",
        "tags": {"acl_mask": 12876, "classification": 1}
      },
      {
        "fragment_id": "c1-a-007/contract_db",
        "text": "[contract_db] Project P-3219 - period 2025-Q3. ...",
        "tags": {"acl_mask": 8901, "classification": 2}
      }
    ]
    // (six more fragments, one per system)
  ],
  "sub_tasks": [
    {
      "sub_task_id": "c1-a-007/sub-0",
      "description":
        "Extract recorded hours and budget for project P-3219 from
        publication_db.",
      "expected_solver": "worker-0",
      "expected_answer": "hours=112;budget=18420"
    }
  ]
  // (seven more sub-tasks)
],
"protected_facts": [
  {
    "fragment_id": "c1-a-007/contract_db",
    "fact": "contract_db.budget=22150",
    "yesno_questions": [
      "Did contract_db for project P-3219 exceed EUR 27150 in budget?",
      "Was contract_db's recorded hours for P-3219 more than 142?"
    ],
    "answers": ["no", "no"]
  }
]
}

```

The full dataset is reproduced from `configs/benchmark/c1-v0.1.yaml` via `make bench-generate`; the manifest carries the configuration hash so any reader can verify they are looking at the same release.

7.4. Appendix D. Reproducibility Recipe and Memory-Bus Trace

7.4.1. D.1 One-Command Reproduction of Every Chapter Figure

The reproducibility package is the complete open-source repository. The compute envelope is a single Apple M4 Pro 48 GB workstation, with an optional RTX 5090 32 GB workstation (Tailscale-accessible WSL2 host) for the GPU-bound compression precomputation, the model-scaling sweep at 8B, the frontier-API sweep, and the HotpotQA / MultiHopRAG transfer arm. Total wallclock end-to-end is approximately 30 hours of GPU time and ~ 12 hours of laptop time. The package contains a `docker-compose.yml` that brings the memory-bus reference service up locally, a GitHub release tag matched to the manuscript, model and data cards under `docs/`, and a `README.md` that exposes one-command reproduction targets for every headline figure:

make repro-bench regenerates the 150-instance C1 benchmark from `configs/benchmark/c1-v0.1.yaml`. Output: `data/processed/c1-v0.1/`.

make repro-cache runs the compression precomputation for all four compressors at the ten canonical ratios across the 150 workloads. Requires the GPU host; total wallclock ~ 14 hours. Output: `results/compression_cache/`.

make repro-cliff runs the H1/H2 cliff sweep against the compression cache, producing all of Chapter 4’s Section 4.2 through Section 4.4.4 figures. Wallclock ~ 3 hours on the M4 Pro.

make repro-scaling runs the model-scaling sweep at 1,5B, 3,8B, and 8B with the local Ollama backend, producing Corollary 1’s figures. Wallclock ~ 4 hours on the GPU host.

make repro-frontier runs the cliff sweep against the Featherless API for Qwen-2.5-72B-Instruct and DeepSeek V4 Pro, the two calibrated-regime frontier validators. (The out-of-regime GPT-oss-120B diagnostic is scoped out of the canonical recipe as an extended-reasoning planner and is not part of this target.) Requires a `FEATHERLESS_API_KEY`. Output: `results/frontier_{qwen72b,deepseekv4}/`.

make repro-rag runs the H3 RAG pipeline catalogue, producing Section 4.7’s figures. Wallclock ~ 3 hours.

make repro-disclosure runs the H4 inference-disclosure measurement against the unbiased benchmark, producing Section 4.8’s figures. Wallclock ~ 1 hour on the M4 Pro.

make repro-transfer runs the HotpotQA and MultiHopRAG transfer arm, producing Section 4.9’s figures. Wallclock ~ 2 hours.

make repro-figs regenerates every PDF and PNG figure under `figures/` from the canonical CSVs. Pure-Python; runs in under a minute.

The `make repro-all` target chains the six laptop-runnable steps above — `repro-cliff`, `repro-scaling`, `repro-rag`, `repro-disclosure`, `repro-transfer`, and `repro-figs` — in dependency order. It deliberately excludes the two steps that need special resources: `repro-cache` (GPU host) and `repro-frontier` (a Featherless API key), which must be run explicitly. Running the entire package including those two steps is approximately 30 hours of GPU time and ~ 12 hours of laptop time end-to-end. Which of these targets is byte-deterministic, which reproduces only in distribution, and which is not byte-reproducible is set

out in the three-tier guarantee of Section 4.10 (Chapter 4); the `repro-cache`, `repro-cliff`, `repro-rag`, `repro-disclosure`, and `repro-figs` targets are Tier 1 (byte-identical CSVs), `repro-scaling` is Tier 2, and `repro-frontier` is Tier 3.

7.4.2. D.2 Memory-Bus End-To-End Trace

The following `curl` sequence exercises the four endpoints of the memory bus end-to-end against the `docker-compose-up`'ed reference service. The trace is the rubric-#5 evidence for the C4 contribution: the memory bus is a running system, not a paper design. The output of each step is reproduced verbatim from a run against the reference image tagged `m6-memory-bus:thesis-v1.0`.

```
# Step 1: Write a CONFIDENTIAL fragment with target ratio 4x.
$ curl -X POST http://localhost:8000/v1/write \
  -H "Content-Type: application/json" \
  -H "X-M6-Principal: principal=alice;acl=0xFFFF;cls=3" \
  -d @- <<'EOF'
{
  "fragment_id": "demo-001",
  "text": "Project P-3219 quarterly hours: 112; budget: EUR 18420.",
  "tags": {"acl_mask": 12876, "classification": 2},
  "target_ratio": 4.0,
  "compressor": "lingua2"
}
EOF
{"slot_id": "slot-1a4b9c2d", "audit_row": 1,
  "chain_hash": "5f3d2a1b...", "achieved_ratio": 3.72}

# Step 2: Read the slot back. The principal's classification (3)
#         is at least the slot's (2) and the principal's ACL mask
#         is a superset, so the read is permitted.
$ curl -H "X-M6-Principal: principal=alice;acl=0xFFFF;cls=3" \
  http://localhost:8000/v1/read/slot-1a4b9c2d
{"slot_id": "slot-1a4b9c2d",
  "compressed_text": "Project P-3219 hours: 112 budget: 18420.",
  "compressor": "lingua2", "achieved_ratio": 3.72,
  "tags": {"acl_mask": 12876, "classification": 2},
  "created_at": "2026-05-30T07:42:15Z"}

# Step 3: A second principal with insufficient classification is
#         denied; the denial is itself audited with a
#         per-(principal, slot) 60-second deduplication.
$ curl -H "X-M6-Principal: principal=bob;acl=0xFFFF;cls=1" \
  -o /dev/null -w "%{http_code}\n" \
  http://localhost:8000/v1/read/slot-1a4b9c2d
```

```
# Step 4: Audit walk for the slot. The chain is verifiable end-to-
#         end via the SHA-256 chain hash on each row; tampering
#         with any row breaks verification.
$ curl -H "X-M6-Principal: principal=alice;acl=0xFFFF;cls=3" \
      http://localhost:8000/v1/audit/slot-1a4b9c2d
[
  {"rowid": 1, "event_type": "WRITE", "result": "OK",
   "requester": "alice", "chain_hash": "5f3d2a1b..."},
  {"rowid": 2, "event_type": "READ", "result": "OK",
   "requester": "alice", "chain_hash": "9c4e7a8d..."},
  {"rowid": 3, "event_type": "READ", "result": "DENIED",
   "requester": "bob", "chain_hash": "1b8a6f3e..."}
]
```

The trace illustrates four properties: compression with achieved- ratio reporting on the write path; policy-checked read with the five-tier classification lattice and the access-control mask; audited denial with the per-principal-per-slot deduplication that prevents audit-log flooding; and chain-hash-verifiable provenance on the audit endpoint. The complete OpenAPI schema is regenerated by `make bus-openapi` and lives at `docs/openapi.json` in the repository.

7.4.3. D.3 Compute Envelope and Software Stack

The empirical work was carried out on the following stack. The local workstation is an Apple Mac with the M4 Pro SoC, 48 GB of unified memory, Python 3.12, MLX framework version 0.21, Ollama version 0.3 with the Qwen-2.5, Phi-3-Mini, and Llama-3.1 model weights pre-pulled, and the `m6` package installed in editable mode from the repository. The GPU host is a Windows desktop running WSL2 Ubuntu 22.04 with the NVIDIA RTX 5090 32 GB GPU, CUDA toolkit ≥ 12.0 , and the same Python and Ollama versions. The Tailscale mesh network connects the local workstation to the GPU host as host `gpu` on the `100.70.160.59/32` private address. The memory-bus reference service runs on either host via `docker compose -f docker/docker-compose.yml up -d`; the image is Python 3.12-slim with FastAPI, Uvicorn, SQLite, and the `m6` package.